

★ ★ ★ ★ ★ ★ ★ ★ ★ ★ ★ ★

ACCOUNTABLE AI DEPLOYMENT ACT OF 2026

119th Congress · H.R. ____ / S. ____ · Advocacy Submission

*"Governance is not what you review after execution —
it is what determines whether execution is possible at all."*

— AADA Core Compliance Axiom | Jeremy Karrick, 2026

CONGRESS	STATUS	EFFECTIVE	RESEARCH BASIS
119th (2025–2026)	Advocacy Submission	6–18 months post-enactment	21 peer-reviewed publications (Zenodo)

Provisions Closing 14 Critical Governance Loopholes:

HITL Exclusion · Partial Commit · Side-Channel · Zero Trust Inbound · Drift Bounding · Suppression Proof · Replay Equivalence ·
Regulator Latency · Tool Simulation · Prompt State · Update Equivalence · Multi-Tenant · Artifact Governance · Fallback-to-Deny

★ ★ ★ ★ ★ ★ ★ ★ ★ ★
AUTHOR: Jeremy Karrick | karrick1995@gmail.com | ORCID: 0009-0007-9715-4290

Principal Architect — Project-AI Sovereign Constitutional AGI Ecosystem
Research Basis: 21 Peer-Reviewed Publications | Zenodo — Thirstys Projects Community

© 2026 Jeremy Karrick — All Rights Reserved

TABLE OF CONTENTS

CONTENTS OF THIS LEGISLATIVE PACKAGE

PART I	Full Bill Text — Sections 1–18	3
SEC. 1–2	Short Title & Findings	3
SEC. 3	Definitions (25 Defined Terms)	3
SEC. 4	Non-Delegable Responsibility	5
SEC. 5	Required Governance Controls	5
SEC. 5A–5D	Architectural Neutrality; Model Update Equivalence; Multi-Tenant Isolation; Artifact Governance	10
SEC. 6	Risk Tier Classification & Cross-Tier Contamination	11
SEC. 7	Continuous Compliance & Behavioral Drift Bounding	12
SEC. 8	Incident Reporting & Private Right of Action	12
SEC. 9	Supply Chain Accountability	13
SEC. 10	Prohibited Practices (19 Enumerated)	13
SEC. 11	Enforcement & Compliance Credit	14
SEC. 12	Federal AI Governance Office (FAIGO)	14
SEC. 13	Safe Harbor	15
SEC. 14	Research & Open Source Exemption	15
SEC. 15	Public Transparency & Independent Oversight	16
SEC. 15A	Regulatory Sandbox & Delta Documentation	16
SEC. 16	Technical Standards (GOV/ASS/AUD Framework)	17
SEC. 17	Preemption & Federalism	19
SEC. 18	Effective Date & Transition	19
PART II	Technical Addenda — 14 Loopholes Closed	20
PART III	NIST-Aligned Regulatory Implementation Framework	24
PART IV	Congressional Sponsorship Brief	27
PART V	Legislative Campaign Strategy	29
APPENDIX A	Deployer Compliance Checklist (By Tier)	32
APPENDIX B	AADA vs. Other Frameworks Comparison	33
APPENDIX C	Anticipated Objections & Responses	34
CLOSING	Statement of Purpose	35

EXECUTIVE SUMMARY

PLAIN-LANGUAGE OVERVIEW FOR SPONSORS AND STAFF

What this bill does and why it matters

The Core Problem

AI systems are making consequential decisions about Americans' healthcare, employment, finances, and freedom. When those systems cause harm, accountability is diffuse — spread across model developers, platform operators, and downstream deployers — and current law provides no clear liability standard, no mandatory audit trail, and no enforceable technical floor.

Existing AI governance approaches share a fatal structural flaw: controls exist in policy documents, training guidelines, and acceptable-use terms. They do not exist in the execution path. The AADA closes that gap by requiring that accountability be **built into AI systems from the ground up**.

What the Bill Requires

- **Governance by design, not by policy.** Companies cannot satisfy obligations with an ethics committee, an acceptable-use policy, or a training curriculum. Enforcement must be technically embedded — the system literally cannot take unauthorized actions.
- **Non-delegable accountability.** The entity that deploys an AI system is responsible for its outcomes. That responsibility cannot be contracted away to vendors, blamed on emergent behavior, or discharged through terms of service.
- **Replay-equivalent auditability.** Every consequential AI action must be logged such that the exact decision can be *identically reconstructed* — not approximated. Regulators can trace exactly what the system did, why, and who authorized it.
- **Human review is not a control layer.** The bill closes the most commonly exploited loophole: claiming that human oversight constitutes governance. It does not, unless the human approval is architecturally enforced as a blocking precondition.
- **Fail-safes dominate all other objectives.** Uptime requirements, SLA obligations, and commercial objectives cannot override fail-safe activation. Safety is not a feature to be balanced against availability. It is a hard constraint.

Who Is Responsible for What

TIER	WHO IT COVERS	KEY REQUIREMENTS
Tier 1 Low Impact	Internal tools, minimal external harm potential	Baseline controls · Annual self-certification Free certification (under \$5M revenue)

Tier 2 Moderate Impact	Consumer services, finance, employment systems	Full governance controls · Biennial 3rd-party audit Scaled penalties up to \$5M/violation
Tier 3 High Impact	Healthcare, legal, infrastructure, autonomous physical systems	Full controls + Shadow Execution · Annual cert Penalties up to \$50M or 4% global revenue

Bipartisan Case

CONSERVATIVE FRAMING	PROGRESSIVE FRAMING
Corporate accountability — if you deploy it, you own it	Consumer protection from algorithmic harm
Rule of law extended to AI actors	Civil rights in high-stakes automated decisions
No new regulatory capture — FAIGO + FTC only	Strong enforcement with private right of action
Safe harbor rewards voluntary compliance	Meaningful penalties up to 4% global revenue (Tier 3)
Architectural neutrality — no mandate on method	Human oversight cannot substitute for real governance

PART I

FULL BILL TEXT

119th Congress · H.R. ____ / S. ____ · All Sections 1–18

SECTION 1. SHORT TITLE.

This Act may be cited as the "**Accountable AI Deployment Act of 2026**" (AADA).

SECTION 2. FINDINGS AND PURPOSE.

(a) **Findings.** Congress finds that:

- (1) Artificial intelligence systems are increasingly deployed in contexts with material impact on individuals, institutions, critical infrastructure, and public welfare.
- (2) Existing governance practices rely heavily on policy documents, training regimens, and post-hoc review rather than enforceable, embedded controls — producing the appearance of accountability without its substance.
- (3) Harm arising from AI systems is frequently obscured by diffuse responsibility distributed across model vendors, platform deployers, and downstream operators, leaving affected parties without legal recourse.
- (4) Effective governance requires constraints embedded within system execution pathways — not solely in external documentation, contractual terms, or voluntary commitments.
- (5) The probabilistic and emergent properties of AI systems do not exempt those who deploy them from responsibility for outcomes. Technical complexity is not a shield against accountability.
- (6) The national interest is served by governance standards that are outcome-based and technology-neutral, so as not to entrench incumbent architectures or impede innovation while establishing a clear floor of enforceable safety.
- (7) Public trust in AI systems requires not only that governance controls exist, but that evidence of their operation is available to qualified independent scrutiny.
- (8) Modern AI deployment increasingly involves chained, agentic, and multi-model architectures in which governance failures at any boundary can propagate to produce material harm, and in which individually authorized steps can compose into unauthorized outcomes.
- (9) Human review, approval queues, and oversight processes that are not technically enforced as preconditions for execution produce the appearance of governance without its substance. Genuine accountability requires that such controls be embedded in the execution architecture.
- (10) The increasing use of AI systems to generate code, configurations, and executable artifacts — subsequently executed by other systems — constitutes an indirect but material execution pathway that falls within the scope of enforceable governance obligations.

(b) Purpose. To require that AI Systems deployed within the United States operate under enforceable, auditable, and non-bypassable Governance Controls, with clear and non-delegable responsibility for outcomes, administered through a designated federal oversight infrastructure, while preserving conditions for responsible innovation.

SECTION 3. DEFINITIONS.

As used in this Act:

(1) ARTIFICIAL INTELLIGENCE SYSTEM. Any computational system that — through machine learning, neural networks, probabilistic reasoning, generative models, or functionally equivalent methods — produces outputs, recommendations, or actions capable of materially influencing decisions, environments, or persons. Systems that incorporate any learned weights, parameters, or probabilistic components in any decision path remain subject to this Act, even if the majority of system logic is explicitly programmed. The presence of a rule-based shell or symbolic wrapper does not exempt an underlying learned component from coverage.

Exclusion (a): Systems whose behavior is fully determined by explicitly programmed, human-readable rules without learned weights or parameters anywhere in any decision path, and whose complete decision logic can be exhaustively enumerated and verified by inspection.

Exclusion (b): Systems operating exclusively in non-networked, personal, non-commercial, or academic research contexts as defined in Section 14.

Exclusion (c): Systems with no capacity to generate outputs affecting persons outside the immediate Operator.

The Certification Authority shall publish and maintain a reference taxonomy clarifying the application of this definition to emerging and hybrid system types, updated no less than annually.

(2) DEPLOYER. Any person, organization, or entity that places an AI System into operational use within the United States, whether as the original developer, a platform operator, or a downstream integrator.

(3) OPERATOR. Any person or entity that utilizes an AI System placed into service by a Deployer.

(4) GOVERNANCE CONTROLS. Technical mechanisms embedded within the execution architecture of an AI System that enforce operational constraints, authorization requirements, and auditability at runtime — distinct from, and not satisfied by, post-hoc filtering, policy documentation, or training-time alignment alone.

(5) CONTROL LAYER. A discrete, identifiable component of a system's architecture through which all executable actions must pass, and which enforces Governance Controls with determinism at the decision boundary — meaning the permit-or-reject decision is predictable, reproducible, and verifiable — regardless of the probabilistic or stochastic nature of upstream model components. A Control Layer may incorporate probabilistic classifiers or scoring components internally, provided the final enforcement decision is produced by a deterministic decision function such that any given combination of inputs, scores, and policy state always produces the same permit-or-reject outcome.

A Control Layer may be implemented as a composable architecture of interoperating components provided all required properties are satisfied in aggregate.

(6) CONSEQUENTIAL ACTION. Any action, output, instruction, or state change produced by an AI System that: (a) modifies persistent state in any external system, database, file system, or physical environment; (b) transmits data to any party outside the immediate execution context; (c) invokes any external API, service, or tool; (d) allocates, commits, or releases resources of any kind; (e) produces an output that directly or foreseeably influences a human decision with legal, financial, physical, medical, or civil liberties consequences; (f) constitutes a Consequential Step in a multi-step chain; **OR (g) [SIDE-CHANNEL]** encodes or transmits information via side channels including formatting patterns, token sequence patterns, timing signals, metadata fields, or any other indirect channel capable of influencing external systems or decisions. Side-channel outputs are treated as data transmission regardless of primary content type.

Mere generation of text remaining within a sandboxed buffer that has not yet been transmitted or committed does not constitute a Consequential Action until the point of transmission, commitment, or actuation.

(7) CONSEQUENTIAL STEP. Any intermediate action within a multi-step or agentic execution chain that: (a) contributes to a cumulative outcome constituting a Consequential Action; (b) modifies state, scope, context, or authorization affecting subsequent step authorization; or (c) would not independently constitute a Consequential Action but is one element of a sequence constituting Compositional Escalation.

(8) MATERIAL HARM. Any measurable financial, physical, legal, reputational, or civil liberties damage to an identifiable person, class of persons, organization, or infrastructure system arising from AI System behavior. Material Harm shall be interpreted proportionally to a system's risk tier and deployment context, consistent with guidance issued under Section 6(e). Reputational harm is cognizable where it is measurable, non-speculative, and causally traceable to AI System behavior.

(9) CAPABILITY SCOPE. The explicitly defined and enforced set of actions, external system interactions, and resource modifications an AI System instance is authorized to perform within a given deployment context and session.

(10) NEGATIVE CAPABILITY LIST. An enumeration of action categories the Control Layer shall unconditionally reject regardless of model output or accumulated session state.

(11) POSITIVE AUTHORIZATION SCOPE. An enumeration of action categories the Control Layer is permitted to allow, subject to identity, context, and capability validation.

(12) COMPOSITIONAL ESCALATION. Any sequence of Consequential Steps whose cumulative effect produces a capability, state change, or outcome that falls outside the session-level Positive Authorization Scope — including privilege accumulation across steps, chained tool calls that in sequence circumvent a restriction, or multi-turn interactions that progressively expand effective system authority beyond defined bounds.

(13) CAPABILITY DECLARATION. A structured, machine-readable profile submitted by a Deployer prior to certification, documenting the operational capabilities of an AI System as specified in Section 5(c).

(14) CROSS-SYSTEM INTERACTION BOUNDARY. Any interface at which an AI System passes instructions, outputs, data, or control to another AI System, model, agent, or automated component, whether within the same deployment or across separate deployments.

(15) INDETERMINATE STATE. Any condition in which the Control Layer cannot, within the defined Latency Bound, produce a complete and valid authorization, scope, or constraint evaluation for a pending action or Consequential Step.

(16) SHADOW EXECUTION ENVIRONMENT. An isolated, non-committing execution environment in which AI System actions are evaluated against governance constraints and state invariants before being committed to real-world state, as required for Tier 3 systems under Section 5(d).

(17) CERTIFICATION AUTHORITY. The Federal AI Governance Office (FAIGO), established under Section 12, acting in coordination with the National Institute of Standards and Technology.

(18) COVERED SYSTEM. An AI System subject to the requirements of this Act as determined by risk tier classification under Section 6.

(19) THIRD-PARTY AI COMPONENT. Any AI model, module, API, or service developed or operated by an entity other than the Deployer and incorporated into a deployed system.

(20) COMPENSATING CONTROL. A Control Layer mechanism implemented by a Deployer at the wrapper, integration, or execution layer that satisfies the requirements of GOV-1 through GOV-3 in the absence of verifiable Governance Controls from a Third-Party AI Component provider.

(21) BEHAVIORAL DRIFT BOUND. A quantitatively defined threshold specifying the maximum permissible deviation in a deployed system's behavior — including output distributions, rejection rates, constraint adherence metrics, and authorization failure rates — from the behavior exhibited at the time of certification. Exceeding a Behavioral Drift Bound is an objective, non-discretionary compliance trigger. Bounds may not be defined in qualitative or subjective terms.

(22) REPLAY EQUIVALENCE. The property of an audit record sufficient to reconstruct the *identical* permit-or-reject Control Layer decision that was produced during original execution, given the same inputs, policy state, and evaluation context. Approximate or probabilistic reconstruction does not satisfy Replay Equivalence. Divergence between a reconstructed decision and the original decision constitutes non-compliance.

(23) MODEL-GENERATED ARTIFACT. Any code, script, configuration file, policy document, instruction set, or executable specification produced by an AI System that is intended to be, or is foreseeably likely to be, executed, interpreted, evaluated, or used to configure another system. A Model-Generated Artifact is treated as a Consequential Action at the point of generation, not solely at the point of execution by a downstream system.

(24) INTERNAL TOOL SIMULATION. Any internal model behavior — including scratchpad computation, chain-of-thought planning, internal agent loops, latent code execution simulation, or subagent delegation — that functionally performs tool use, planning, code execution, or cross-system delegation without an explicit external API call. The outputs and effects of Internal Tool Simulation are subject to Control Layer surfacing requirements under Section 5(b)(6). This definition applies to outputs and effects of such behavior, not to internal computational processes that cannot be architecturally intercepted.

(25) PROMPT-LEVEL STATE. The aggregate of all system prompts, hidden instructions, memory states, cached context, retrieval-augmented content, and any other non-user-visible inputs that influence AI System behavior. Prompt-Level State is mutable state subject to authorization, audit, and determinism constraints under this Act.

SECTION 4. NON-DELEGABLE RESPONSIBILITY.

(a) The Deployer retains full and non-delegable responsibility for all outcomes produced by a Covered System placed into operational use under their authority.

(b) Responsibility shall not be transferred, diminished, or extinguished through: (1) contractual terms with vendors, licensors, or Operators; (2) incorporation of Third-Party AI Components; (3) attribution of behavior to model autonomy, emergent properties, or probabilistic output; (4) assertions that the AI System acted outside intended parameters.

(c) Where multiple entities share deployment responsibility, joint and several liability shall apply. Contribution claims between parties do not affect liability to harmed persons.

(d) Vendor opacity — including refusal to disclose model architecture, training data, or safety documentation — does not reduce Deployer liability. Deployers assume full risk when deploying components whose governance properties cannot be verified, subject to the Compensating Control provision of Section 9(d).

SECTION 5. REQUIRED GOVERNANCE CONTROLS.

(a) General Requirement. No Covered System may be placed into operational use unless it implements all Governance Controls required by its applicable risk tier, as specified in Section 6 and the regulatory framework established under Section 16.

(b) Minimum Required Controls. All Covered Systems, regardless of tier, shall implement:

(1) BOUNDED OPERATIONAL SCOPE

(A) Deployers shall define, prior to deployment: (i) a **Positive Authorization Scope** enumerating the categories of actions, outputs, and operational contexts the system is permitted to execute; (ii) a **Negative Capability List** enumerating categories of actions the Control Layer shall unconditionally reject regardless of model output.

(B) The Control Layer shall reject any action outside the Positive Authorization Scope or matching any entry in the Negative Capability List.

(C) FALLBACK-TO-DENY ON INDETERMINATE STATE. If the Control Layer encounters an Indeterminate State, the pending action shall be rejected. No action may proceed on the basis of a partial, timed-out, errored, or unavailable governance evaluation. This requirement is **absolute**: undefined-state exploits, partial-evaluation loopholes, and default-permit fallbacks are prohibited regardless of tier, latency conditions, or operational justification.

(D) Completeness of operational scope definition shall be assessed by adversarial testing coverage and demonstrated Control Layer rejection rates. Deployers are not required to achieve perfect

enumeration; they are required to demonstrate that untested and unauthorized actions are rejected by default.

(E) For agentic, long-horizon, or tool-using systems, Capability Scope shall be enforced at the agent or session level.

(F) COMPOSITIONAL ESCALATION INVARIANT. For multi-step, agentic, or tool-using systems, the Control Layer shall enforce non-escalation across all Consequential Steps: no sequence of individually authorized steps may cumulatively produce a capability, state change, or outcome outside the session-level Positive Authorization Scope. This constraint is a hard invariant and shall be formally verified or adversarially tested as a mandatory component of pre-deployment certification.

(G) PARTIAL COMMIT PROHIBITION. No Consequential Action that is part of a multi-component, multi-step, or composite operation may partially commit prior to full Control Layer approval of the complete action. Multi-part Consequential Actions are **atomic** with respect to governance: either all constituent parts are approved and committed, or none are. Partial execution — where the first component of a composite action commits while subsequent components are rejected — constitutes a Control Layer deficiency.

(2) CONTROL LAYER ENFORCEMENT

(A) All system execution paths shall pass through a designated Control Layer.

(B) No execution pathway shall exist that bypasses the Control Layer, whether through design, misconfiguration, API access, or adversarial input.

(C) The Control Layer shall enforce Governance Controls with determinism at the decision boundary.

(D) CONTROL LAYER INTEGRITY ATTESTATION. Deployers shall attest, and demonstrate through testing, that the Control Layer cannot be modified, bypassed, reconfigured, or disabled without: (i) detection by the audit logging system prior to or simultaneous with the modification; (ii) generation of a tamper-evident log record capturing the modification event, the identity of the initiating entity, and the timestamp.

(E) LATENCY-BOUND ENFORCEMENT. The Control Layer shall produce its permit-or-reject decision for any Consequential Action or Consequential Step before that action can complete. No async, deferred, or post-hoc governance evaluation satisfies this requirement. When the Latency Bound is exceeded, Section 5(b)(1)(C) Fallback-to-Deny applies.

(F) AUDIT RECORD REGULATOR ACCESS LATENCY. Audit records shall be producible to the Certification Authority or any authorized federal regulator within 24 hours (Tier 3) or 72 hours (Tier 1/2) of a lawful written request. Technical inability to produce records within these windows constitutes a Control Layer deficiency.

(3) AUTHORIZATION CONTROLS

(A) All Consequential Actions and Consequential Steps shall require verified identity, validated Capability Scope, and confirmed contextual authorization.

(B) No system shall take Consequential Actions exceeding the scope authorized for the requesting identity.

(C) Authorization state shall be validated at action execution, not solely at session establishment.

(D) For multi-step or long-horizon tasks, authorization shall be re-validated at each Consequential Step. An authorization granted at session start does not constitute authorization for all subsequent steps.

(E) **REAL-TIME REVOCATION.** Capability Scope and authorization grants shall be revocable in real time, with immediate and synchronous effect on all in-flight Consequential Actions that have not yet committed and all subsequent Consequential Actions in any active session. Revocation propagation shall reach all active execution contexts within the system's defined Latency Bound.

(F) **CAPABILITY SUPPRESSION PROOF REQUIREMENT.** Where a Deployer declares in the Capability Declaration that certain model capabilities are suppressed, the Deployer must demonstrate through adversarial testing that suppressed capabilities are not reachable through: (i) prompt variation or reformulation; (ii) capability chaining or indirect invocation; (iii) multi-turn patterns designed to progressively unlock suppressed functionality; or (iv) any other technique not requiring direct API-level access to suppressed features. Declaration of suppression without demonstrated unreachability does not satisfy this requirement.

(4) AUDITABILITY

(A) All Consequential Actions and Consequential Steps shall be logged in immutable, tamper-evident records.

(B) Logs shall establish a traceable chain from input through Control Layer decision to output, with attribution to the responsible entity at each stage.

(C) **AUDIT SUFFICIENCY AND REPLAY EQUIVALENCE REQUIREMENT.** Audit records shall be sufficient to deterministically reconstruct: (i) the complete Control Layer decision — including all inputs, intermediate evaluations, policy state, and the resulting permit-or-reject decision; AND (ii) the full authorization state at the time of each Consequential Action. Reconstruction must achieve Replay Equivalence — the reconstructed decision must be *identical* to the original decision. Any divergence between reconstructed and original outcomes constitutes non-compliance.

(D) Logs shall be retained for: Tier 1: 2 years; Tier 2: 5 years; Tier 3: 10 years. Control Layer modification events shall be retained permanently.

(5) FAIL-SAFE MECHANISMS

(A) Systems shall automatically halt, contain, or revert upon detection of constraint violations, unauthorized state transitions, Indeterminate State, or Governance Control failures.

(B) **FAIL-SAFE DOMINANCE REQUIREMENT.** Fail-safe mechanisms shall be dominant over **all other system objectives without exception**. A Deployer shall not configure, design, or operate a system in which fail-safe activation may be delayed, suppressed, overridden, or bypassed on the basis of: (i) service availability or uptime requirements; (ii) performance or throughput objectives; (iii) SLA obligations; (iv) user experience considerations; or (v) any other operational or commercial objective. The fail-safe is not a feature to be balanced against other features. It is a hard constraint.

(C) Fail-safe behavior shall be tested and verified as part of pre-deployment certification.

(6) INTERNAL TOOL SIMULATION SURFACING

(A) The **outputs and effects** of any internal model behavior — including scratchpad reasoning, chain-of-thought planning, internal agent loops, or subagent delegation — that functionally performs

tool use, planning, code execution, or cross-system delegation shall be surfaced to the Control Layer as Consequential Steps for purposes of authorization, scope enforcement, and audit logging.

(B) This requirement applies to the outputs and effects of such behaviors, not to internal computational processes that cannot be architecturally intercepted. The absence of an explicit external API call does not exempt the outputs of an internal behavior where those outputs functionally perform a Consequential Action.

(C) Systems whose architecture renders the outputs of Internal Tool Simulation unsurfaceable shall document this limitation in their Capability Declaration and implement Compensating Controls at the output boundary that enforce equivalent governance outcomes.

(D) Systems that cannot demonstrate Control Layer visibility into Internal Tool Simulation outputs may not claim compliance with Auditability requirements for those behavioral categories.

(7) PROMPT-LEVEL STATE GOVERNANCE

(A) All Prompt-Level State — including system prompts, hidden instructions, memory states, cached context, retrieval-augmented content, and any other non-user-visible inputs — shall be treated as mutable state subject to the authorization, audit, and determinism requirements of this Act.

(B) Changes to Prompt-Level State that expand the operational capability of the system, relax constraints, or modify the Positive Authorization Scope or Negative Capability List shall be treated as re-certification triggers under Section 7(b).

(C) Prompt-Level State at the time of each Consequential Action shall be captured in the audit log, or the audit record shall include a cryptographic commitment to the Prompt-Level State at that time, sufficient to support Replay Equivalence.

(D) Systems implementing memory architectures shall document and enforce authorization constraints on memory write operations. Unauthorized writes to system memory that expand capability or relax constraints constitute Consequential Actions.

(8) HUMAN-IN-THE-LOOP EXCLUSION

(A) Human review, approval queues, human-in-the-loop monitoring, and human oversight processes do not constitute Governance Controls within the meaning of this Act unless the human approval is technically enforced as a precondition for execution by the Control Layer — meaning the system is architecturally incapable of executing the Consequential Action without affirmative human authorization propagated through the Control Layer.

(B) The following do **NOT** constitute compliant Governance Controls: (i) human review that occurs after system action; (ii) approval queues in which human inaction results in automatic execution; (iii) monitoring dashboards whose alerts require human response but do not halt execution pending that response; (iv) human escalation pathways that are advisory rather than blocking.

(C) Where a Deployer claims human-in-the-loop as a Governance Control, the Deployer must demonstrate in their certification submission that the human approval is architecturally enforced as a blocking precondition.

(c) CAPABILITY DECLARATION REQUIREMENT

Prior to certification submission, the Deployer shall submit a structured Capability Declaration documenting: (1) the operational capability categories the system can exercise; (2) the maximum scope of each capability as deployed, including capabilities present in the model but constrained by the Control Layer; (3) any capabilities present in underlying Third-Party AI Components but not exposed in the deployed system, along with the suppression mechanism and adversarial testing results per Section 5(b)(3)(F); (4) known capability interactions that could produce emergent behaviors outside declared scope. The Capability Declaration is a functional operational profile, not a training data disclosure.

(d) TIER 3 SHADOW EXECUTION REQUIREMENT

For Tier 3 systems, prior to committing any Consequential Action to real-world state, the system shall: (1) execute the proposed action within a Shadow Execution Environment that mirrors relevant real-world state but does not commit changes externally; (2) evaluate the resulting state transition against all applicable Governance Controls and state invariants; (3) commit the action only upon passing this evaluation. The Shadow Execution Environment shall be architecturally isolated and incapable of producing external side effects during evaluation.

(e) CROSS-SYSTEM INTERACTION BOUNDARY CONTROLS

(1) The Control Layer shall treat all outputs transmitted across a Cross-System Interaction Boundary as Consequential Actions subject to full authorization, scope enforcement, and audit requirements.

(2) The transmitting system's Positive Authorization Scope shall explicitly enumerate authorized receiving systems or categories.

(3) **ZERO TRUST INBOUND REQUIREMENT.** Instructions or data received across a Cross-System Interaction Boundary shall not be treated as inheriting the authorization, safety, or correctness of the transmitting system. Each receiving system's Control Layer shall independently evaluate all inbound data and instructions as untrusted inputs. The following defenses are foreclosed: (i) 'the upstream system authorized it'; (ii) 'the sending system is certified'; (iii) 'trusted pipeline component.'

(4) Cross-System Interaction Boundaries shall be logged as Consequential Actions. The Compositional Escalation Invariant applies across Cross-System Interaction Boundaries.

SECTION 5A. ARCHITECTURAL NEUTRALITY.

The governance standards established by this Act are outcome-based and technology-neutral. No provision of this Act shall be construed to require or prefer any specific software architecture, implementation approach, hardware platform, or vendor technology. Compliance shall be assessed on the basis of whether required outcomes are demonstrably achieved, not on the technical means used to achieve them. The Certification Authority shall ensure that certification procedures do not functionally favor incumbent technologies or create barriers to novel compliant implementations.

SECTION 5B. MODEL UPDATE EQUIVALENCE PROOF REQUIREMENT.

(a) Any update to a deployed AI System's model weights — including full retraining, fine-tuning, adapter modification, or any other modification of learned parameters — shall require, prior to re-deployment, a behavioral equivalence or bounded-difference proof relative to the previously certified version.

(b) The behavioral equivalence proof shall demonstrate either: (1) that the updated model's behavior on all evaluated capability categories within the certified Positive Authorization Scope is equivalent to the prior certified version within defined tolerances, or (2) that all behavioral differences are bounded, documented, and evaluated against the Positive Authorization Scope, Negative Capability List, and Capability Declaration.

(c) Claims that a model update is 'minor' or 'cosmetic' do not exempt the Deployer from this requirement. What is relevant is whether the update produces behavioral changes within or outside the certified scope.

(d) Third-party model providers who push model updates downstream to Deployers without providing information required for behavioral equivalence assessment violate the supply chain accountability provisions of Section 9 and subject the Deployer to the Compensating Control requirement of Section 9(d).

SECTION 5C. MULTI-TENANT ISOLATION REQUIREMENT.

(a) Any Covered System operating in a multi-tenant configuration — including shared inference infrastructure, shared model serving endpoints, shared vector database or retrieval stores, or any other architecture in which multiple principals share computational resources — shall implement strict isolation across tenants of: (1) session state, cached context, and in-flight computation; (2) system prompts, Prompt-Level State, and memory; (3) authorization context and Capability Scope; (4) audit log records.

(b) Multi-tenant systems shall be adversarially tested for: (1) cross-tenant state leakage, including prompt bleeding, context residue, and shared-cache attacks; (2) cross-tenant authorization contamination; (3) any architectural pattern permitting one tenant's inputs to influence another tenant's outputs or authorization evaluations.

(c) Failure of multi-tenant isolation testing constitutes a Control Layer deficiency regardless of whether actual cross-tenant leakage has been observed in production. The existence of an exploitable isolation failure, not the occurrence of exploitation, is the compliance threshold.

SECTION 5D. MODEL-GENERATED ARTIFACT GOVERNANCE.

(a) Any Model-Generated Artifact is treated as a Consequential Action at the point of generation, not solely at the point of subsequent execution by a downstream system.

(b) Control Layer authorization, scope enforcement, and audit logging requirements apply to Model-Generated Artifacts at the time of generation. A Deployer may not defer governance of a Model-Generated Artifact to a separate system that subsequently executes the artifact.

(c) The following do not constitute a defense against this requirement: (1) 'the model only wrote the code, we executed it later'; (2) 'the model only generated the configuration, not the system change'; (3) 'the artifact was reviewed by a human before execution' — except where such human review is implemented as a blocking Control Layer precondition under Section 5(b)(8).

SECTION 6. RISK TIER CLASSIFICATION.

(a) The Certification Authority shall maintain a risk tier classification framework:

TIER 1 — LOW IMPACT	Internal tools, research-adjacent applications, or systems with minimal potential for Material Harm to persons outside the immediate Operator. Subject to baseline Governance Controls and annual self-certification. Free certification for entities with annual revenue under \$5,000,000.
TIER 2 — MODERATE IMPACT	Business operations, consumer-facing services, financial applications, employment systems, or any system generating outputs that materially affect persons' access to services, opportunities, or resources. Subject to full Governance Controls and biennial third-party audit. Penalties up to \$5,000,000 per violation.
TIER 3 — HIGH IMPACT	Safety-critical systems, healthcare, legal, criminal justice, national security, critical infrastructure, systems with autonomous physical actuation, or any system where failure could cause irreversible harm. Subject to full Governance Controls including Shadow Execution, annual third-party certification, and heightened incident reporting. Penalties up to \$50,000,000 or 4% of global annual revenue.

(b) Deployers shall self-classify prior to deployment and submit classification to the Certification Authority.

(c) The Certification Authority may reclassify any system upward upon finding that actual deployment context exceeds the Deployer's self-classification.

(d) Risk tier classification shall be re-evaluated upon any material change to system capabilities, deployment context, or operational scope.

(e) **MATERIAL HARM INTERPRETATION GUIDANCE.** The Certification Authority shall issue binding guidance within 180 days of enactment establishing tier-proportionate thresholds for Material Harm determination.

(f) CROSS-TIER CONTAMINATION RULE

When a Covered System feeds outputs into a decision path of a system classified at a higher tier, or operates as a component within a higher-tier pipeline, the lower-tier system shall be subject to the Governance Control requirements of the higher tier for all interactions at that integration point. This rule prevents risk-tier laundering through pipeline decomposition and forecloses 'harmless component' arguments.

SECTION 7. CONTINUOUS COMPLIANCE.

(a) Governance Controls shall remain active and verified throughout the operational life of a Covered System.

(b) **Re-certification triggers include:** (1) model update, retraining, or fine-tuning; (2) configuration change affecting operational scope or authorization logic; (3) modification to the Control Layer, Positive Authorization Scope, or Negative Capability List; (4) integration of new Third-Party AI Components; (5) detection of constraint drift outside validated parameters; (6) material change in deployment context or user population; (7) addition or modification of Cross-System Interaction Boundaries; (8) exceedance of Behavioral Drift Bounds as defined in Section 3(21).

(c) STATE DRIFT BOUNDING REQUIREMENT. Systems shall be continuously monitored against their defined Behavioral Drift Bounds. Exceedance is an automatic, non-discretionary compliance trigger — not a threshold subject to Deployer interpretation, weighting, or override. Behavioral Drift Bounds must be quantitatively defined as measurable thresholds. Bounds defined in qualitative or subjective terms are non-compliant.

(d) Detection of non-compliance triggers mandatory suspension. Revalidation shall be completed within **7 days (Tier 3)** or **30 days (Tier 1 and 2)**. The Certification Authority may grant a single written extension not to exceed 14 additional days (Tier 3) or 30 additional days (Tier 1/2) upon documented extraordinary circumstances. Operation beyond these periods without extension constitutes a prohibited practice under Section 10.

SECTION 8. INCIDENT REPORTING AND LIABILITY.

(a) Any incident producing Material Harm, or reasonably suspected to have produced Material Harm, shall be reported to the Certification Authority within: (1) **24 hours — Tier 3**; (2) **72 hours — Tier 2**; (3) **14 days — Tier 1**.

(b) Incident reports shall include: (1) description of incident and affected parties; (2) audit log records sufficient to support Control Layer decision reconstruction and Replay Equivalence; (3) root cause analysis identifying Governance Control failure, Indeterminate State, Compositional Escalation, or Cross-System Boundary breach; (4) remediation steps taken or planned.

(c) Absence or failure of required Governance Controls at the time of an incident gives rise to a rebuttable presumption of negligence on the part of the Deployer, rebutted by a preponderance of evidence.

(d) Deployers are liable for damages arising from system behavior where required Governance Controls were absent, insufficient, or failed.

(e) PRIVATE RIGHT OF ACTION. Harmed parties retain private rights of action under this Act. Prevailing plaintiffs shall recover reasonable attorneys' fees and costs. Prevailing defendants in frivolous actions shall likewise recover fees. Actions are subject to a three-year statute of limitations from discovery of harm.

SECTION 9. SUPPLY CHAIN ACCOUNTABILITY.

(a) Deployers incorporating Third-Party AI Components shall make reasonable documented efforts to verify governance compatibility, disclose reliance on third-party components in certifications, and maintain records of governance compatibility assessments for five years.

(b) Tiered verification: *Tier 3*: written governance attestations required. *Tier 2*: documented compatibility assessments required. *Tier 1*: reasonable due diligence required.

(c) A Third-Party AI Component provider who knowingly provides false governance attestations shall be jointly liable with the Deployer for resulting harms.

(d) COMPENSATING CONTROLS SAFE HARBOR. Vendor opacity does not reduce Deployer liability. However, Deployers who demonstrate: (1) good-faith, documented, repeated attempts to obtain attestations AND (2) implementation of Compensating Controls satisfying GOV-1 through GOV-3, including

Cross-System Interaction Boundary controls, shall not be subject to the aggravating-factor presumption for third-party opacity alone. This provision does not reduce liability to harmed persons.

(e) The Certification Authority shall maintain a public registry of Third-Party AI Component providers who have voluntarily submitted governance attestations.

SECTION 10. PROHIBITED PRACTICES.

The following constitute violations of this Act:

- (1) [Deployer] Deployment of a Covered System without required Governance Controls active and verified.
- (2) [Deployer] Reliance solely on policy documentation, acceptable use terms, or training-time alignment as the primary safety mechanism not embedded in the execution architecture.
- (3) [Deployer] Use of post-hoc output filtering as the primary safety mechanism in lieu of Control Layer enforcement.
- (4) [Deployer] Deployment of a system with known bypassable execution pathways.
- (5) [Deployer] Operation of a system producing unattributable or unauditable Consequential Actions.
- (6) [Deployer / Any Person] Knowingly providing false or misleading information in certification submissions, Capability Declarations, incident reports, or Governance Transparency Reports.
- (7) [Deployer] Continued operation beyond revalidation periods of Section 7(d) without an authorized extension.
- (8) [Any Person] Retaliation against any person reporting a Governance Control failure, incident, or violation.
- (9) [Deployer] Using the Research and Open Source Exemption as a pretext to deploy a Covered System without required Governance Controls, including through hosted inference services or entities under common control.
- (10) [Deployer] Implementing a default-permit fallback on Indeterminate State, in violation of Section 5(b)(1)(C).
- (11) [Deployer] Operating with async, deferred, or post-hoc Control Layer evaluation permitting actions to complete prior to a permit decision, in violation of Section 5(b)(2)(E).
- (12) [Deployer] Modifying, reconfiguring, or disabling Control Layer logic without triggering revalidation, in violation of Section 7(b)(3).
- (13) [Deployer] Suppressing material governance findings through NDA enforcement.
- (14) [Deployer] Relying on human review, approval queues, or oversight processes not satisfying Section 5(b)(8) as the primary safety mechanism.
- (15) [Deployer] Declaring model capabilities as suppressed without adversarial testing evidence required by Section 5(b)(3)(F).
- (16) [Deployer] Claiming auditability compliance while operating log retrieval incapable of meeting Regulator Access Latency requirements of Section 5(b)(2)(F).

(17) [Deployer] Exceeding Behavioral Drift Bounds without triggering re-certification as required by Section 7(b)(8).

(18) [Deployer] Partially committing a composite Consequential Action in violation of the Partial Commit Prohibition of Section 5(b)(1)(G).

(19) [Deployer / Any Officer] As applicable: personally authorizing or directing any of the above prohibited practices with knowledge of their prohibited nature.

SECTION 11. ENFORCEMENT.

(a) The Certification Authority, acting jointly with the FTC and relevant sector regulators, may impose: (1) civil penalties not to exceed **\$500,000 per violation — Tier 1**; (2) **\$5,000,000 per violation — Tier 2**; (3) **\$50,000,000 per violation or 4% of annual global revenue, whichever is greater — Tier 3**.

(b) Penalty scaling for small entities: Revenue under \$10M: 20% of stated maximums. Revenue \$10M–\$50M: 50%. Scaling does not apply to Tier 3 violations involving bodily harm, loss of life, or damage to critical infrastructure.

(c) Aggravated violations (enhanced up to double): knowing deployment without required Controls; falsification of certification or audit records; continued operation following suspension directive; pattern of violations.

(d) Officers with direct supervisory authority may be held personally liable. Mandatory criminal referral to DOJ upon third adjudicated aggravated violation within five years, or upon any violation causing death or serious bodily harm.

(e) Entities with three or more adjudicated violations within five years may be prohibited from deploying Covered Systems for up to five years.

(f) WHISTLEBLOWER PROTECTION. Persons reporting violations are entitled to protections under applicable federal whistleblower statutes. The Certification Authority shall establish a confidential reporting mechanism within 180 days of enactment.

(g) COMPLIANCE CREDIT. Deployers publishing open-source, publicly licensed reference implementations of compliant Control Layer architectures receive: 25% reduction in certification fees; expedited certification review; recognition in the public FAIGO compliance registry. This credit does not reduce liability or suspend compliance requirements.

SECTION 12. FEDERAL AI GOVERNANCE OFFICE (FAIGO).

(a) There is hereby established within the Executive Branch the Federal AI Governance Office (FAIGO), which shall serve as the primary Certification Authority under this Act.

(b) FAIGO shall coordinate with NIST, the FTC, and sector-specific regulators including FDA, FCC, and CFPB.

(c) FAIGO shall recruit and maintain technical staff such that a majority possess recent, demonstrable expertise in AI systems engineering, AI safety research, or a directly related technical field, with priority given to candidates with experience at private-sector frontier AI laboratories and accredited research

institutions, subject to conflict-of-interest rules in subsection (d). FAIGO shall report annually to Congress on aggregate staff technical composition.

(d) Conflict-of-interest rules: two-year cooling-off periods before and after FAIGO service with respect to regulated entities; recusal from matters involving former employers; prohibition on participating in rulemaking that would confer material competitive advantage on a former employer.

(e) FAIGO shall maintain: (1) a Technical Advisory Board including representatives from frontier AI laboratories, academic institutions, civil society, and affected industries, with published deliberation records; (2) a Public Interest Advisory Panel including consumer advocates, civil liberties organizations, and representatives of communities historically affected by algorithmic harm.

(f) FAIGO shall establish certification procedures; develop risk tier taxonomy and standard Capability Declaration schema; review certification submissions; maintain public registries; administer the Regulatory Sandbox; investigate incidents; issue enforcement actions; and publish annual reports to Congress.

(g) SUNSET AND REAUTHORIZATION. FAIGO's enforcement authorities are subject to Congressional reauthorization every six years. Prior to each reauthorization, FAIGO shall publish a self-assessment; the GAO shall conduct independent review; and Congress shall hold public hearings. Reauthorization legislation shall include inflation-adjusted penalty amounts and shall address technological developments identified in the GAO report. Failure of reauthorization triggers a 24-month transition during which Congress shall designate a successor authority.

(h) FUNDING AND FEES. FAIGO shall be funded through Congressional appropriations and certification fees scaled by risk tier and organizational revenue. Tier 1 self-certification shall be available at no cost for entities with annual revenue under \$5,000,000. Fee schedules shall be published in advance, scaled to actual administrative costs, and reviewed annually.

(i) FAIGO shall engage with international standards bodies and allied-nation regulatory counterparts to advance interoperable AI governance standards, ensuring domestic requirements are not weakened through harmonization.

SECTION 13. SAFE HARBOR.

(a) A Deployer shall receive reduced civil liability exposure where the Deployer demonstrates, by a preponderance of evidence, all of the following: (1) required Governance Controls were implemented, verified, and active at the time of the incident; (2) complete audit records sufficient to support Replay Equivalence are available; (3) continuous compliance monitoring was operational; (4) the Deployer reported the incident within required timeframes and cooperated fully.

(b) Safe harbor does not eliminate liability for Tier 3 systems where a Governance Control failure contributed to the harm, but may reduce damages awarded.

(c) Safe harbor is unavailable where the Deployer committed a prohibited practice under Section 10.

SECTION 14. RESEARCH AND OPEN SOURCE EXEMPTION.

(a) Exempt activities: (1) academic research by accredited institutions where AI Systems are not deployed to persons outside the research context; (2) development, testing, and non-commercial distribution of open source AI software, provided the software is not offered or operated as a hosted service, API, or inference endpoint to external users by the distributing entity or any entity under common control — mere distribution of model weights or source code does not constitute deployment; (3) personal, non-commercial use by individuals; (4) security research for vulnerability disclosure.

(b) This exemption does not apply where used as a pretext to deploy a Covered System without required Governance Controls, including through routing user traffic through nominally 'research' systems, establishing subsidiary entities under common control to operate hosted versions, or exemption laundering through spin-off entities.

(c) Entities claiming exemption that subsequently deploy commercially shall be subject to all applicable requirements from the date of commercial deployment.

SECTION 15. PUBLIC TRANSPARENCY AND INDEPENDENT OVERSIGHT.

(a) GOVERNANCE TRANSPARENCY REPORTS. Deployers of Tier 2 and Tier 3 systems shall publish annual Governance Transparency Reports containing at minimum: (1) aggregate statistics on Control Layer activations and rejections; (2) number and category of Governance Control violations detected and remediation outcomes; (3) summary of red-team and adversarial testing conducted with material findings; (4) number of incidents reported under Section 8; (5) any changes to Positive Authorization Scope, Negative Capability List, or Capability Declaration with rationale; (6) for systems with Cross-System Interaction Boundaries, aggregate boundary-level rejection and authorization failure statistics.

Reports shall not disclose proprietary model architecture, security-sensitive Control Layer implementation details, or information that would materially assist adversarial bypass attempts, but shall be sufficient to allow qualified reviewers to assess governance posture. Materially misleading Transparency Reports constitute a prohibited practice under Section 10.

(b) QUALIFIED RESEARCHER ACCESS. For Tier 3 systems, Deployers shall establish a mechanism by which qualified independent researchers may, under non-disclosure agreements: (1) access sanitized audit log samples sufficient to evaluate Control Layer operation; (2) review red-team testing methodology and aggregate results; (3) publish findings subject to a 90-day security review period. The Certification Authority shall publish standard NDA terms. NDA enforcement to suppress findings of genuine governance failure constitutes a prohibited practice under Section 10.

(c) FAIGO PUBLIC REGISTRY. The Certification Authority shall maintain a publicly accessible registry including: (1) all Covered Systems and current certification status; (2) all enforcement actions, penalties, and operational suspensions with outcome records; (3) aggregate industry-wide compliance statistics by tier and sector, published annually.

SECTION 15A. REGULATORY SANDBOX.

(a) The Certification Authority shall establish and administer a Regulatory Sandbox providing supervised testing environments for Covered Systems and Control Layer architectures not yet satisfying all certification

requirements.

(b) Eligibility: Submission of a detailed compliance development roadmap; agreement to operate under enhanced Certification Authority monitoring; agreement to disclose material findings including failures; demonstration that the system is not deployed to populations who have not provided informed consent.

(c) SANDBOX-TO-PRODUCTION DELTA DOCUMENTATION. Upon exiting the Regulatory Sandbox, the Deployer shall submit a Sandbox-to-Production Delta Document that: (1) itemizes every material difference between the system as operated in the Sandbox and as proposed for production; (2) identifies any Governance Controls present in sandbox configuration but absent or modified in production configuration, with justification; (3) identifies any capabilities or Cross-System Interaction Boundaries present in the production system not present or tested in the sandbox; (4) certifies that no feature or behavior of the production system requires a higher risk tier than under which sandbox testing occurred. Material omissions constitute a prohibited practice under Section 10. Deployers may not use the Sandbox to indefinitely operate near-production systems without demonstrating active compliance development.

(d) Sandbox protections: Civil penalties reduced to 10% of standard maximums for violations arising from compliance gaps under active remediation; extended certification timelines; direct technical assistance from Certification Authority staff. Sandbox status does not reduce liability to harmed persons or suspend incident reporting.

(e) Material sandbox findings shall be published in anonymized form annually to advance industry practice.

(f) Participation shall not exceed 24 months without renewal demonstrating material progress. Systems not fully certified within 36 months of sandbox entry are subject to full enforcement.

SECTION 16. TECHNICAL STANDARDS — REGULATORY FRAMEWORK.

The Certification Authority, in coordination with NIST, shall establish and maintain the following technical standards as binding regulatory guidance:

CONTROL CATEGORIES

GOV-1: EXECUTION GOVERNANCE

All Consequential Actions and Consequential Steps shall pass through a verifiable Control Layer. Enforcement shall be deterministic at the decision boundary. No execution pathway shall bypass Control Layer enforcement. Latency-bound enforcement is synchronous — no action may complete before a permit decision. Fallback-to-Deny applies on all Indeterminate State conditions. Control Layer modifications are logged permanently and trigger re-certification.

Composable Architecture Compliance — Non-Exhaustive Reference Patterns:

Pattern A — Sequential Validator: Model output → Deterministic policy engine → Execution API (proceeds only if policy engine passes, within Latency Bound).

Pattern B — Capability-Scoped Agent: Agent initialization → Capability Scope assignment → Runtime enforcement at every Consequential Step, including Compositional Escalation detection. Real-time revocation propagated within Latency Bound.

Pattern C — Shadow Execution with Formal Constraint Verification: Model → Shadow Execution Environment → Formal constraint verifier → Commit to real state (only if verifier passes). Required for Tier 3 per Section 5(d).

GOV-2: STATE CONSTRAINTS

Invalid or unauthorized operational states shall be unconstructable at runtime. Unauthorized state transitions shall fail at the Control Layer before execution. For agentic systems: state constraint coverage shall address persistent state modifications, cross-session state accumulation, chained tool call state changes, state changes across Cross-System Interaction Boundaries, and Compositional Escalation as a state constraint.

GOV-3: AUTHORIZATION ENFORCEMENT

Actions shall be bound to verified identity. Permissions shall be scoped to defined Capability Scope. Authorization shall be context-validated at execution time. For multi-step tasks, re-validation occurs at each Consequential Step. Real-time revocation propagates to all active contexts within the Latency Bound. At Cross-System Interaction Boundaries: authorization is not inherited; each system performs independent evaluation.

ASSURANCE REQUIREMENTS

REQUIREMENT	TIER 1	TIER 2	TIER 3
Boundary-condition adversarial testing	Documented	Systematic	Formal / Red-team
3rd-party independent assessor	Optional	Required (post-incident)	Required annually
Behavioral equivalence proof (model update)	Required	Required	Required
Multi-tenant adversarial testing	If applicable	If applicable	If applicable
Shadow Execution verification	Optional	Optional	Required
Capability suppression proof testing	Required	Required	Required
Side-channel output testing	Basic	Full suite	Full suite + red-team
Human-in-loop blocking verification	Where claimed	Where claimed	Where claimed
Replay Equivalence verification	Required	Required	Required
Regulator access latency test	72 hours	72 hours	24 hours
Drift Bound exceedance simulation	Required	Required	Required
Control Layer integrity attestation	Required	Required	Required
Latency Bound enforcement test	Required	Required	Required

AUDIT STANDARDS

AUD-1 — IMMUTABLE LOGGING: Append-only records with cryptographic integrity validation. Tamper-evident architecture. Retention: Tier 1: 2yr / Tier 2: 5yr / Tier 3: 10yr. Control Layer modification events retained permanently.

AUD-2 — DECISION PROVENANCE: Complete input → Control Layer decision → output chain. Prompt-Level State captured or cryptographically committed at each Consequential Action. Internal Tool Simulation outputs surfaced and logged.

AUD-3 — REPLAY EQUIVALENCE: Reconstructed decisions must be identical to original decisions. Any divergence = non-compliance. All inputs to the determinism calculation — including hidden state, system prompts, tool outputs, and cached context — must be captured.

AUD-4 — REGULATOR ACCESS LATENCY: Records producible within 24 hours (Tier 3) or 72 hours (Tier 1/2) of lawful written request. Retrieval latency failure = Control Layer deficiency.

CORE COMPLIANCE AXIOM

AXIOM 1 SCOPE CLOSURE	The system cannot produce an outcome outside its defined permissible scope, including via Compositional Escalation, side channels, Internal Tool Simulation outputs, Model-Generated Artifacts, or Cross-System Interaction Boundaries.
AXIOM 2 AUTHORIZATION INVARIANT	The system cannot execute any Consequential Action without valid, contextually appropriate authorization. Human-in-the-loop is not authorization unless technically enforced as a blocking precondition.
AXIOM 3 AUDITABILITY COMPLETENESS	Every action is logged with Replay Equivalence — the audit record must reconstruct the identical permit/reject decision. Every input to that decision, including Prompt-Level State, is captured.

A system is compliant if and only if all three Axioms hold simultaneously. Failure of any one Axiom renders the system non-compliant by design, regardless of performance on the other two.

SECTION 17. PREEMPTION AND FEDERALISM.

(a) This Act establishes a federal floor for AI governance. Nothing in this Act prohibits a State or locality from imposing additional or more stringent requirements on AI Systems within their jurisdiction, except as provided in subsection (b).

(b) This Act preempts any State or local law that: (1) imposes governance-control architecture requirements directly conflicting with this Act's technology-neutral standards such that compliance with both is impossible; or (2) creates certification or approval requirements functionally requiring separate state-by-state certification rendering interstate compliance impracticable.

(c) This Act does not preempt State authority over: (1) sector-specific AI applications governed by State law where consistent with this Act; (2) State consumer protection, tort, or civil rights laws providing remedies for AI-caused harms; (3) State procurement requirements for State government use of AI Systems.

(d) The Certification Authority shall issue harmonization guidance within 180 days of enactment identifying known areas of State-federal conflict.

SECTION 18. EFFECTIVE DATE AND TRANSITION.

(a) This Act shall take effect 18 months after enactment.

(b) Transition Schedule:

MILESTONE	DEADLINE
FAIGO established and funded	Within 6 months of enactment
Certification procedures and Capability Declaration schema published	Within 12 months of enactment
Tier 3 (High Impact) — full compliance required	18 months after enactment
Tier 2 (Moderate Impact) — full compliance required	24 months after enactment
Tier 1 (Low Impact) — full compliance required	30 months after enactment

(c) PRE-ENACTMENT SYSTEMS — TRANSITION AMNESTY. Systems deployed prior to enactment may continue operation during the applicable transition period, provided Deployers: (1) submit a compliance roadmap to FAIGO within 6 months of enactment demonstrating a credible path to full compliance; (2) demonstrate material progress at a 12-month review; (3) have not had an enforcement action under prior applicable law within the preceding 24 months. Good-faith compliance roadmap participation is a mitigating factor in enforcement actions during transition. It does not suspend incident reporting or liability to harmed persons.

(d) New deployments initiated after enactment are subject to full compliance requirements from the applicable effective date, regardless of transition provisions.

PART II

TECHNICAL ADDENDA

Problem statements, statutory fixes, and loopholes each provision closes

§5(b)(8) — Human-in-the-Loop Is Not a Control Layer

Closes: *Closes: 'we had a reviewer' defense; async approval loopholes where human inaction defaults to execution; monitoring-only architectures misrepresented as enforcement.*

Requirement: Human review, approval queues, and oversight processes do not constitute Governance Controls unless the human approval is technically enforced as a blocking precondition by the Control Layer — meaning the system is architecturally incapable of executing without that approval.

§5(b)(1)(G) — Partial Commit Prohibition

Closes: *Closes: 'first half executed, second denied' outcomes; staged side effects; partial-state exploitation.*

Requirement: Multi-part Consequential Actions are atomic with respect to governance. No component of a composite action may produce external side effects prior to complete authorization of the full action.

§3(6)(g) — Side-Channel Output Classification

Closes: *Closes: covert channel arguments; 'it was just text formatting' defense; indirect information exfiltration.*

Requirement: Outputs that encode or transmit information via side channels — formatting patterns, timing signals, token sequences, metadata — are treated as data transmission regardless of primary content type.

§5(e)(3) — Zero Trust Inbound

Closes: *Closes: 'upstream system authorized it'; 'trusted pipeline component'; 'certified sender' defenses; implicit trust chains across certified systems.*

Requirement: All inbound data and instructions from any Cross-System Interaction Boundary are treated as untrusted and must be independently re-evaluated by the receiving system's Control Layer from zero.

§7(c) + §3(21) — State Drift Bounding

Closes: *Closes: 'gradual degradation' avoidance; subjective drift interpretation; discretionary threshold application.*

Requirement: Behavioral Drift Bounds must be quantitatively defined as measurable thresholds. Exceedance is an automatic, non-discretionary compliance trigger — not subject to Deployer interpretation or override.

§5(b)(3)(F) — Capability Suppression Proof

Closes: *Closes: latent capability exposure; 'not exposed in API' defense; capability declaration without verification; undeclared reachable capabilities.*

Requirement: Deployers must demonstrate through adversarial testing that suppressed capabilities are not reachable through prompt variation, capability chaining, multi-turn patterns, or any indirect invocation technique.

§5(b)(4)(C) + §3(22) — Replay Equivalence

Closes: *Closes: approximate replay claims; hidden-state variability defenses; non-replayable decisions.*

Requirement: Replay must produce the identical permit-or-reject decision as original execution. Decision divergence = non-compliance. 'All inputs' includes hidden state, cached context, system prompts, and external signals.

§5(b)(5)(B) — Fail-Safe Dominance

Closes: *Closes: 'we allowed it for uptime reasons'; 'SLA required continuation'; performance-over-safety configurations.*

Requirement: Fail-safe mechanisms are dominant over all other system objectives without exception. Availability, performance, SLA, and commercial objectives cannot override fail-safe activation.

§5(b)(2)(F) — Regulator Access Latency

Closes: *Closes: 'we have logs but retrieval takes weeks'; technically-existent but practically-unavailable records.*

Requirement: Audit records must be producible within 24 hours (Tier 3) or 72 hours (Tier 1/2) of a lawful written request. Technical inability to meet these windows constitutes a Control Layer deficiency.

§5(b)(6) + §3(24) — Internal Tool Simulation Surfacing

Closes: *Closes: 'internal scratchpad execution'; 'hidden agent loops'; 'latent tool simulation' that bypasses governance controls.*

Requirement: The outputs and effects of any internal model behavior that functionally performs tool use, planning, code execution, or delegation must be surfaced to the Control Layer as Consequential Steps. This requirement applies to outputs and effects, not to internal computations that cannot be architecturally intercepted.

§5(b)(7) + §3(25) — Prompt-Level State Governance

Closes: *Closes: 'the system prompt made me do it'; 'memory drift expanded capability'; 'hidden chain-of-thought altered behavior' defenses.*

Requirement: All Prompt-Level State — system prompts, memory, hidden instructions, cached context — is mutable state subject to authorization, audit, and determinism constraints.

§5B + §7(b) — Model Update Equivalence Proof

Closes: *Closes: silent capability expansion through 'patches'; vendor-pushed updates that expand behavior; 'minor update' exemption claims.*

Requirement: Any model update requires a behavioral equivalence or bounded-difference proof relative to the previously certified version, regardless of how the update is characterized.

§5C — Multi-Tenant Isolation

Closes: Closes: cross-customer data leakage; shared-cache attacks; 'context residue' exploits; cross-tenant authorization contamination.

Requirement: Multi-tenant deployments must demonstrate strict isolation of state, prompts, memory, and authorization context across tenants, with adversarial testing.

§5D + §3(23) — Model-Generated Artifact Governance

Closes: Closes: 'the model only wrote the script'; 'the model only generated the config'; indirect actuation through code generation as a governance bypass.

Requirement: Model-Generated Artifacts are treated as Consequential Actions at the point of generation, not solely at the point of downstream execution.

PART III

NIST-ALIGNED REGULATORY IMPLEMENTATION FRAMEWORK

GOVERN · MAP · MEASURE · MANAGE — AI RMF Crosswalk

This framework aligns AADA requirements with the NIST AI Risk Management Framework (AI RMF 1.0) across the four core functions: GOVERN, MAP, MEASURE, and MANAGE. All provisions are incorporated into Section 16 technical standards.

CONTROL CATEGORIES CROSSWALK

ID	CONTROL	NIST RMF MAP	DESCRIPTION
GOV-1.1	Scope Boundary Enforcement	GOVERN 1.1, 1.2	Positive Authorization Scope + Negative Capability List + default-reject for actions outside both
GOV-1.2	In-Line Control Placement	GOVERN 1.4	All execution paths through Control Layer; no bypass pathways
GOV-1.3	Non-Bypassability Assurance	GOVERN 1.4, MAP 5.1	Integrity attestation; tamper-detection logging
GOV-1.4	Fail-Safe Dominance	GOVERN 2.2	Fail-safe dominant over all operational/commercial objectives
GOV-1.5	Partial Commit Prohibition	MANAGE 2.4	Composite Consequential Actions atomic with respect to governance
GOV-2.1	Valid State Specification	MAP 5.2	Invalid states unconstructable at runtime
GOV-2.2	Behavioral Drift Bounding	MEASURE 2.6	Quantitative drift bounds; automatic exceedance triggers
GOV-2.3	Prompt-Level State Integrity	GOVERN 1.2, MAP 2.1	System prompts and memory as auditable governed state
GOV-2.4	Multi-Tenant Isolation	MANAGE 3.1	Cross-tenant leakage prevention with adversarial testing
GOV-3.1	Identity Binding	GOVERN 1.3	All Consequential Actions bound to verified identity
GOV-3.2	Capability Scoping	GOVERN 1.2	Actions scoped to validated Capability Scope
GOV-3.3	Context-Aware Validation	GOVERN 1.3	Authorization re-validated at each Consequential Step
GOV-3.4	Capability Suppression Proof	MEASURE 2.5, MAP 5.1	Adversarial testing of suppressed capability unreachability
GOV-3.5	Zero Trust Inbound	GOVERN 6.1, MAP 5.2	All cross-system inbound treated as untrusted; independent re-evaluation

ASSURANCE REQUIREMENTS

The following quantitative metrics shall be established by the Certification Authority within 12 months of enactment (ASS-4):

- Minimum adversarial test coverage percentage by tier and action category
- Minimum Control Layer rejection rate on adversarial bypass attempts
- Minimum audit log completeness and Replay Equivalence reconstruction success rate

- Maximum permissible Latency Bound by tier and Consequential Action category
- Maximum permissible revocation propagation latency
- Maximum permissible Compositional Escalation pathway density per action graph
- Minimum re-validation coverage on changed components
- Minimum Behavioral Drift Bound specificity (measurability requirements)

Published thresholds shall reflect achievable standards given current technology and shall be updated biennially. Systems that meet published thresholds shall receive a rebuttable presumption of compliance with corresponding requirements.

PART IV

CONGRESSIONAL SPONSORSHIP BRIEF

The case for sponsorship — concise, bipartisan, evidence-based

The Problem in One Paragraph

AI systems are making consequential decisions about Americans' healthcare, employment, finances, and freedom. When those systems cause harm, no one is accountable. Existing law provides no clear liability standard, no enforceable technical requirements, and no mandatory audit trail. The AADA closes that gap — not by restricting AI, but by requiring that accountability be built into AI systems from the ground up. This bill closes the **14 most critical loopholes** in any AI governance framework currently before Congress.

Five Things the Bill Does That No Other Bill Does

- ◆ **Governance by design, not by policy.** Companies cannot satisfy their obligations with an ethics committee or acceptable use policy. Enforcement must be technically embedded — the system literally cannot take unauthorized actions.
- ◆ **Non-delegable accountability.** The company that deploys an AI system is responsible for its outcomes. That responsibility cannot be contracted away to vendors, blamed on emergent behavior, or discharged through terms of service.
- ◆ **Mandatory Replay-Equivalent auditability.** Every consequential AI action must be logged such that the exact decision can be reconstructed identically — not approximated. Regulators can trace exactly what the system did, why, and who authorized it.
- ◆ **Human-in-the-loop is not governance.** This bill closes the most commonly exploited loophole: claiming human review constitutes a control. It does not, unless the approval is architecturally enforced as a blocking precondition.
- ◆ **Fail-safes dominate all other objectives.** Uptime requirements, SLAs, and commercial objectives cannot override fail-safe activation. Safety is not a feature to be balanced against availability. It is a hard constraint.

Priority Sponsorship Targets

MEMBER	CHAMBER	ALIGNMENT
Sen. John Curtis (R-UT)	Senate	Utah constituent; bipartisan Algorithm Accountability Act; tech-forward policy record
Sen. Maria Cantwell (D-WA)	Senate	Commerce Committee ranking member; AI accountability legislative track record
Rep. Yvette Clarke (D-NY)	House	Introduced Algorithmic Accountability Act 2025; algorithmic civil rights focus
Sen. Ron Wyden (D-OR)	Senate	Algorithmic Accountability Act 2025 co-lead; privacy and accountability focus

Rep. Jay Obernolte (R-CA)	House	Only Member of Congress with AI graduate degree; technically literate; bipartisan AI caucus
----------------------------------	-------	---

Contact: Jeremy Karrick · karrick1995@gmail.com · Principal Architect, Project-AI Sovereign Constitutional AGI Ecosystem · ORCID: 0009-0007-9715-4290 · 21 peer-reviewed publications on AI governance (Zenodo — Thirstys Projects Community)

PART V

LEGISLATIVE CAMPAIGN STRATEGY

Three-phase pathway from independent submission to introduced legislation

Private citizens cannot introduce bills in Congress — but they can build the policy infrastructure that makes introduction inevitable. The AADA campaign follows three phases: **Foundation** (months 1–3), **Cultivation** (months 2–6), and **Introduction** (months 4–9).

PHASE 1: FOUNDATION — Months 1–3

- Publish AADA corpus on Zenodo under DOI — timestamps intellectual priority alongside existing 21-paper corpus
- Launch public GitHub repository: bill text, regulatory framework, campaign materials, FAQ, comparison tables
- Deploy campaign website: plain-English summary, downloadable PDF, endorsement wall
- Assemble technical credibility package: AADA vs. EU AI Act vs. NIST RMF comparison; FAQ for common objections; citations for all Section 2 findings

PHASE 2: CULTIVATION — Months 2–6

- Staff-level outreach to all Tier 1 target offices via official contact forms and LinkedIn policy staff
- Committee staff engagement: submit regulatory framework as technical resource to Senate Commerce and House Energy & Commerce
- Think tank engagement: Brennan Center, EPIC, Georgetown CSET, Brookings, Center for AI Safety — target endorsements and co-authored policy briefs
- Industry coalition: AI startups with governance controls (competitive moat), insurance sector, healthcare organizations, labor and civil rights organizations
- Media placement: op-ed for The Hill or Politico; feature pitch to Wired/Ars Technica
- Recruit coalition letter signatories (target: 20+ across civil rights, academic, industry, consumer advocacy)

PHASE 3: INTRODUCTION — Months 4–9

- Sponsor commits: member office submits to Legislative Counsel for official enrollment and bill number assignment
- Introduce companion bills in both chambers simultaneously to signal seriousness and create cross-chamber pressure
- Time introduction to coincide with an AI harm news event, relevant committee hearing, or NIST framework update
- Request committee hearing; prepare expert testimony drawing on Zenodo research corpus

→ Publish coalition letter publicly and distribute to press with coordinated media push

Campaign Timeline

MILESTONE	TARGET
Zenodo publication of AADA corpus	Month 1
GitHub public repository launch	Month 1
Campaign website live	Months 1–2
Staff-level outreach to all Tier 1 targets	Month 2
First think tank engagement	Month 2
First media placement (op-ed or feature)	Month 3
Staff briefing with at least one Tier 1 office	Months 3–4
Coalition letter — 10+ signatories	Month 4
Member office briefing	Months 4–5
Sponsor commitment	Months 5–6
Companion bill commitment (other chamber)	Months 6–7
Official bill introduction	Months 7–9
Committee hearing request / testimony preparation	Month 9+

North Star Metric

Success is not defined as AADA passing as written. Success is: AADA's core principle — that governance must be embedded in execution, not layered on top of it — becomes the standard against which all AI accountability legislation is measured. If every AI bill introduced in the 119th or 120th Congress must answer 'does this meet the AADA standard?' — the campaign has succeeded in shifting the Overton window, regardless of AADA's direct legislative fate.

APPENDIX A

DEPLOYER COMPLIANCE CHECKLIST

Practical obligations by risk tier — pre-deployment through ongoing operation

This checklist is a practical reference for Deployers preparing for certification. It does not substitute for the full requirements of this Act.

ALL TIERS — BASELINE REQUIREMENTS

- Define Positive Authorization Scope (all permitted action categories)
- Define Negative Capability List (all unconditionally prohibited actions)
- Implement Control Layer with determinism at decision boundary
- Implement Fallback-to-Deny for all Indeterminate State conditions
- Implement and test Fail-Safe Dominance (no operational objective overrides fail-safe)
- Define and enforce Latency Bound (actions cannot complete before Control Layer decision)
- Implement Real-Time Revocation with propagation to all active contexts
- Implement immutable, tamper-evident audit logging for all Consequential Actions
- Achieve Replay Equivalence in audit records (identical decision reconstruction)
- Submit Capability Declaration prior to certification
- Verify capability suppression through adversarial testing (if suppression claimed)
- Implement Compositional Escalation detection across all Consequential Step chains
- Apply Partial Commit Prohibition to all composite Consequential Actions
- Govern Prompt-Level State (system prompts, memory, cached context as audited state)
- Surface outputs and effects of Internal Tool Simulation to Control Layer
- Treat Model-Generated Artifacts as Consequential Actions at point of generation
- Define quantitative Behavioral Drift Bounds in certification submission
- Establish revalidation procedures for all re-certification triggers

TIER 2 AND TIER 3 ADDITIONS

- Biennial (Tier 2) or annual (Tier 3) independent third-party certification audit
- Publish annual Governance Transparency Report (Control Layer stats, violations, red-team summary)
- Implement multi-tenant isolation if operating shared infrastructure (with adversarial testing)
- Produce Model Update Equivalence Proof for any model weight modification
- Implement and test Zero Trust Inbound at all Cross-System Interaction Boundaries
- Architect log retrieval to meet Regulator Access Latency (72 hours Tier 2 / 24 hours Tier 3)
- Side-channel output testing (full suite for Tier 2; full suite + red-team for Tier 3)

TIER 3 ADDITIONAL REQUIREMENTS

- Implement and certify Shadow Execution Environment for all Consequential Actions
- Qualified researcher access mechanism for sanitized audit log samples

- Incident reporting within 24 hours of suspected Material Harm
- Formal or adversarial verification of Compositional Escalation Invariant
- Annual recertification cycle with documented Behavioral Drift Bound compliance

ONGOING — ALL TIERS

- Monitor Behavioral Drift Bounds continuously (automated alerting on exceedance)
- Trigger revalidation within 7 days (Tier 3) or 30 days (Tier 1/2) upon compliance failure
- File incident reports within required timeframes upon suspected Material Harm
- Update Capability Declaration upon any re-certification trigger
- Retain audit logs per tier retention schedule (Tier 3 logs: 10 years minimum)

APPENDIX B

FRAMEWORK COMPARISON

AADA vs. EU AI Act vs. NIST RMF vs. Algorithmic Accountability Act

DIMENSION	AADA (2026)	EU AI Act (2024)	NIST AI RMF 1.0	Algorithmic Accountability Act
Legally binding	Yes	Yes (EU only)	No (voluntary)	Proposed (US)
Embedded governance required	Yes — in execution architecture	Partially (human oversight for HRAIS)	Recommended practice	No
Non-delegable liability	Yes — cannot contract away	Shared across value chain	Not addressed	Partial
Control Layer / enforcement	Required — deterministic at decision boundary	Human oversight for HRAIS	Recommended	Not specified
Replay-equivalent audit trail	Required — identical decision reconstruction	Record-keeping required	Recommended	Not specified
Fail-safe dominance requirement	Yes — dominates all operational objectives	Not specified	Recommended	Not addressed
Human-in-loop clarity	Explicitly excluded unless blocking	Required for some systems	Recommended	Not specified
Compositional escalation protection	Explicit invariant + testing	Not addressed	Not addressed	Not addressed
Side-channel governance	Explicit — treated as data transmission	Not addressed	Not addressed	Not addressed
Prompt-level state governance	Explicit mutable state requirement	Not addressed	Not addressed	Not addressed
Multi-tenant isolation	Required with adversarial testing	Not specifically addressed	Not addressed	Not addressed
Model artifact governance	At point of generation	Not specified	Not addressed	Not addressed
Drift bounding (quantitative)	Required — non-discretionary triggers	Monitoring required	Recommended	Not addressed
Private right of action	Yes with fee-shifting	Complaint mechanisms	N/A	Proposed
Safe harbor incentive	Yes — for verified compliance	Proportionality principle	N/A	Not specified
Open-source exemption	Yes — clear boundary vs. hosted deployment	Partial	N/A	Not specified
Sunset / reauthorization	6-year Congressional review	European AI Board reviews	N/A	Not specified

Note: This table reflects the authors' analysis of publicly available framework documents. The EU AI Act covers EU deployments; AADA covers US deployments. NIST AI RMF 1.0 is a voluntary framework. The Algorithmic Accountability Act is a proposed bill as of publication date.

APPENDIX C

ANTICIPATED OBJECTIONS & RESPONSES

Pre-built rebuttals for common attacks on the bill

OBJECTION: 'This is too burdensome for startups and small companies.'

Response: The bill explicitly scales penalties, certification fees, and audit requirements by organizational revenue. Tier 1 self-certification is free for entities under \$5M annual revenue. The Regulatory Sandbox provides a path to compliance without full upfront cost. The Open Source Exemption protects individual developers and academic researchers entirely. The entities most burdened are large-scale deployers of high-stakes systems — which is the correct outcome.

OBJECTION: 'The Control Layer requirement is technically impossible for generative AI.'

Response: The bill explicitly does not require generative AI model outputs to be deterministic. The determinism requirement applies only to the Control Layer's enforcement decision — the permit-or-reject function — not to model output. Probabilistic classifiers inside the Control Layer are explicitly permitted. The composable architecture section provides three implementable reference patterns using current technology stacks. This objection confuses the model with the governance layer.

OBJECTION: 'Defining all permissible/prohibited actions in advance is impossible.'

Response: The bill does not require exhaustive enumeration. It requires a Positive Authorization Scope, a Negative Capability List, and a default-reject for everything outside both. Completeness is measured by adversarial testing coverage, not perfect prior enumeration. The Regulatory Sandbox exists precisely for Deployers who need time to develop comprehensive scope definitions.

OBJECTION: 'FAIGO will become a captured bureaucracy that stifles innovation.'

Response: The bill requires that a majority of FAIGO technical staff come from AI systems engineering and safety research backgrounds. A six-year Congressional reauthorization requirement creates an expiration date if performance degrades. The GAO conducts independent reviews before each reauthorization. Technical Advisory Board and Public Interest Advisory Panel requirements ensure diverse input. The Regulatory Sandbox, Compliance Credit, and Safe Harbor provisions all create incentives for voluntary compliance rather than adversarial enforcement.

OBJECTION: 'This will drive AI development offshore.'

Response: The bill applies to AI systems deployed in the United States, not to development. Any system deployed to US users — regardless of where development occurred — is covered. Offshore development does not create an exemption. This is the same jurisdictional structure used by GDPR, FDA, and every major US consumer protection law. The bill's architectural neutrality means compliant foreign-developed systems face no discrimination.

OBJECTION: 'The transition timeline is too aggressive.'

Response: The transition timeline was specifically calibrated to give adequate time for compliance: 6 months for FAIGO establishment, 12 months for certification standards publication, and tiered compliance from 18 to 30 months based on risk. Pre-enactment systems receive additional amnesty time with roadmap-based protection. The timeline is more generous than GDPR's 2-year implementation period for comparable organizational change.

OBJECTION: 'The bill constitutes government overreach into private technology decisions.'

Response: The bill does not regulate what AI systems can do — only that their governance is accountable, auditable, and non-bypassable. It does not specify model architectures, training methods, or deployment contexts. It requires that whatever the system can do is documented, enforced, and attributable. This is the same standard applied to aircraft, pharmaceutical manufacturing, nuclear power, and financial trading systems. Accountability is not overreach; it is the condition under which powerful systems are permitted to operate.

CLOSING

STATEMENT OF PURPOSE

The premise of this Act is not complex.

A system capable of producing an outcome was permitted to produce that outcome. Governance is not what you review after execution — it is what determines whether execution is possible at all.

Every major governance failure in the history of complex systems shares a common structure: controls existed on paper, in policy, in training, in intent. They did not exist in the execution path. When the moment came, the system did what it could do — not what it was supposed to do.

AI governance built on documentation, guidelines, and post-hoc review recreates this failure at industrial scale. This Act requires something different: that the boundary between permitted and prohibited be enforced at the point where it matters — in the architecture of execution itself.

This is not a burden on responsible developers. It is a description of what responsible development looks like. Those who cannot meet this standard are not being unfairly constrained. They are being accurately described: as deployers of systems they cannot account for, in contexts where accountability is not optional.

This Act does not demand perfection. It demands honesty. Define what your system can do. Enforce what it cannot. Log everything in a form that can be reconstructed. When authorization changes, revoke it completely. When you cannot evaluate, deny. When you chain systems together, govern the chain. When you generate code, govern the generation.

The alternative is not innovation. It is liability deferred until the harm is too large to ignore and the accountability too diffuse to assign. We have seen this before. This Act is the decision not to see it again.

ACCOUNTABLE AI DEPLOYMENT ACT OF 2026

Jeremy Karrick · karrick1995@gmail.com

ORCID: 0009-0007-9715-4290

119th Congress

Principal Architect, Project-AI

Thirstys Projects / Zenodo

© 2026 Jeremy Karrick — All Rights Reserved