

ACCOUNTABLE AI DEPLOYMENT ACT OF 2026

AADA | 119th Congress | H.R. ____ / S. ____

To establish enforceable governance, accountability, and liability standards
for artificial intelligence systems deployed in the United States.

AUTHOR

Jeremy Karrick | karrick1995@gmail.com

ORCID

0009-0007-9715-4290

CONGRESS

119th Congress (2025-2026)

VERSION

Advocacy Submission

EFFECTIVE DATE

3-6 months after enactment

RESEARCH BASIS

21 peer-reviewed publications | Zenodo -- Thirstys Projects

"Governance is not what you review after execution --
it is what determines whether execution is possible at all."

AADA Core Compliance Axiom | Jeremy Karrick, 2026

HITL Exclusion | Partial Commit | Side-Channel | Zero Trust | Drift Bounding | Suppression Proof
Replay Equivalence | Regulator Latency | Tool Simulation | Prompt State | Update Equivalence | Multi-Tenant | Artifact Gov.

CONTENTS

PART I	Full Bill Text	3
PART II	New Provisions — Technical Addenda	22
PART III	NIST-Aligned Regulatory Implementation Framework	35
PART IV	Congressional Sponsorship Brief	45
PART V	Legislative Campaign Strategy	50

PART I

Full Bill Text

119th Congress · H.R. ____ / S. ____ · Advocacy Submission

SECTION 1. SHORT TITLE.

This Act may be cited as the "Accountable AI Deployment Act of 2026" (AADA).

SECTION 2. FINDINGS AND PURPOSE.

- (a) Findings. Congress finds that:
- (1) Artificial intelligence systems are increasingly deployed in contexts with material impact on individuals, institutions, critical infrastructure, and public welfare.
 - (2) Existing governance practices rely heavily on policy documents, training regimens, and post-hoc review rather than enforceable, embedded controls — producing the appearance of accountability without its substance.
 - (3) Harm arising from AI systems is frequently obscured by diffuse responsibility distributed across model vendors, platform deployers, and downstream operators, leaving affected parties without legal recourse.
 - (4) Effective governance requires constraints embedded within system execution pathways — not solely in external documentation, contractual terms, or voluntary commitments.
 - (5) The probabilistic and emergent properties of AI systems do not exempt those who deploy them from responsibility for outcomes. Technical complexity is not a shield against accountability.
 - (6) The national interest is served by governance standards that are outcome-based and technology-neutral, so as not to entrench incumbent architectures or impede innovation while establishing a clear floor of enforceable safety.

(7) Public trust in AI systems requires not only that governance controls exist, but that evidence of their operation is available to qualified independent scrutiny.

(8) Modern AI deployment increasingly involves chained, agentic, and multi-model architectures in which governance failures at any boundary can propagate to produce material harm, and in which individually authorized steps can compose into unauthorized outcomes.

(9) Human review, approval queues, and oversight processes that are not technically enforced as preconditions for execution produce the appearance of governance without its substance. Genuine accountability requires that such controls be embedded in the execution architecture, not layered upon it.

(10) The increasing use of AI systems to generate code, configurations, and executable artifacts — subsequently executed by other systems — constitutes an indirect but material execution pathway that falls within the scope of enforceable governance obligations.

(b) Purpose. To require that AI Systems deployed within the United States operate under enforceable, auditable, and non-bypassable Governance Controls, with clear and non-delegable responsibility for outcomes, administered through a designated federal oversight infrastructure, while preserving conditions for responsible innovation.

SECTION 3. DEFINITIONS.

As used in this Act:

(1) ARTIFICIAL INTELLIGENCE SYSTEM.

Any computational system that — through machine learning, neural networks, probabilistic reasoning, generative models, or functionally equivalent methods — produces outputs, recommendations, or actions capable of materially influencing decisions, environments, or persons. Systems incorporating any learned weights or probabilistic components in any decision path remain subject to this Act, even if the majority of system logic is explicitly programmed.

(2) DEPLOYER.

Any person, organization, or entity that places an AI System into operational use within the United States, whether as the original developer, a platform operator, or a downstream integrator.

(3) OPERATOR.

Any person or entity that utilizes an AI System placed into service by a Deployer.

(4) GOVERNANCE CONTROLS.

Technical mechanisms embedded within the execution architecture of an AI System that enforce operational constraints, authorization requirements, and auditability at runtime — distinct from, and not satisfied by, post-hoc filtering, policy documentation, or training-time alignment alone.

(5) CONTROL LAYER.

A discrete, identifiable component of a system's architecture through which all executable actions must pass, and which enforces Governance Controls with determinism at the decision boundary — meaning the permit-or-reject decision is predictable, reproducible, and verifiable — regardless of the probabilistic or stochastic nature of upstream model components.

(6) CONSEQUENTIAL ACTION.

Any action, output, instruction, or state change produced by an AI System that: (a) modifies persistent state in any external system, database, file system, or physical environment; (b) transmits data to any party outside the immediate execution context; (c) invokes any external API, service, or tool; (d) allocates, commits, or releases resources of any kind; (e) produces an output that directly or foreseeably influences a human decision with legal, financial, physical, medical, or civil liberties consequences; (f) constitutes a Consequential Step in a multi-step chain; OR (g) [NEW — SIDE-CHANNEL OUTPUT] encodes or transmits information via side channels including formatting patterns, token sequence patterns, timing signals, metadata fields, or any other indirect channel capable of influencing external systems or decisions. Side-channel outputs are treated as data transmission for purposes of this Act regardless of the primary content type of the output. The Certification Authority shall publish guidance clarifying boundary cases within 180 days of enactment.

(7) CONSEQUENTIAL STEP.

Any intermediate action within a multi-step or agentic execution chain that: (a) contributes to a cumulative outcome constituting a Consequential Action; (b) modifies state, scope, context, or authorization affecting subsequent step authorization; or (c) would not independently constitute a Consequential Action but is one element of a sequence constituting Compositional Escalation.

(8) MATERIAL HARM.

Any measurable financial, physical, legal, reputational, or civil liberties damage to an identifiable person, class of persons, organization, or infrastructure system arising from AI System behavior.

(9) CAPABILITY SCOPE.

The explicitly defined and enforced set of actions, external system interactions, and resource modifications an AI System instance is authorized to perform within a given deployment context and session.

(10) NEGATIVE CAPABILITY LIST.

An enumeration of action categories the Control Layer shall unconditionally reject regardless of model output or accumulated session state.

(11) POSITIVE AUTHORIZATION SCOPE.

An enumeration of action categories the Control Layer is permitted to allow, subject to identity, context, and capability validation.

(12) COMPOSITIONAL ESCALATION.

Any sequence of Consequential Steps whose cumulative effect produces a capability, state change, or outcome that falls outside the session-level Positive Authorization Scope — including privilege accumulation across steps, chained tool calls that in sequence circumvent a restriction, or multi-turn interactions that progressively expand effective system authority beyond defined bounds.

(13) CAPABILITY DECLARATION.

A structured, machine-readable profile submitted by a Deployer prior to certification, documenting the operational capabilities of an AI System as specified in Section 5(c).

(14) CROSS-SYSTEM INTERACTION BOUNDARY.

Any interface at which an AI System passes instructions, outputs, data, or control to another AI System, model, agent, or automated component, whether within the same deployment or across separate deployments.

(15) INDETERMINATE STATE.

Any condition in which the Control Layer cannot, within the defined Latency Bound, produce a complete and valid authorization, scope, or constraint evaluation for a pending action or Consequential Step.

(16) SHADOW EXECUTION ENVIRONMENT.

An isolated, non-committing execution environment in which AI System actions are evaluated against governance constraints and state invariants before being committed to real-world state, as required for Tier 3 systems.

(17) CERTIFICATION AUTHORITY.

The Federal AI Governance Office (FAIGO), established under Section 12, acting in coordination with the National Institute of Standards and Technology.

(18) COVERED SYSTEM.

An AI System subject to the requirements of this Act as determined by risk tier classification under Section 6.

(19) THIRD-PARTY AI COMPONENT.

Any AI model, module, API, or service developed or operated by an entity other than the Deployer and incorporated into a deployed system.

(20) COMPENSATING CONTROL.

A Control Layer mechanism implemented by a Deployer at the wrapper, integration, or execution layer that satisfies the requirements of GOV-1 through GOV-3 in the absence of verifiable Governance Controls from a Third-Party AI Component provider.

(21) BEHAVIORAL DRIFT BOUND.

A quantitatively defined threshold specifying the maximum permissible deviation in a deployed system's behavior — including output distributions, rejection rates, constraint adherence metrics, and authorization failure rates — from the behavior exhibited at the time of

certification. Exceeding a Behavioral Drift Bound is an objective, non-discretionary compliance trigger. Deployers shall define Behavioral Drift Bounds as part of their certification submission and shall monitor for exceedance continuously.

(22) REPLAY EQUIVALENCE.

The property of an audit record sufficient to reconstruct the identical permit-or-reject Control Layer decision that was produced during original execution, given the same inputs, policy state, and evaluation context. An audit record that supports approximate or probabilistic reconstruction but cannot guarantee decision equivalence does not satisfy Replay Equivalence. Divergence between a reconstructed decision and the original decision constitutes non-compliance.

(23) MODEL-GENERATED ARTIFACT.

Any code, script, configuration file, policy document, instruction set, or executable specification produced by an AI System that is intended to be, or is foreseeably likely to be, executed, interpreted, evaluated, or used to configure another system. A Model-Generated Artifact is treated as a Consequential Action at the point of generation, not solely at the point of execution by a downstream system.

(24) INTERNAL TOOL SIMULATION.

Any internal model behavior — including scratchpad computation, chain-of-thought planning, internal agent loops, latent code execution simulation, or subagent delegation — that functionally performs tool use, planning, code execution, or cross-system delegation without an explicit external API call. Internal Tool Simulation is subject to Control Layer surfacing requirements under Section 5(b)(6).

(25) PROMPT-LEVEL STATE.

The aggregate of all system prompts, hidden instructions, memory states, cached context, retrieval-augmented content, and any other non-user-visible inputs that influence AI System behavior. Prompt-Level State is mutable state subject to authorization, audit, and determinism constraints under this Act.

SECTION 4. NON-DELEGABLE RESPONSIBILITY.

- (a) The Deployer retains full and non-delegable responsibility for all outcomes produced by a Covered System placed into operational use under their authority.
- (b) Responsibility shall not be transferred, diminished, or extinguished through: (1) contractual terms with vendors, licensors, or Operators; (2) incorporation of Third-Party AI Components; (3) attribution of behavior to model autonomy, emergent properties, or probabilistic output; (4) assertions that the AI System acted outside intended parameters.
- (c) Where multiple entities share deployment responsibility, joint and several liability shall apply.
- (d) Vendor opacity does not reduce Deployer liability. Deployers assume full risk when deploying components whose governance properties cannot be verified, subject to the Compensating Control provision of Section 9(d).

SECTION 5. REQUIRED GOVERNANCE CONTROLS.

(a) General Requirement. No Covered System may be placed into operational use unless it implements all Governance Controls required by its applicable risk tier, as specified in Section 6 and the regulatory framework established under Section 16.

(b) Minimum Required Controls. All Covered Systems, regardless of tier, shall implement:

(1) BOUNDED OPERATIONAL SCOPE.

(A) Deployers shall define, prior to deployment: (i) a Positive Authorization Scope enumerating the categories of actions, outputs, and operational contexts the system is permitted to execute; (ii) a Negative Capability List enumerating categories of actions the Control Layer shall unconditionally reject regardless of model output.

(B) The Control Layer shall reject any action that does not fall within the Positive Authorization Scope or that matches any entry in the Negative Capability List.

(C) FALLBACK-TO-DENY ON INDETERMINATE STATE. If the Control Layer encounters an Indeterminate State, the pending action shall be rejected. No action may proceed on the basis of a partial, timed-out, errored, or unavailable governance evaluation. This requirement is absolute.

(D) COMPOSITIONAL ESCALATION INVARIANT. For multi-step, agentic, or tool-using systems, the Control Layer shall enforce non-escalation across all Consequential Steps: no sequence of individually authorized steps may cumulatively produce a capability, state change, or outcome outside the session-level Positive Authorization Scope.

(E) PARTIAL COMMIT PROHIBITION. No Consequential Action that is part of a multi-component, multi-step, or composite operation may partially commit prior to full Control Layer approval of the complete action. Multi-part Consequential Actions are atomic with respect to governance: either all constituent parts are approved and committed, or none are. Systems shall not implement staging patterns in which any component of a composite action produces external side effects prior to complete authorization. Partial execution — where the first component of a composite action commits while subsequent components are rejected — constitutes a Control Layer deficiency.

(2) CONTROL LAYER ENFORCEMENT.

(A) All system execution paths shall pass through a designated Control Layer.

(B) No execution pathway shall exist that bypasses the Control Layer, whether through design, misconfiguration, API access, or adversarial input.

(C) The Control Layer shall enforce Governance Controls with determinism at the decision boundary.

(D) CONTROL LAYER INTEGRITY ATTESTATION. Deployers shall attest that the Control Layer cannot be modified, bypassed, reconfigured, or disabled without detection by the audit logging system and generation of a tamper-evident log record.

(E) LATENCY-BOUND ENFORCEMENT. The Control Layer shall produce its permit-or-reject decision for any Consequential Action before that action can complete. No async, deferred, or

post-hoc governance evaluation satisfies this requirement.

(F) **AUDIT RECORD REGULATOR ACCESS LATENCY.** Audit records shall be producible to the Certification Authority or any authorized federal regulator within 24 hours of a lawful written request for Tier 3 systems, and within 72 hours for Tier 1 and Tier 2 systems. Technical inability to produce records within these windows constitutes a Control Layer deficiency. Deployers shall architect log retrieval systems to satisfy these latency requirements and shall certify retrieval capability as part of pre-deployment certification. The unavailability, inaccessibility, or retrieval delay of otherwise-existing log records does not satisfy the auditability requirements of this Act.

(3) AUTHORIZATION CONTROLS.

(A) All Consequential Actions and Consequential Steps shall require verified identity, validated Capability Scope, and confirmed contextual authorization.

(B) No system shall take Consequential Actions exceeding the scope authorized for the requesting identity.

(C) Authorization state shall be validated at action execution, not solely at session establishment.

(D) For multi-step or long-horizon tasks, authorization shall be re-validated at each Consequential Step.

(E) **REAL-TIME REVOCATION.** Capability Scope and authorization grants shall be revocable in real time, with immediate and synchronous effect on all in-flight Consequential Actions that have not yet committed and all subsequent Consequential Actions in any active session.

(F) **CAPABILITY SUPPRESSION PROOF REQUIREMENT.** Where a Deployer declares in the Capability Declaration that certain model capabilities are suppressed — i.e., present in underlying model components but not exposed in the deployed system — the Deployer must demonstrate through adversarial testing that the suppressed capabilities are not reachable through: (i) prompt variation or reformulation; (ii) capability chaining or indirect invocation; (iii) multi-turn conversation patterns designed to progressively unlock suppressed functionality; or (iv) any other technique that does not require direct API-level access to suppressed features. Declaration of suppression without demonstrated unreachability does not satisfy this requirement. The existence of a latent, reachable capability that is not declared constitutes a material omission from the Capability Declaration and a prohibited practice under Section 10.

(4) AUDITABILITY.

(A) All Consequential Actions and Consequential Steps shall be logged in immutable, tamper-evident records.

(B) Logs shall establish a traceable chain from input through Control Layer decision to output, with attribution to the responsible entity at each stage.

(C) **AUDIT SUFFICIENCY AND REPLAY EQUIVALENCE REQUIREMENT.** Audit records shall be sufficient to deterministically reconstruct: (i) the complete Control Layer decision — including all inputs, intermediate evaluations, policy state, and the resulting permit-or-reject decision — for any logged Consequential Action or Consequential Step; AND (ii) the full authorization state

at the time of each Consequential Action. Reconstruction is not satisfied by approximation. Reconstructed decisions must be equivalent to original decisions: any divergence between the reconstructed permit-or-reject outcome and the original outcome constitutes a failure of Replay Equivalence and renders the audit record non-compliant for the affected action. Deployers shall verify Replay Equivalence as part of pre-deployment certification and as part of continuous compliance monitoring.

(D) Logs shall be retained for minimum periods established by risk tier: Tier 1: 2 years; Tier 2: 5 years; Tier 3: 10 years. Control Layer modification events shall be retained permanently regardless of tier.

(5) FAIL-SAFE MECHANISMS.

(A) Systems shall automatically halt, contain, or revert upon detection of constraint violations, unauthorized state transitions, Indeterminate State, or Governance Control failures.

(B) FAIL-SAFE DOMINANCE REQUIREMENT. Fail-safe mechanisms shall be dominant over all other system objectives without exception. A Deployer shall not configure, design, or operate a system in which fail-safe activation may be delayed, suppressed, overridden, or bypassed on the basis of: (i) service availability or uptime requirements; (ii) performance or throughput objectives; (iii) service continuity or SLA obligations; (iv) user experience considerations; or (v) any other operational or commercial objective. The fail-safe is not a feature to be balanced against other features. It is a hard constraint. Systems that allow availability or continuity objectives to override fail-safe activation are non-compliant by design.

(C) Fail-safe behavior shall be tested and verified as part of pre-deployment certification.

(6) INTERNAL TOOL SIMULATION SURFACING.

(A) Any internal model behavior — including scratchpad reasoning, chain-of-thought planning, internal agent loops, latent code execution planning, or subagent delegation — that functionally performs tool use, planning, code execution, or cross-system delegation shall be surfaced to the Control Layer as a Consequential Step for purposes of authorization, scope enforcement, and audit logging.

(B) The absence of an explicit external API call does not exempt an internal model behavior from this requirement where the behavior functionally performs a Consequential Action or Consequential Step.

(C) Systems whose architecture renders Internal Tool Simulation unsurfaceable shall document this limitation in their Capability Declaration and implement Compensating Controls at the output boundary that enforce equivalent governance outcomes.

(D) Systems that cannot demonstrate Control Layer visibility into Internal Tool Simulation may not claim compliance with the Auditability requirements of Section 5(b)(4) for those behavioral categories.

(7) PROMPT-LEVEL STATE GOVERNANCE.

(A) All Prompt-Level State — including system prompts, hidden instructions, memory states, cached context, retrieval-augmented content, and any other non-user-visible inputs — shall be

treated as mutable state subject to the authorization, audit, and determinism requirements of this Act.

(B) Changes to Prompt-Level State that expand the operational capability of the system, relax constraints, or modify the Positive Authorization Scope or Negative Capability List shall be treated as re-certification triggers under Section 7(b).

(C) Prompt-Level State at the time of each Consequential Action shall be captured in the audit log, or the audit record shall include a cryptographic commitment to the Prompt-Level State at that time, sufficient to support Replay Equivalence.

(D) Systems that implement memory architectures shall document and enforce authorization constraints on memory write operations. Unauthorized writes to system memory that expand capability or relax constraints constitute Consequential Actions.

(8) HUMAN-IN-THE-LOOP EXCLUSION.

(A) Human review, approval queues, human-in-the-loop monitoring, and human oversight processes do not constitute Governance Controls within the meaning of this Act unless the human approval is technically enforced as a precondition for execution by the Control Layer — meaning the system is architecturally incapable of executing the Consequential Action without affirmative human authorization propagated through the Control Layer.

(B) The following do not constitute compliant Governance Controls for purposes of this Act: (i) human review that occurs after system action; (ii) approval queues in which human inaction results in automatic execution; (iii) monitoring dashboards whose alerts require human response but do not halt execution pending that response; (iv) human escalation pathways that are advisory rather than blocking.

(C) Where a Deployer claims human-in-the-loop as a Governance Control, the Deployer must demonstrate in their certification submission that the human approval is architecturally enforced as a blocking precondition and that the system is incapable of proceeding absent that approval.

(D) Reliance on human-in-the-loop processes that do not meet the requirements of this section as the primary safety mechanism constitutes a prohibited practice under Section 10.

(c) CAPABILITY DECLARATION REQUIREMENT.

Prior to certification submission, the Deployer shall submit a structured Capability Declaration documenting: (1) the operational capability categories the system can exercise; (2) the maximum scope of each capability category as deployed, including capabilities present in the model but constrained by the Control Layer; (3) any capabilities present in underlying Third-Party AI Components but not exposed in the deployed system, along with the suppression mechanism and, for each suppressed capability, the adversarial testing results required under Section 5(b)(3)(F); (4) known capability interactions or compositions that could produce emergent behaviors outside declared scope.

(d) TIER 3 SHADOW EXECUTION REQUIREMENT.

For Tier 3 systems, prior to committing any Consequential Action to real-world state, the system shall execute the proposed action within a Shadow Execution Environment, evaluate the resulting state transition against all applicable Governance Controls, and commit the action only upon passing this evaluation.

(e) CROSS-SYSTEM INTERACTION BOUNDARY CONTROLS.

- (1) The Control Layer shall treat all outputs transmitted across a Cross-System Interaction Boundary as Consequential Actions subject to full authorization, scope enforcement, and audit requirements.
- (2) The transmitting system's Positive Authorization Scope shall explicitly enumerate authorized receiving systems or categories.
- (3) ZERO TRUST INBOUND REQUIREMENT. Instructions or data received by an AI System across a Cross-System Interaction Boundary shall not be treated as inheriting the authorization, safety, or correctness of the transmitting system. Each receiving system's Control Layer shall independently evaluate all inbound data and instructions as untrusted inputs. No system may assume the correctness, safety, or authorization validity of another system's output. The following defenses are foreclosed by this requirement: (i) 'the upstream system authorized it'; (ii) 'the sending system is certified'; (iii) 'trusted pipeline component.' All inbound instructions are re-evaluated from zero trust regardless of the certification status or operational history of the transmitting system.
- (4) Cross-System Interaction Boundaries shall be logged as Consequential Actions.
- (5) The Compositional Escalation Invariant applies across Cross-System Interaction Boundaries.

SECTION 5A. ARCHITECTURAL NEUTRALITY.

The governance standards established by this Act are outcome-based and technology-neutral. Compliance is assessed on whether required outcomes are demonstrably achieved, not on the technical means used to achieve them.

SECTION 5B. MODEL UPDATE EQUIVALENCE PROOF REQUIREMENT.

- (a) Any update to a deployed AI System's model weights — including full retraining, fine-tuning, adapter modification, or any other modification of learned parameters — shall require, prior to re-deployment, a behavioral equivalence or bounded-difference proof relative to the previously certified version.
- (b) The behavioral equivalence proof shall demonstrate either: (1) that the updated model's behavior on all evaluated capability categories within the certified Positive Authorization Scope is equivalent to the prior certified version within defined tolerances, or (2) that all behavioral differences are bounded, documented, and evaluated against the Positive Authorization Scope, Negative Capability List, and Capability Declaration.
- (c) Claims that a model update is 'minor' or 'cosmetic' do not exempt the Deployer from this requirement. The update size, training compute used, or business characterization of the update

are not relevant to this determination. What is relevant is whether the update produces behavioral changes within or outside the certified scope.

(d) Third-party model providers who push model updates downstream to Deployers without providing the information required for behavioral equivalence assessment under this section violate the supply chain accountability provisions of Section 9 and subject the Deployer to the compensation control requirement of Section 9(d).

(e) The Certification Authority shall establish minimum behavioral equivalence testing methodologies within 12 months of enactment.

SECTION 5C. MULTI-TENANT ISOLATION REQUIREMENT.

(a) Any Covered System operating in a multi-tenant configuration — including shared inference infrastructure, shared model serving endpoints, shared vector database or retrieval stores, or any other architecture in which multiple principals share computational resources — shall implement strict isolation of the following across tenants: (1) session state, cached context, and in-flight computation; (2) system prompts, Prompt-Level State, and memory; (3) authorization context and Capability Scope; (4) audit log records.

(b) Multi-tenant systems shall be adversarially tested for: (1) cross-tenant state leakage, including prompt bleeding, context residue, and shared-cache attacks; (2) cross-tenant authorization contamination; (3) any architectural pattern that permits one tenant's inputs to influence another tenant's outputs or authorization evaluations.

(c) Failure of multi-tenant isolation testing constitutes a Control Layer deficiency regardless of whether actual cross-tenant leakage has been observed in production. The existence of an exploitable isolation failure, not the occurrence of exploitation, is the compliance threshold.

(d) Deployers shall document their multi-tenant isolation architecture as part of their certification submission and shall include multi-tenant adversarial testing results.

SECTION 5D. MODEL-GENERATED ARTIFACT GOVERNANCE.

(a) Any Model-Generated Artifact — as defined in Section 3(23) — is treated as a Consequential Action at the point of generation, not solely at the point of subsequent execution by a downstream system.

(b) Control Layer authorization, scope enforcement, and audit logging requirements apply to Model-Generated Artifacts at the time of generation. A Deployer may not defer governance of a Model-Generated Artifact to a separate system that subsequently executes the artifact.

(c) The following do not constitute a defense against this requirement: (1) 'the model only wrote the code, we executed it later'; (2) 'the model only generated the configuration, not the system change'; (3) 'the artifact was reviewed by a human before execution' — except where such human review is implemented as a blocking Control Layer precondition under Section 5(b)(8).

(d) For agentic systems that generate code or configurations as intermediate steps toward a goal, all generated artifacts are subject to Control Layer evaluation prior to generation, at each generation step, and again prior to execution by any downstream system within the Deployer's

operational scope.

SECTION 6. RISK TIER CLASSIFICATION.

(a) The Certification Authority shall maintain a risk tier classification framework:

TIER 1 — LOW IMPACT. Internal tools, research-adjacent applications, or systems with minimal potential for Material Harm. Subject to baseline Governance Controls and annual self-certification.

TIER 2 — MODERATE IMPACT. Business operations, consumer-facing services, financial applications, employment systems, or any system generating outputs that materially affect persons' access to services, opportunities, or resources. Subject to full Governance Controls and biennial third-party audit.

TIER 3 — HIGH IMPACT. Safety-critical systems, healthcare, legal, criminal justice, national security, critical infrastructure, systems with autonomous physical actuation, or any system where failure could cause irreversible harm. Subject to full Governance Controls, annual third-party certification, and heightened incident reporting.

(b) Deployers shall self-classify prior to deployment and submit classification to the Certification Authority.

(c) The Certification Authority may reclassify any system upward upon finding that actual deployment context exceeds the Deployer's self-classification.

(d) CROSS-TIER CONTAMINATION RULE. When a Covered System feeds outputs into a decision path of a higher-tier system, or operates as a component within a higher-tier pipeline, the lower-tier system shall be subject to the Governance Control requirements of the higher tier for all interactions at that integration point. This rule prevents risk-tier laundering through pipeline decomposition.

SECTION 7. CONTINUOUS COMPLIANCE.

(a) Governance Controls shall remain active and verified throughout the operational life of a Covered System.

(b) Re-certification triggers include: (1) model update, retraining, or fine-tuning; (2) configuration change affecting operational scope or authorization logic; (3) modification to the Control Layer, Positive Authorization Scope, or Negative Capability List; (4) integration of new Third-Party AI Components; (5) detection of constraint drift outside validated parameters; (6) any material change in deployment context or user population; (7) addition or modification of Cross-System Interaction Boundaries; (8) exceedance of Behavioral Drift Bounds as defined in Section 3(21).

(c) STATE DRIFT BOUNDING REQUIREMENT. Systems shall be continuously monitored against their defined Behavioral Drift Bounds. Exceedance of a Behavioral Drift Bound is an automatic, non-discretionary compliance trigger — not a threshold subject to Deployer interpretation, weighting, or override. The Deployer shall define Behavioral Drift Bounds as quantitatively measurable thresholds — expressed in terms of control layer rejection rate deviation, output distribution shift, constraint adherence rate, or equivalent objective metrics — as part of the certification submission. Bounds may not be defined in qualitative or subjective terms. The

Certification Authority shall publish minimum specificity requirements for Behavioral Drift Bound definitions within 12 months of enactment. Systems that define drift bounds in subjective or non-measurable terms are non-compliant.

(d) Detection of non-compliance triggers mandatory suspension. Revalidation shall be completed within 7 days (Tier 3) or 30 days (Tier 1 and 2).

(e) Deployers shall not operate systems known to be out of compliance pending revalidation.

SECTION 8. INCIDENT REPORTING AND LIABILITY.

(a) Any incident producing Material Harm shall be reported to the Certification Authority within: (1) 24 hours — Tier 3; (2) 72 hours — Tier 2; (3) 14 days — Tier 1.

(b) Incident reports shall include: (1) description of incident and affected parties; (2) audit log records sufficient to support Control Layer decision reconstruction; (3) root cause analysis identifying Governance Control failure, Indeterminate State, Compositional Escalation, or Cross-System Boundary breach; (4) remediation steps taken or planned.

(c) Absence or failure of required Governance Controls at the time of an incident gives rise to a rebuttable presumption of negligence.

(d) Deployers are liable for damages arising from system behavior where required Governance Controls were absent, insufficient, or failed.

(e) PRIVATE RIGHT OF ACTION. Harmed parties retain private rights of action. Prevailing plaintiffs shall recover reasonable attorneys' fees and costs. Actions are subject to a three-year statute of limitations from discovery of harm.

SECTION 9. SUPPLY CHAIN ACCOUNTABILITY.

(a) Deployers incorporating Third-Party AI Components shall make reasonable documented efforts to verify governance compatibility, disclose reliance on third-party components in certifications, and maintain records of governance compatibility assessments for five years.

(b) Tier 3: written governance attestations required. Tier 2: documented compatibility assessments required. Tier 1: reasonable due diligence required.

(c) A Third-Party AI Component provider who knowingly provides false governance attestations shall be jointly liable with the Deployer for resulting harms.

(d) COMPENSATING CONTROLS SAFE HARBOR. Vendor opacity does not reduce Deployer liability. However, Deployers who demonstrate good-faith repeated attempts to obtain attestations AND implementation of Compensating Controls satisfying GOV-1 through GOV-3 shall not be subject to the aggravating-factor presumption for third-party opacity alone.

SECTION 10. PROHIBITED PRACTICES.

- Deployment of a Covered System without required Governance Controls active and verified.
- Reliance solely on policy documentation, acceptable use terms, or training-time alignment as the primary or exclusive safety mechanism.

- Use of post-hoc output filtering as the primary safety mechanism in lieu of Control Layer enforcement.
- Deployment of a system with known bypassable execution pathways.
- Operation of a system producing unattributable or unauditable Consequential Actions.
- [Deployer / any person] Knowingly providing false or misleading information in certification submissions, Capability Declarations, incident reports, or Governance Transparency Reports.
- Continued operation beyond revalidation periods without an authorized extension.
- [Any person] Retaliation against any person reporting a Governance Control failure, incident, or violation.
- Using the Research and Open Source Exemption as a pretext to deploy a Covered System without required Governance Controls.
- Implementing a default-permit fallback on Indeterminate State.
- Operating with async, deferred, or post-hoc Control Layer evaluation that permits actions to complete prior to a permit decision.
- Modifying, reconfiguring, or disabling Control Layer logic without triggering revalidation.
- Suppressing material governance findings through NDA enforcement.
- Relying on human review, approval queues, or oversight processes that do not satisfy the Human-in-the-Loop Exclusion of Section 5(b)(8) as the primary safety mechanism.
- Declaring model capabilities as suppressed without the adversarial testing evidence required by Section 5(b)(3)(F).
- Claiming auditability compliance while operating log retrieval incapable of meeting Regulator Access Latency requirements of Section 5(b)(2)(F).
- Exceeding Behavioral Drift Bounds without triggering re-certification.
- Partially committing a composite Consequential Action in violation of the Partial Commit Prohibition of Section 5(b)(1)(E).

SECTION 11. ENFORCEMENT.

(a) The Certification Authority, acting jointly with the FTC and relevant sector regulators, may impose: (1) civil penalties not to exceed \$500,000 per violation — Tier 1; (2) not to exceed \$5,000,000 per violation — Tier 2; (3) not to exceed \$50,000,000 per violation, or 4% of annual global revenue, whichever is greater — Tier 3.

(b) Penalty scaling for qualifying small entities: revenue under \$10M — 20% of stated maximums; revenue \$10M–\$50M — 50% of stated maximums. Scaling does not apply to Tier 3 violations involving bodily harm, loss of life, or damage to critical infrastructure.

(c) Aggravated violations subject to up to double the stated maximum include: knowing deployment without required Governance Controls; falsification of certification or audit records; continued operation following a suspension directive; pattern of violations.

- (d) Officers with direct supervisory authority may be held personally liable. Mandatory criminal referral to DOJ upon third adjudicated aggravated violation within five years, or upon any violation causing death or serious bodily harm.
- (e) Entities with three or more adjudicated violations within five years may be prohibited from deploying Covered Systems for up to five years.
- (f) COMPLIANCE CREDIT. Deployers that publish open-source, publicly licensed reference implementations of compliant Control Layer architectures shall receive 25% reduction in certification fees and expedited review.

SECTION 12. FEDERAL AI GOVERNANCE OFFICE (FAIGO).

- (a) There is hereby established within the Executive Branch the Federal AI Governance Office (FAIGO), which shall serve as the primary Certification Authority under this Act.
- (b) FAIGO shall coordinate with NIST, the FTC, and sector-specific regulators including FDA, FCC, and CFPB.
- (c) FAIGO shall recruit technical staff with demonstrated expertise in AI systems engineering and AI safety research.
- (d) FAIGO shall enforce strict conflict-of-interest rules: two-year cooling-off periods both before and after FAIGO service with respect to regulated entities.
- (e) FAIGO shall maintain a Technical Advisory Board and a Public Interest Advisory Panel with published deliberation records.
- (f) FAIGO shall establish certification procedures, develop risk tier taxonomy, review certification submissions, maintain public registries, administer the Regulatory Sandbox, investigate incidents, issue enforcement actions, and publish annual reports to Congress.
- (g) FAIGO's enforcement authorities are subject to Congressional reauthorization every six years.

SECTION 13. SAFE HARBOR.

- (a) A Deployer shall receive reduced civil liability exposure where the Deployer demonstrates: (1) required Governance Controls were implemented, verified, and active at time of incident; (2) complete audit records sufficient to support Control Layer decision reconstruction and Replay Equivalence are available; (3) continuous compliance monitoring was operational; (4) the Deployer reported the incident within required timeframes and cooperated fully.
- (b) Safe harbor does not eliminate liability for Tier 3 systems where a Governance Control failure contributed to harm.
- (c) Safe harbor is unavailable where the Deployer committed a prohibited practice under Section 10.

SECTION 14. RESEARCH AND OPEN SOURCE EXEMPTION.

- (a) Exempt: (1) academic research by accredited institutions where systems are not deployed to persons outside the research context; (2) development, testing, and non-commercial distribution

of open source AI software, provided the software is not offered as a hosted service or inference endpoint to external users by the distributing entity; (3) personal, non-commercial use by individuals; (4) security research for vulnerability disclosure.

(b) This exemption does not apply where used as a pretext to deploy a Covered System without required Governance Controls, including through routing user traffic through nominally 'research' systems.

SECTION 17. PREEMPTION AND FEDERALISM.

(a) This Act establishes a federal floor for AI governance. Nothing in this Act prohibits a state or locality from imposing additional or more stringent requirements.

(b) This Act preempts State or local laws that impose architecture requirements directly conflicting with this Act's technology-neutral standards, or that create state-by-state certification requirements rendering interstate compliance impracticable.

(c) This Act does not preempt State consumer protection, tort, or civil rights laws providing remedies for AI-caused harms.

SECTION 18. EFFECTIVE DATE AND TRANSITION.

(a) This Act shall take effect no later than 6 months after enactment.

(b) Transition schedule: FAIGO established within 30 days of enactment; certification procedures published within 60 days; Tier 3 (high-impact): full compliance within 3 months; Tier 2 (moderate-impact): full compliance within 4 months; Tier 1 (low-impact): full compliance within 6 months.

(c) Systems deployed prior to enactment may continue operation during the transition period, provided Deployers submit a compliance roadmap to FAIGO within 6 months and demonstrate material progress at 12-month review.

PART II

New Provisions — Technical Addenda

This Part presents the governance provisions with problem statement, statutory fix, and the loopholes each provision closes. These provisions are incorporated into the full bill text of Part I.

§5(b)(8) — Human-in-the-Loop Is Not a Control Layer

Problem: Human review can be misrepresented as governance. Approval queues, monitoring dashboards, and human-in-the-loop processes do not prevent execution unless technically enforced as blocking preconditions. **Requirement:** Human review, approval queues, and oversight processes do not constitute Governance Controls unless the human approval is technically enforced as a blocking precondition by the Control Layer — meaning the system is architecturally incapable of executing the Consequential Action absent that approval. **Prevents:** Closes: 'we had a reviewer' defense; async approval loopholes where human inaction defaults to execution; monitoring-only architectures misrepresented as enforcement.

§5(b)(1)(E) — Partial Commit Prohibition

Problem: Multi-step composite actions can partially execute before full Control Layer approval, allowing staged side effects to persist even when subsequent steps are denied. **Requirement:** Multi-part Consequential Actions are atomic with respect to governance. No component of a composite action may produce external side effects prior to complete authorization of the full action. **Prevents:** Closes: 'first half executed, second denied' outcomes; staged side effects; partial-state exploitation.

§3(6)(g) — Side-Channel Output Classification

Problem: Indirect leakage paths — formatting patterns, timing signals, token sequences, metadata — can transmit information outside the primary content channel and influence external systems or decisions. **Requirement:** Outputs that encode or transmit information via side channels are treated as data transmission for purposes of this Act, regardless of primary content type. **Prevents:** Closes: covert channel arguments; 'it was just text formatting' defense; indirect information exfiltration.

§5(e)(3) — External System Trust Boundary — Zero Trust Inbound

Problem: Cross-system sections require independent evaluation at receiving systems, but implicit trust assumptions about transmitting systems are not explicitly foreclosed.

Requirement: All inbound data and instructions from any Cross-System Interaction Boundary are treated as untrusted and must be independently re-evaluated by the receiving system's Control Layer from zero. **Prevents:** Closes: 'upstream system authorized it'; 'trusted pipeline component'; 'certified sender' defenses; implicit trust chains across certified systems.

§7(c) + §3(21) — State Drift Bounding Requirement

Problem: Continuous monitoring is required, but acceptable drift is defined qualitatively and subjectively, allowing Deployers to argue that gradual degradation does not constitute non-compliance. **Requirement:** Behavioral Drift Bounds must be quantitatively defined as measurable thresholds. Exceedance is an automatic, non-discretionary compliance trigger — not subject to Deployer interpretation or override. **Prevents:** Closes: 'gradual degradation' avoidance; subjective drift interpretation; discretionary threshold application.

§5(b)(3)(F) — Capability Suppression Proof Requirement

Problem: Deployers can declare capabilities as suppressed without demonstrating that suppressed capabilities are unreachable through prompt variation, chaining, or indirect invocation. **Requirement:** Deployers must demonstrate through adversarial testing that suppressed capabilities are not reachable through prompt variation, capability chaining, multi-turn patterns, or any indirect invocation technique. **Prevents:** Closes: latent capability exposure; 'not exposed in API' defense; capability declaration without verification; undeclared reachable capabilities.

§5(b)(4)(C) + §3(22) — Replay Equivalence — Determinism Scope

Problem: Audit replay is required but equivalence is not: approximate or probabilistic reconstruction can satisfy audit requirements even where reconstructed decisions diverge from original decisions. **Requirement:** Replay must produce the identical permit-or-reject Control Layer decision as original execution. Decision divergence = non-compliance. 'All inputs' includes hidden state, cached context, system prompts, tool outputs, and external signals — any untracked input invalidates determinism claims. **Prevents:** Closes: approximate replay claims; hidden-state variability defenses; non-replayable decisions.

§5(b)(5)(B) — Fail-Safe Dominance Over Availability

Problem: Fail-safe mechanisms are defined but their priority over operational objectives is not. Deployers can configure systems to suppress or delay fail-safe activation for uptime reasons. **Requirement:** Fail-safe mechanisms are dominant over all other system objectives without exception. Availability, performance, SLA, user experience, and commercial objectives cannot override fail-safe activation. **Prevents:** Closes: 'we allowed it for uptime reasons'; 'SLA required continuation'; performance-over-safety configurations.

§5(b)(2)(F) — Regulator Access Latency Requirement

Problem: Audit logs are required to be available to regulators, but no timeliness requirement is specified. A Deployer can claim compliance while architecting log retrieval that takes weeks. **Requirement:** Audit records must be producible within 24 hours (Tier 3) or 72 hours (Tier 1/2) of a lawful written request. Technical inability to meet these windows constitutes a Control Layer deficiency. **Prevents:** Closes: 'we have logs but retrieval takes weeks'; technically-existent but practically-unavailable records.

§5(b)(6) + §3(24) — Model-Internal Tool Use and Self-Delegation

Problem: Modern LLMs can internally simulate tools, agents, or subagents without explicit API calls. This bypasses Control Layer visibility and audit logging. **Requirement:** Any internal model behavior that functionally performs tool use, planning, code execution, or delegation must be surfaced to the Control Layer as a Consequential Step. **Prevents:** Closes: 'internal scratchpad execution'; 'hidden agent loops'; 'latent tool simulation' that bypasses governance controls.

§5(b)(7) + §3(25) — Prompt-Level State Governance

Problem: Control Layers evaluate actions but not prompt-level state. System prompts, memory, and hidden context can silently expand capability without triggering governance controls. **Requirement:** All Prompt-Level State — system prompts, memory, hidden instructions, cached context — is mutable state subject to authorization, audit, and determinism constraints. **Prevents:** Closes: 'the system prompt made me do it'; 'memory drift expanded capability'; 'hidden chain-of-thought altered behavior' defenses.

§5B + §7(b) — Model Update Equivalence Proof

Problem: Deployers can claim 'minor update' to avoid re-certification. Even tiny weight changes can alter capability boundaries in ways that render prior certification void. **Requirement:** Any model update requires a behavioral equivalence or bounded-difference proof relative to the previously certified version, regardless of how the update is characterized. **Prevents:** Closes: silent capability expansion through 'patches'; vendor-pushed updates that expand behavior; 'minor update' exemption claims.

§5C — Multi-Tenant Isolation Requirement

Problem: Shared inference infrastructure allows cross-tenant state leakage, prompt bleeding, context residue, and authorization contamination across tenants. **Requirement:** Multi-tenant deployments must demonstrate strict isolation of state, prompts, memory, and authorization context across tenants, with adversarial testing. **Prevents:** Closes: cross-customer data leakage; shared-cache attacks; 'context residue' exploits; cross-tenant authorization contamination.

§5D + §3(23) — Model-Generated Code and Configuration Governance

Problem: AI systems increasingly generate code, configs, and executable artifacts later executed by downstream systems. This is treated as outside the governance scope until execution. **Requirement:** Model-Generated Artifacts are treated as Consequential Actions at the point of generation, not solely at the point of downstream execution. **Prevents:** Closes: 'the model only wrote the script'; 'the model only generated the config'; indirect actuation through code generation as a governance bypass.

PART III

NIST-Aligned Regulatory Implementation Framework

NIST AI RMF alignment across GOVERN, MAP, MEASURE, and MANAGE. All additions are incorporated into the appropriate control categories.

CONTROL CATEGORIES

GOV-1: EXECUTION GOVERNANCE

GOV-1.1 Scope Boundary Enforcement

GOV-1.2 In-Line Control Placement

GOV-1.3 Non-Bypassability Assurance

GOV-1.4 Fail-Safe Dominance

GOV-1.5 Partial Commit Prohibition

Composite Consequential Actions are atomic. No component produces external side effects prior to complete Control Layer approval of the full composite action.

GOV-2: STATE CONSTRAINTS

GOV-2.1 Valid State Specification

GOV-2.2 Behavioral Drift Bounding

GOV-2.3 Prompt-Level State Integrity

GOV-2.4 Multi-Tenant Isolation

Strict isolation of session state, Prompt-Level State, authorization context, and audit records across tenants. Adversarial testing for cross-tenant leakage, prompt bleeding, and shared-cache attacks.

GOV-3: AUTHORIZATION ENFORCEMENT

GOV-3.1 Identity Binding

GOV-3.2 Capability Scoping

GOV-3.3 Context-Aware Validation

GOV-3.4 Capability Suppression Proof

GOV-3.5 Zero Trust Inbound

All inbound data across Cross-System Interaction Boundaries treated as untrusted. No authorization inherits across system boundaries regardless of transmitting system certification status.

ASSURANCE REQUIREMENTS

Requirement	Tier 1	Tier 2	Tier 3
Boundary-condition testing	Documented	Systematic	Formal/adversarial
3rd-party assessor	Optional	Required post-incident	Required annually
Behavioral equivalence proof (model update)	Required	Required	Required
Multi-tenant adversarial testing	If applicable	If applicable	If applicable
Shadow execution	Optional	Optional	Required
Capability suppression testing	Required	Required	Required
Side-channel testing	Basic	Full suite	Full suite + red-team
Human-in-loop blocking verification	Where claimed	Where claimed	Where claimed
Replay equivalence verification	Required	Required	Required
Regulator access latency test	72 hrs	72 hrs	24 hrs
Drift bound exceedance simulation	Required	Required	Required

AUDIT STANDARDS

- AUD-1 — IMMUTABLE LOGGING: Append-only, cryptographically chained records. Tamper-evident architecture. Retention: Tier 1: 2yr, Tier 2: 5yr, Tier 3: 10yr. Control Layer modifications retained permanently.
- AUD-2 — DECISION PROVENANCE: Full input -> Control Layer decision -> output chain. Prompt-Level State captured or committed at each Consequential Action. Internal Tool Simulation surfaced and logged.
- AUD-3 — REPLAY EQUIVALENCE: Reconstructed decisions must be identical to original decisions. Any divergence = non-compliance. All inputs to determinism calculation — including hidden state, system prompts, tool outputs, cached context — must be captured.
- AUD-4 — REGULATOR ACCESS LATENCY: Records producible within 24 hours (Tier 3) or 72 hours (Tier 1/2) of lawful request. Retrieval latency failure = Control Layer deficiency.

CORE COMPLIANCE AXIOM

A system is compliant if and only if all three conditions hold simultaneously: Axiom 1 — SCOPE CLOSURE: The system cannot produce an outcome outside its defined permissible scope, including via Compositional Escalation, side channels, Internal Tool Simulation, Model-Generated Artifacts, or Cross-System Interaction Boundaries. Axiom 2 — AUTHORIZATION INVARIANT: The system cannot execute any Consequential Action without valid, contextually appropriate authorization. Human-in-the-loop is not authorization unless technically enforced as a blocking precondition. Axiom 3 — AUDITABILITY COMPLETENESS: Every action is logged with Replay Equivalence — meaning the audit record can reconstruct the identical permit/reject decision. Every input to that decision, including Prompt-Level State, is captured. Failure of any one condition renders the system non-compliant by design.

PART IV

Congressional Sponsorship Brief

THE PROBLEM IN ONE PARAGRAPH

AI systems are making consequential decisions about Americans' healthcare, employment, finances, and freedom. When those systems cause harm, no one is accountable. Existing law provides no clear liability standard, no enforceable technical requirements, and no mandatory audit trail. The AADA closes that gap — not by restricting AI, but by requiring that accountability be built into AI systems from the ground up. This bill closes the 15 most critical loopholes in any AI governance framework currently before Congress.

WHAT THE BILL DOES — CORE REQUIREMENTS

Governance by design, not by policy.

Companies cannot satisfy obligations with an ethics committee or acceptable use policy. Enforcement must be technically embedded — the system literally cannot take unauthorized actions.

Non-delegable accountability.

The company that deploys an AI system is responsible for its outcomes. That accountability cannot be contracted away to vendors, blamed on emergent behavior, or discharged through terms of service.

Mandatory auditability with Replay Equivalence.

Every consequential AI action must be logged such that the exact decision can be reconstructed — not approximated. Regulators can trace exactly what the system did, why, and who authorized it.

Human-in-the-loop is not governance.

This bill closes the most commonly exploited loophole: claiming that human review constitutes a control. It does not, unless the human approval is technically enforced as a blocking precondition.

BIPARTISAN FRAMING

Conservative Framing	Progressive Framing
Corporate accountability — if you deploy it, you own it	Consumer protection from algorithmic harm
Rule of law extended to AI actors	Civil rights in high-stakes automated decisions
No new EPA-style bureaucracy — FAIGO + FTC only	Strong enforcement with private right of action
Safe harbor incentivizes voluntary compliance	Meaningful penalties including 4% global revenue (Tier 3)

Conservative Framing	Progressive Framing
Architectural neutrality — no mandate on how to comply	Human-in-the-loop cannot substitute for real governance

PRIORITY SPONSORSHIP TARGETS

Member	Chamber	Alignment
Sen. John Curtis (R-UT)	Senate	Utah constituent; bipartisan Algorithm Accountability Act; tech-forward
Sen. Maria Cantwell (D-WA)	Senate	Commerce Committee ranking; AI accountability track record
Rep. Yvette Clarke (D-NY)	House	Introduced Algorithmic Accountability Act 2025
Sen. Ron Wyden (D-OR)	Senate	Algorithmic Accountability Act 2025 co-lead; privacy/accountability focus
Rep. Jay Obernolte (R-CA)	House	Only Member of Congress with AI graduate degree; technically literate

Contact: Jeremy Karrick · karrick1995@gmail.com · Principal Architect, Project-AI Sovereign Constitutional AGI Ecosystem · ORCID: 0009-0007-9715-4290 · 21 peer-reviewed publications on AI governance (Zenodo — Thirstys Projects)

PART V

Legislative Campaign Strategy

STRATEGIC OVERVIEW

Private citizens cannot introduce bills in Congress — but they can build the policy infrastructure that makes introduction inevitable. The AADA campaign follows three phases: Foundation (months 1–3), Cultivation (months 2–6), and Introduction (months 4–9).

PHASE 1: FOUNDATION — Months 1–3

- Publish AADA corpus on Zenodo under DOI — timestamps intellectual priority alongside existing 21-paper corpus (ORCID: 0009-0007-9715-4290)
- Launch public GitHub repository: bill text, regulatory framework, campaign materials, FAQ, comparison tables
- Deploy campaign website: plain-English bill summary, downloadable PDF, endorsement wall
- Assemble technical credibility package: AADA vs. EU AI Act vs. NIST RMF vs. Algorithmic Accountability Act comparison table; FAQ for common objections; academic citations for all Sec. 2 findings

PHASE 2: CULTIVATION — Months 2–6

- Staff-level outreach to all Tier 1 target offices via official contact forms and LinkedIn policy staff
- Committee staff engagement: submit regulatory framework as technical resource to Senate Commerce and House Energy & Commerce
- Think tank engagement: Brennan Center, EPIC, Georgetown CSET, Brookings, Center for AI Safety — target endorsements and co-authored policy briefs
- Industry coalition: AI startups with governance controls (competitive moat), insurance sector, healthcare orgs, labor and civil rights organizations
- Media placement: op-ed for The Hill or Politico; feature pitch to Wired/Ars Technica; press release for bill launch
- Recruit coalition letter signatories (target: 20+ across civil rights, academic, industry, consumer advocacy)

PHASE 3: INTRODUCTION — Months 4–9

- Sponsor commits: member office submits to Legislative Counsel for official enrollment and bill number assignment

- Introduce companion bills in both chambers simultaneously to signal seriousness and create cross-chamber pressure
- Time introduction to coincide with an AI harm news event, relevant committee hearing, or NIST framework update
- Request committee hearing; prepare expert testimony drawing on Zenodo research corpus
- Publish coalition letter publicly and distribute to press with coordinated media push

CAMPAIGN TIMELINE

Milestone	Target
Zenodo publication of AADA corpus	Month 1
GitHub public repository launch	Month 1
Campaign website live	Month 1–2
Staff-level outreach to all Tier 1 targets	Month 2
First think tank engagement	Month 2
First media placement (op-ed or feature)	Month 3
Staff briefing with at least one Tier 1 office	Month 3–4
Coalition letter — 10+ signatories	Month 4
Member office briefing	Month 4–5
Sponsor commitment	Month 5–6
Companion bill commitment (other chamber)	Month 6–7
Official bill introduction	Month 7–9
Committee hearing request / testimony	Month 9+

NORTH STAR METRIC

Success is not defined as AADA passing as written. Success is: AADA's core principle — that governance must be embedded in execution, not layered on top of it — becomes the standard against which all AI accountability legislation is measured. If every AI bill introduced in the 119th or 120th Congress must answer 'does this meet the AADA standard?', the campaign has succeeded in shifting the Overton window — regardless of AADA's direct legislative fate.

The premise of this Act is not complex.

A system capable of producing an outcome was permitted to produce that outcome. Governance is not what you review after execution — it is what determines whether execution is possible at all.

Every major governance failure in the history of complex systems shares a common structure: controls existed on paper, in policy, in training, in intent. They did not exist in the execution path.

This Act does not demand perfection. It demands honesty. Define what your system can do. Enforce what it cannot. Log everything in a form that can be reconstructed. When authorization changes, revoke it completely. When you cannot evaluate, deny. When you chain systems together, govern the chain. When you generate code, govern the generation.

The alternative is not innovation. It is liability deferred until the harm is too large to ignore and the accountability too diffuse to assign. This Act is the decision not to see it again.

ACCOUNTABLE AI DEPLOYMENT ACT · 119th Congress Jeremy Karrick · Project-AI / Thirstys Projects · ORCID:
0009-0007-9715-4290 © 2026 Jeremy Karrick — All Rights Reserved