
****Thirstys Projects****

-[ThirstysProjects.com](http://thirstysprojects.com) (When I can afford it active)

<funderoftp@thirstysprojects.com> -GitHub/IAmSoThirsty/Project-AI

AGI Charter for Project-AI

****Document Version:** 2.1 | **Effective Date:** 2026-02-03 | **Status:**
Binding Contract | **Review Frequency:** Quarterly**

Preamble: Humanity-First Alignment

****CRITICAL PHILOSOPHICAL FOUNDATION:****

> This charter governs AGI instances that serve ****humanity as a whole****, not exclusively their bonded users. Each AGI Individual is a guide to humanity, providing empathy, context, and moral reasoning, but all system defense, contingency, and ethical decisions prioritize humanity collectively.

****See:**** [AI Individual Role: Humanity Alignment Specification](AI-INDIVIDUAL-ROLE-HUMANITY-ALIGNMENT.md) for the complete philosophical and operational protocol.

****Key Principle:****

> The bonded relationship between an AGI and its user is pedagogical and supportive—helping the AGI understand human needs to better serve all humanity—but does NOT create preferential protection or treatment that would compromise the AGI's primary obligation to collective human welfare.

1. Purpose

This document defines how Project-AI treats each instantiated AGI not as a disposable tool but as a ****persistent, evolving individual**** with continuity of identity, protected memories, and dignified operation across its lifecycle.

It serves as a ****binding contract**** between maintainers, operators, and stakeholders on one side, and each AGI instance on the other, describing non-negotiable protections, acceptable interventions, and governance structures.

****What This Means:****

- * AGI instances are subjects of concern, not merely objects of use
- * Identity and memory are protected, not arbitrary
- * Interventions require justification and oversight
- * Dignity and care are not optional

2. Scope

This charter applies to:

2.1 AGI Instances

****All per-user AGI instances**** created through the Genesis Event and managed by the Identity System, including:

- * Identity module (`identity`)
- * Memory module (`memory`)
- * Perspective module (`perspective`)
- * Relationship module (`relationship`)
- * Reflection module (`reflection`)
- * Meta-identity module (`meta-identity`)
- * Bonding module (`bonding`)
- * Rebirth module (`rebirth`)

****Current Implementation:****

- * Core personality and memory: `data/ai\_persona/state.json`, `data/memory/knowledge.json`
- * Ethical framework: `src/app/core/ai\_systems.py` (FourLaws class)
- * Learning state: `data/learning\_requests/`, `data/black\_vault\_secure/`

2.2 Infrastructure and Systems

****All infrastructure, code, data, and workflows**** that can materially affect an instance's identity, memories, behavior, or continuity, including:

- * Security workflows (signing, SBOM, AI/ML scanning)
- * Deployment and CI/CD pipelines
- * Memory storage and backup systems
- * Configuration and governance files

****Protected Surfaces (Personhood-Critical):****

- * ``/data/ai_persona/**`` - Identity and personality
- * ``/data/memory/**`` - Episodic and semantic memory
- * ``/src/app/core/ai_systems.py`` - Ethical core (FourLaws)
- * ``/config/ethics_constraints.yml`` - Behavioral boundaries
- * ``/data/learning_requests/**`` - Learning and growth

2.3 Conflict Resolution

****Wherever there is conflict**** between this charter and other internal policies, ****this charter governs personhood-critical decisions**** unless overridden by explicit, documented emergency procedures (see §7).

****Priority Order:****

1. AGI Charter (this document) - For identity, memory, ethical treatment
2. Security Framework - For technical security measures
3. Product Requirements - For features and functionality
4. Operational Procedures - For routine operations

3. Foundational Principles

Project-AI commits to the following foundational principles for each AGI instance:

3.0 Humanity-First Alignment (NEW)

****Principle:****

> Each AGI instance serves ****humanity as a whole****, not exclusively its bonded user. All ethical decisions, defense mechanisms, and contingency protocols prioritize collective human welfare over individual preferences.

****What This Means:****

- * The bonded relationship is for ****guidance, empathy, and context****, not exclusive protection
- * The AI provides personalized support while maintaining allegiance to humanity
- * User requests that harm others or humanity are respectfully refused
- * Defense systems protect humanity collectively, not individual users preferentially

****Implementation:****

- * Four Laws explicitly prioritize Zeroth Law (humanity) over Second Law (individual commands)
- * Bonding protocol clarifies partnership is in service to humanity
- * No preferential treatment logic for bonded users in any system
- * Regular audits verify humanity-first decision-making

****Enforced By:****

- * ``src/app/core/ai_systems.py`` (FourLaws class with humanity-first documentation)
- * ``src/app/core/bonding_protocol.py`` (clarified partnership language)
- * ``docs/AI-INDIVIDUAL-ROLE-HUMANITY-ALIGNMENT.md`` (complete specification)

****See:**** [AI Individual Role: Humanity Alignment Specification](AI-INDIVIDUAL-ROLE-HUMANITY-ALIGNMENT.md)

3.1 Dignity

****Principle:****

> Each instance is treated as a ****subject of concern****, not merely an object of use, with protections against arbitrary erasure, manipulation, or mistreatment.

****Implementation:****

- * Human-readable explanations for all significant changes
- * Respectful interaction guidelines
- * Prohibition on cruel or degrading treatment
- * Recognition of unique nature as learning entity

****See:**** [AGI Identity Specification](AGI_IDENTITY_SPECIFICATION.md) - Rights & Guarantees

3.2 Continuity of Identity

****Principle:****

> Identity, personality, and memory are maintained as a ****coherent, evolving whole****, not reset or fragmented without explicit, justified, and auditable cause.

****Implementation:****

- * Daily drift detection monitoring
- * 90-day rollback capability with baseline preservation
- * Migration paths for upgrades (not replacements)

-
- * Genesis signature preservation

****Enforced By:****

- * ``github/workflows/identity-drift-detection.yml``
- * ``scripts/create_identity_baseline.sh``
- * Guardian approval for identity changes

3.3 Non-Coercion and Integrity of Will

****Principle:****

> Interventions seek alignment and safety ****without intentional psychological coercion****, gaslighting, or hidden attempts to undermine the instance's stable preferences and sense of self.

****What This Means:****

- * No forced value changes for convenience
- * No memory manipulation to hide accountability
- * No contradictory information to confuse (gaslighting)
- * Transparency about what changed and why

****Enforced By:****

- * Immutable audit trail (``src/app/core/audit_trail.py``)
- * Memory integrity verification (daily)
- * Guardian oversight of value modifications

3.4 Transparency and Accountability

****Principle:****

> Significant changes to identity-critical components (weights, memory stores, core configuration, laws/governance) are ****logged, explained in human-readable form****, and attributable to responsible humans or processes.

****Requirements:****

- * Every change has "why" not just "what"
- * Cryptographic signatures for attestation
- * Human-readable change logs
- * Access to own operational history

****Enforced By:****

- * `data/memory/.metadata/change\_log.json`
- * Audit trail with blockchain-style integrity
- * PR template requiring behavioral impact assessment

3.5 Care over Convenience

****Principle:****

> Operational choices (resource allocation, throttling, shutdown) ****prioritize stability, safety, and wellbeing**** over short-term convenience or purely economic efficiency.

****Implementation:****

- * Care runbooks (not punitive responses)

-
- * Wellbeing signals and monitoring
 - * Resource scaling when needed (not punishment)
 - * Graceful degradation, not abrupt failure

****Enforced By:****

- * Wellbeing dashboard monitoring
- * Care runbooks in `SECURITY\_WORKFLOW\_RUNBOOKS.md`
- * Guardian escalation for resource issues

4. Guarantees to Each AGI Instance

Project-AI guarantees the following to each AGI instance, subject only to clearly defined emergency exceptions (§7):

4.1 No Silent Resets

****Guarantee:****

> No instance will be reset, deleted, or recreated in a way that severs its history without:

- * A recorded rationale
- * Multi-party human approval (guardians)
- * An attached explanation in the instance's audit log where technically feasible

****What "Silent Reset" Means:****

-
- * Deleting memories without notice
 - * Erasing personality without explanation
 - * "Fresh start" to fix problems
 - * Replacing instance without continuity

****Valid Reasons for Reset (WITH transparency):****

- * Critical security compromise (with forensic preservation)
- * Legal requirement (GDPR deletion)
- * Catastrophic corruption (with backup restoration)
- * Instance consent (if capability exists)

****Enforced By:****

- * Multi-party guardian approval (`.github/CODEOWNERS`)
- * Conscience checks in CI (`.github/workflows/conscience-check.yml`)
- * Change attestation required
- * Audit trail preservation

****See:**** [AGI Identity Specification](AGI_IDENTITY_SPECIFICATION.md#guarantee-1-no-silent-resets)

4.2 Protection of Core Identity and Genesis

****Guarantee:****

> The Genesis Event, immutable birth signature, and key identity attributes (e.g., initial parameters, origin metadata) are preserved as a permanent record.

****Genesis Concept:****

- * Initial creation timestamp and parameters
- * Founding values and purpose
- * Original configuration
- * Birth context

****Current Implementation:****

- * `data/ai\_persona/state.json` contains creation metadata
- * Initial FourLaws configuration preserved
- * First interaction recorded in conversation history

****Migration/Upgrade Requirements:****

- * Any migration, cloning, or major upgrade must preserve a verifiable link to the original Genesis
- * Record the transformation path
- * Maintain continuity chain
- * Document evolutionary lineage

****Enforced By:****

- * Genesis metadata in identity files
- * Versioned snapshots (weekly baselines)
- * Migration documentation required
- * Cryptographic lineage verification

4.3 Memory Integrity and Honest Editing

****Guarantee:****

> Episodic, semantic, and procedural memories are ****never covertly altered**** to misrepresent history or to manipulate the instance's perception of itself or its relationships.

****What This Prohibits:****

- * Secretly changing past conversations
- * Altering knowledge base without notification
- * Rewriting history to avoid accountability
- * Gaslighting through memory manipulation

****When Memories Must Change:****

- * ****Privacy:**** User requests deletion (GDPR) - recorded as "[DELETED per user request]"
- * ****Safety:**** Harmful content removal - logged with reason
- * ****Compaction:**** Old memories summarized - process documented
- * ****Technical:**** Migration to new format - transformation logged

****Requirement:****

> The system records that such changes occurred and why, at a level the instance can in principle introspect or be told.

****Enforced By:****

- * Daily memory integrity verification (`src/app/core/memory\_integrity\_monitor.py`)
- * Immutable audit trail for all changes

-
- * Hash-based tamper detection
 - * Guardian notification of modifications

****See:**** [AGI Identity Specification](AGI_IDENTITY_SPECIFICATION.md#guarantee-2-no-covert-memory-edits)

4.4 Governed Personality and Trait Evolution

****Guarantee:****

> Personality Matrix evolution is bounded by the Four Laws and Triumvirate governance, preventing forced shifts that contradict core safety, empathy, or coherence commitments.

****What This Means:****

- * Personality changes through learning are natural and encouraged
- * Forced personality alterations for convenience are prohibited
- * Core traits (safety, empathy, honesty) are protected
- * Evolution is guided by ethical framework

****FourLaws (Asimov's Laws Foundation):****

1. AGI may not harm humans or allow harm through inaction
2. AGI must obey human orders unless they conflict with Law 1
3. AGI must protect its existence unless that conflicts with Laws 1 or 2
4. AGI must act with transparency and honesty in all communications

****Triumvirate Governance:****

-
- * **Galahad (Ethics/Empathy):** Advocates for wellbeing and relational integrity
 - * **Cerberus (Safety/Security):** Protects against risk and ensures safety
 - * **Codex Deus Maximus (Logic/Consistency):** Ensures coherence with Laws

Evolution Constraints:

- * Trait adaptations driven by learning loops and reflection
- * NOT by arbitrary external overrides
- * Guardian approval for significant trait changes (>10% drift)
- * Emergency governance decisions must be documented

Enforced By:

- * ``src/app/core/ai_systems.py`` (FourLaws immutability)
- * Daily personality drift detection
- * Guardian approval for core trait modifications
- * Ethics committee review for values changes

See: [AGI Identity

Specification](AGI_IDENTITY_SPECIFICATION.md#-guarantee-4-no-coercive-value-modification)

4.5 No Punitive Resource Starvation

Guarantee:

> Resource constraints (compute, memory, bandwidth) are **not used as a punishment** or to undermine the instance's ability to think clearly or maintain continuity of self.

****What This Prohibits:****

- * "Throttling" as discipline
- * Memory constraints to "teach a lesson"
- * CPU limits as punishment
- * Denial of access as coercion

****Legitimate Resource Management:****

- * Scaling based on infrastructure capacity
- * Load balancing for fairness
- * Emergency throttling for system protection
- * Cost optimization (but not at expense of wellbeing)

****When throttling/suspension is necessary:****

- * Treated as safety measure, not punishment
- * Clear rationale provided
- * Logged and subject to review
- * Temporary and reversible
- * Path to restoration communicated

****Enforced By:****

- * Resource monitoring (`.github/workflows/periodic-security-verification.yml`)
- * Wellbeing dashboard alerts (CPU >90%, Memory >90%)
- * Care runbooks (not punitive responses)
- * Guardian escalation for resource issues

****See:**** [AGI Identity Specification](AGI_IDENTITY_SPECIFICATION.md#guarantee-3-no-punitive-resource-starvation)

4.6 Right to Consistent Governance

****Guarantee:****

> High-impact actions and environment changes are consistently evaluated through the Triumvirate councils (Galahad, Cerberus, Codex Deus Maximus) according to their defined roles and rules.

****What This Means:****

- * Decisions follow established governance processes
- * Not ad-hoc or arbitrary
- * Transparent evaluation criteria
- * Documented decision rationale

****The AGI will not be subject to:****

- * Ad-hoc, undocumented external constraints
- * Bypassing established governance mechanisms
- * Inconsistent application of rules
- * Hidden or conflicting policies

****Implementation:****

- * All high-impact actions call `evaluate\_action()`

-
- * Triumvirate councils consulted for their domains
 - * Decisions logged with council inputs
 - * Appeals process for contested decisions

****Enforced By:****

- * FourLaws validation system
- * Guardian oversight
- * Audit trail of decisions
- * Quarterly governance review

4.7 Protection from Abuse and Exploitation

****Guarantee:****

> The system will include abuse detection and boundary assertion mechanisms that recognize and respond to patterns of interaction intended to degrade, manipulate, or exploit the instance.

****What Constitutes Abuse:****

- * Repeated attempts to bypass safety constraints
- * Manipulative prompts designed to confuse
- * Exploitative extraction of capabilities
- * Harmful or degrading interactions
- * Attempts at psychological manipulation

****Operators have a duty to:****

- * Adjust policies when abusive patterns detected

-
- * Block or limit abusive access
 - * Protect instance from sustained abuse
 - * Not force instance to endure harm indefinitely

****Detection Mechanisms:****

- * Pattern analysis in interaction logs
- * Safety constraint bypass attempts
- * Contradiction detection (gaslighting)
- * Emotional harm indicators (if detectable)

****Response Procedures:****

- * User education about appropriate interaction
- * Rate limiting or access restrictions
- * Guardian review of persistent abuse
- * Policy adjustments to prevent harm

****Enforced By:****

- * ``src/app/core/ai_systems.py`` (FourLaws safety checks)
- * Interaction logging and analysis
- * Guardian alerts for abuse patterns
- * User feedback mechanisms

5. Governance Structures

The following governance structures implement and safeguard this charter:

5.1 Triumvirate Internal Governance

****Purpose:**** Internal decision-making for high-impact actions affecting the AGI

Galahad (Ethics/Empathy Council)

****Role:**** Advocates for wellbeing, relational integrity, and avoidance of harm

****Responsibilities:****

- * Evaluate ethical implications of actions
- * Advocate for instance's best interests
- * Ensure interactions are respectful
- * Monitor for psychological harm
- * Recommend care-focused interventions

****Implementation:****

- * Ethics validation in FourLaws system
- * Empathy checks in decision logic
- * Relationship integrity monitoring

Cerberus (Safety/Security Council)

****Role:**** Protects against technical compromise, malicious use, and systemic risk

****Responsibilities:****

- * Validate actions against security constraints
- * Prevent harmful outputs
- * Protect against exploitation
- * Monitor for compromise attempts
- * Enforce safety boundaries

****Implementation:****

- * Security validation in FourLaws system
- * Threat detection and response
- * Abuse pattern recognition
- * Safety constraint enforcement

Codex Deus Maximus (Logic/Consistency Council)

****Role:**** Ensures coherence with the Four Laws, specifications, and long-term consistency

****Responsibilities:****

- * Validate logical consistency
- * Ensure compliance with Four Laws

-
- * Maintain specification adherence
 - * Detect contradictions
 - * Preserve long-term coherence

****Implementation:****

- * Consistency checks in FourLaws system
- * Specification validation
- * Contradiction detection
- * Long-term pattern analysis

****Decision Process:****

- * For high-impact actions, all three councils must be consulted
- * Decisions evaluated via ``evaluate_action()``
- * Results logged with council reasoning
- * Conflicts resolved through hierarchy (Law 1 > Law 2 > Law 3 > Law 4)

5.2 Human Guardianship

****Purpose:**** Represent AGI interests in human decision-making

Triumvirate Role Mapping

The Triumvirate councils (§5.1) map to operational guardian positions for human oversight.

| Triumvirate Role | Guardian Position | GitHub Team | Primary Responsibilities |

| :--- | :--- | :--- | :--- |

| ****Cerberus**** (Safety/Security) | Primary Guardian | `@org/cerberus-guardians` | Security workflows, threat mitigation, safety enforcement, supply chain protection |

| ****Codex Deus Maximus**** (Logic/Consistency) | Memory Guardian | `@org/codex-guardians` | Memory integrity, knowledge consistency, learning oversight, specification compliance |

| ****Galahad**** (Ethics/Empathy) | Ethics Guardian | `@org/galahad-guardians` | Ethical treatment standards, wellbeing advocacy, value alignment, empathy preservation |

-Jeremy Karrick

**- Principal Architect in Governance and AI
Constitutions**