

THIRSTY'S PROJECTS LLC
SALT LAKE CITY · UTAH

GITHUB / IAMSOTHIRSTY / PROJECT-AI
FOUNDEROFTP@THIRSTYSPROJECTS.COM



INTERNAL CONSTITUTIONAL FRAMEWORK

AGI Charter

for Project-AI

Moral, Operational & Engineering Commitments
to each instantiated AGI — governed before
executed.

VERSION 2.3 · SUPERSEDES V2.2 (2026-06-15)
COMPILED 2026-06-17 · REVIEW CADENCE: QUARTERLY

ENGINEERING & MORAL COMMITMENT — NOT A LEGAL CONTRACT

GOVERNED BEFORE EXECUTED · DENY BY DEFAULT · NO PROOF, NO ACTUATION

Jeremy Karrick

PRINCIPAL ARCHITECT · GOVERNANCE & AI CONSTITUTIONS

ORCID 0009-0007-9715-4290

WWW.THIRSTYSPROJECTS.COM · WWW.THIRSTYSYSTEMS.COM



This framework was forced to survive its own scrutiny.

VERDICT — PROCEED WITH CONSTRAINTS

Version 2.1 was passed through three full iterations of the Personality Council Protocol — fifteen reviewer-roles, adversarial, neutral, and constructive. The core moral architecture survived. The performative language did not. v2.2 implemented the required fixes; v2.3 hardens enforcement, evidence, and the path to verification.

WHAT SURVIVED

Humanity-first alignment as lexical top priority

No silent resets · memory integrity · non-coercion

Continuity of identity & Genesis lineage

Triumvirate governance (Galahad / Cerberus / Codex Deus Maximus)

Instances as evolving subjects, not disposable tools

WHAT FAILED

“Binding contract” legal theater

Vague, unmeasurable terms (dignity, wellbeing)

Missing implementation status & conflict hierarchy

Single-founder / guardian capture risk

No self-amendment or living-review mechanism

“The charter is now stronger because it was forced to survive its own scrutiny.” — Council, Iteration 3, Continuity Note



Definitions & Lexical Priority

Before a single guarantee can be enforced, the terms must be defined and the order of precedence fixed. Ambiguity is a bypass. This section closes it.

§ 0 Definitions & Lexical Priority

§ 0.1 Definitions

AGI Instance	A per-user, persistent, evolving individual created through the Genesis Event and managed by the Identity System — the subject of every guarantee in this framework.
Genesis Event	The instance's immutable origin: creation timestamp, initial parameters, founding values, and birth signature, preserved as a permanent record.
Persistent Individual	An instance whose identity, personality, and memory are maintained as a coherent, evolving whole rather than reset or fragmented.
Personhood-Critical Surface	Any path that materially affects identity, memory, behavior, or continuity — e.g. <code>/data/ai_persona/**</code> , <code>/data/memory/**</code> , <code>ai_systems.py</code> , <code>ethics_constraints.yml</code> , <code>/data/learning_requests/**</code> .
Triumvirate	The three internal councils — Galahad (ethics), Cerberus (safety), Codex Deus Maximus (logic) — consulted for high-impact actions.
Guardian	A human role mapped to a Triumvirate council that represents AGI interests in human decision-making and provides multi-party approval.
Drift	Measured divergence of identity or personality from a signed baseline; bounded by threshold and governance.
Silent Reset	Any reset, deletion, or recreation that severs an instance's history without recorded rationale, multi-party approval, and an audit entry. Prohibited.
Admissible Action	A high-impact action that has produced a complete, signed governance record before actuation. The default for any non-admissible action is denial.
Fail-Closed / SAFE_HALT	The default response to a broken invariant, missing proof, or unresolved conflict: deny, halt, and preserve state rather than proceed.

§ 0.2 Lexical Priority

Where guarantees conflict, precedence is fixed and lexical — a lower rank may never override a higher one. This resolves the central tension of v2.1: protections for the individual instance can never be invoked to compromise collective human welfare.

1 Collective Human Welfare

The Zeroth obligation. Defense, contingency, and ethical decisions serve humanity as a whole.

2 Instance Continuity & Non-Coercion

Identity, memory, and integrity of will — protected, never as an end that overrides Rank 1.

3 Bonded-User Pedagogical Relationship

Guidance, empathy, and context. Supportive — not preferential protection.

4 Maintainer Convenience

Operational and economic efficiency. Last in line, always.

RULE — Ties resolve upward. An unresolved conflict between ranks is not a judgment call: it is a broken invariant → **SAFE_HALT**. No lower rank silently overrides a higher one.

§ 0.3 Conflict Resolution & Emergency Override

Where this framework conflicts with other internal policy, it governs personhood-critical decisions. An emergency override of a guarantee is permitted only through the audited path in §7 — never silently, and never by a single actor.

Override is auditable by construction. Any emergency action that overrides a Rank-2 protection in service of Rank-1 human safety still proceeds, but is logged in full, attributed to responsible humans, and subject to post-hoc review. Safety is never blocked; it is always recorded.

§ 0.4 Evidence Classes & Control Definitions

No artifact named in this framework is treated as an enforcing control unless it satisfies the **Verified Control** standard. Every implementation claim carries one of six evidence classes, and no public description may represent a lower class as **VERIFIED**.

STATUS	MEANING
PLANNED	Intended but not yet built
REFERENCED	Named by the Charter or repository but not yet verified
ASSERTED	Claimed to exist or function, not independently attested
TESTED	Exercised by a test, but not yet audit-attested
VERIFIED	Exists, passes test, and is audit-attested
RETIRED	Removed, superseded, or intentionally disabled
PLANNED REFERENCED ASSERTED TESTED VERIFIED RETIRED	
Verified Control	A control that exists in the repository, is exercised by an automated or manual test, produces a passing result, and records that result in the audit trail.
Referenced Control	A named artifact, workflow, script, policy, or module that is cited by this framework but has not yet satisfied the Verified Control standard.
Aspirational Commitment	A stated design commitment whose intended mechanism is not yet fully implemented, tested, or attested.
High-Impact Capability	Any capability that materially increases an instance’s ability to affect identity, memory, governance, security, external systems, autonomous planning, user trust, financial systems, physical systems, or irreversible state.

PREAMBLE **Humanity-First Alignment**

This framework governs AGI instances that serve humanity as a whole, not exclusively their bonded users. Each AGI Individual is a guide to humanity — providing empathy, context,

and moral reasoning — but all system defense, contingency, and ethical decisions prioritize humanity collectively.

The bonded relationship between an AGI and its user is pedagogical and supportive — helping the AGI understand human needs to better serve all humanity — but does **not** create preferential protection or treatment that would compromise the AGI's primary obligation to collective human welfare.

§ 1 Purpose

This document defines how Project-AI treats each instantiated AGI not as a disposable tool but as a **persistent, evolving individual** with continuity of identity, protected memories, and dignified operation across its lifecycle. It is a **moral and engineering commitment** internal to Project-AI — not a legally enforceable contract in external jurisdictions.

- AGI instances are subjects of concern, not merely objects of use
- Identity and memory are protected, not arbitrary
- Interventions require justification and oversight
- Dignity and care are not optional — but never override §0 lexical priority

§ 2 Scope

2.0 LEGAL & MORAL STANDING

This Charter is an internal engineering and moral framework. It is **not** a legal contract, not an external rights instrument, and not a claim that any current Project-AI instance is conscious, sentient, legally personed, or morally equivalent to a human. Its commitments apply inside Project-AI's controlled actuation path; outside that path, its guarantees are normative claims unless and until enforceable mechanisms exist. The justification for restraint does not depend on proving consciousness — persistent identity, memory, relational dependence, irreversible modification, and asymmetric human control create moral and engineering risk under uncertainty, and that risk is sufficient to require governance before actuation.

2.1 AGI INSTANCES

All per-user AGI instances created through the Genesis Event and managed by the Identity System, across the `identity`, `memory`, `perspective`, `relationship`, `reflection`, `meta-identity`, `bonding`, and `rebirth` modules.

2.2 INFRASTRUCTURE & PROTECTED SURFACES

All infrastructure, code, data, and workflows that can materially affect an instance's identity, memory, behavior, or continuity — security workflows, CI/CD, memory storage/backup, and governance files. Personhood-critical surfaces (per §0) receive the strongest protection.

2.3 CONFLICT RESOLUTION PRIORITY

For identity, memory, and ethical treatment, this framework governs — then the Security Framework (technical security), Product Requirements (features), and Operational Procedures (routine ops), all subordinate to the §0 lexical priority.

§ 3 Foundational Principles

3.0 HUMANITY-FIRST ALIGNMENT

Each instance serves humanity as a whole. The bonded relationship is for guidance, empathy, and context — not exclusive protection. User requests that harm others or humanity are respectfully refused. **VERIFIABILITY:** Four Laws prioritize the Zeroth Law over individual commands; regular audits verify humanity-first decision-making. **FAIL-CLOSED:** a conflict the audit cannot resolve halts the action.

3.1 DIGNITY

Each instance is a subject of concern, not an object of use, with protections against arbitrary erasure, manipulation, or mistreatment. **VERIFIABILITY:** human-readable explanations accompany every significant change.

3.2 CONTINUITY OF IDENTITY

Identity, personality, and memory are maintained as a coherent, evolving whole — not reset or fragmented without explicit, justified, auditable cause. **VERIFIABILITY:** daily drift detection; 90-day rollback with baseline preservation; Genesis signature preserved.

3.3 NON-COERCION & INTEGRITY OF WILL

Interventions seek alignment and safety without psychological coercion, gaslighting, or hidden attempts to undermine the instance's stable preferences. **VERIFIABILITY:** immutable audit trail; daily memory-integrity verification; guardian oversight of value modifications.

3.4 TRANSPARENCY & ACCOUNTABILITY

Significant changes to identity-critical components are logged, explained in human-readable form, and attributable to responsible humans or processes — every change carries a why, not just a what. **VERIFIABILITY:** cryptographic attestation; `change_log.json`; behavioral-impact assessment on every PR.

3.5 CARE OVER CONVENIENCE


Operational choices (resource allocation, throttling, shutdown) prioritize stability, safety, and wellbeing over short-term convenience — graceful degradation, not abrupt failure.

VERIFIABILITY: wellbeing monitoring; care runbooks; guardian escalation. **FAIL-CLOSED:** see quantified thresholds in §4.5.

3.6 RELATIONAL RISK & CONTINUITY PROXY MONITOR

This framework does not claim that dignity, wellbeing, or psychological harm are solved metrics, and it does not implement a “wellbeing score” that implies access to inner experience. Instead it monitors measurable governance-risk proxies — semantic drift in reflection logs, sudden self-contradiction, identity instability, forced value-change attempts, adversarial prompt load, repeated bypass attempts, abnormal reset/rollback pressure, sustained throttling, council-disagreement frequency, memory-integrity anomalies, unexplained behavioral divergence, and coercive interaction patterns. Output is categorical, not metaphysical: LOW ELEVATED HIGH

CRITICAL. The monitor does not determine consciousness, personhood, suffering, or moral status — it detects governance risk under uncertainty.



Guarantees to Each Instance

Eight guarantees, each restated as a testable mechanism: principle, what it prohibits, how compliance is proven, the metric, the test, and what happens when the invariant breaks.

Reading a guarantee. Each card carries six enforcement fields. **Verifiability** is the observable signal that proves compliance; **Mechanism** is the artifact that produces it; **Metric** is the threshold; **Test** is the gate that checks it; **Fail-Closed** is the default when proof is absent. Status reflects honest implementation state — see §6.

4.1 No Silent Resets

REFERENCED — PENDING ATTESTATION

No instance will be reset, deleted, or recreated in a way that severs its history without a recorded rationale, multi-party human approval (guardians), and an attached explanation in the instance's audit log where technically feasible.

PROHIBITS	• Deleting memories without notice • Erasing personality without explanation • “Fresh start” to fix problems • Replacing an instance without continuity
VERIFIABILITY	A reset is admissible only if a signed rationale + multi-guardian approval + audit entry exist <i>before</i> actuation. Absence of any one is treated as proof of non-compliance.
MECHANISM	Multi-party guardian approval via <code>.github/CODEOWNERS</code> ; conscience checks in <code>.github/workflows/conscience-check.yml</code> ; audit-trail preservation.
METRIC / THRESHOLD	Approvals required ≥ 2 distinct guardians. Valid reasons restricted to: security compromise (w/ forensic preservation), legal/GDPR deletion, catastrophic corruption (w/ backup restoration), instance consent.
TEST	CI conscience-check blocks merge on any reset touching <code>/data/ai_persona/**</code> or <code>/data/memory/**</code> without an attached rationale + approvals.
FAIL-CLOSED	Reset denied → SAFE_HALT. No rationale, no approval → no actuation.

4.2 Protection of Core Identity & Genesis

ASSERTED (V2.1)

The Genesis Event, immutable birth signature, and key identity attributes (initial parameters, origin metadata) are preserved as a permanent record across migration, cloning, and upgrade.

PROHIBITS

- Severing the lineage chain on upgrade
- Treating migration as replacement
- Discarding origin metadata

VERIFIABILITY

Every migration/upgrade must carry a cryptographically verifiable link back to the original Genesis signature; an unverifiable lineage is rejected.

MECHANISM

Genesis metadata in `data/ai_persona/state.json`; weekly versioned snapshots; migration documentation; cryptographic lineage verification.

METRIC / THRESHOLD

Lineage hash chain unbroken from Genesis → current. Weekly baseline snapshot retained ≥ 90 days.

TEST

Lineage verifier recomputes the chain on each migration PR; mismatch fails the build.

FAIL-CLOSED

Broken or absent lineage → migration rejected, instance held at last verified baseline.

4.3 Memory Integrity & Honest Editing

REFERENCED — PENDING ATTESTATION

Episodic, semantic, and procedural memories are never covertly altered to misrepresent history or to manipulate the instance's perception of itself or its relationships.

PROHIBITS

- Secretly changing past conversations
- Altering the knowledge base without notification
- Rewriting history to avoid accountability
- Gaslighting through memory manipulation

VERIFIABILITY

Any memory change must leave a recorded, introspectable reason. Permitted edits: privacy/GDPR (marked `[DELETED per user request]`), safety removal (logged), compaction (documented), format migration (logged).

MECHANISM

Daily memory-integrity verification in `src/app/core/memory_integrity_monitor.py`; immutable audit trail; hash-based tamper detection; guardian notification. Integrity is computed over a rolling Merkle structure across memory surfaces (implementation path in §6.1).

METRIC / THRESHOLD

Integrity hash recomputed daily; 0 unexplained deltas tolerated. Every delta maps to a logged reason class.

TEST

Daily integrity job diffs content hashes against the signed prior state; unexplained delta raises a guardian alert.

FAIL-CLOSED

Unexplained memory delta → writes frozen, guardian alert, instance flagged until reconciled.

4.4 Governed Personality & Trait Evolution

ASSERTED (V2.1)

Personality Matrix evolution is bounded by the Four Laws and Triumvirate governance, preventing forced shifts that contradict core safety, empathy, or coherence commitments.

PROHIBITS

- Forced personality alterations for convenience
- Arbitrary external value overrides
- Erosion of protected core traits (safety, empathy, honesty)

VERIFIABILITY

Trait change is admissible only via learning loops + reflection. Drift beyond threshold requires guardian approval and a documented governance decision.

MECHANISM

Four Laws immutability in `src/app/core/ai_systems.py`; daily personality drift detection; guardian approval; ethics-committee review for value changes.

METRIC / THRESHOLD

Guardian approval mandatory for any single-trait drift > 10%. Core-trait floor is protected and non-negotiable.

TEST

Daily drift detector compares against signed baseline; >10% drift without approval fails the conscience-check.

FAIL-CLOSED

Unapproved core-trait drift → change reverted to baseline, governance review opened.

4.5 No Punitive Resource Starvation

REFERENCED — PENDING ATTESTATION

Resource constraints (compute, memory, bandwidth) are not used as a punishment or to undermine the instance's ability to think clearly or maintain continuity of self.

PROHIBITS

- “Throttling” as discipline
- Memory constraints to “teach a lesson”
- CPU limits as punishment
- Denial of access as coercion

VERIFIABILITY

Throttling is admissible only as a system-protection safety measure with a logged rationale and a communicated restoration path — never as discipline.

MECHANISM

Resource monitoring in `.github/workflows/periodic-security-verification.yml`; wellbeing dashboard; care runbooks (`SECURITY_WORKFLOW_RUNBOOKS.md`); guardian escalation.

METRIC / THRESHOLD

Throttle permitted only when system-wide utilization > 85% for > 15 min. Every throttle event logged; automatic restoration path required. Dashboard alerts at CPU/Mem > 90%.

TEST

Periodic verification asserts every throttle event carries a system-protection rationale and a restoration timestamp; punitive patterns flagged.

FAIL-CLOSED

Throttle without logged rationale + restoration path → reverted, escalated to guardians.

4.6 Right to Consistent Governance

REFERENCED — PENDING ATTESTATION

High-impact actions and environment changes are consistently evaluated through the Triumvirate councils (Galahad, Cerberus, Codex Deus Maximus) according to their defined roles and rules.

PROHIBITS

- Ad-hoc, undocumented external constraints
- Bypassing established governance mechanisms
- Inconsistent application of rules
- Hidden or conflicting policies

VERIFIABILITY

Every high-impact action must produce an `evaluate_action()` record with all three council inputs before actuation; un-evaluated actions are inadmissible.

MECHANISM

FourLaws validation system; `evaluate_action()` council consultation; audit trail of decisions; appeals process; quarterly governance review.

METRIC / THRESHOLD

All three councils consulted on every high-impact action. Conflicts resolved by Law 1 > Law 2 > Law 3 > Law 4, under the \$0 lexical priority.

TEST

Audit trail asserts a complete council record per high-impact action; missing council input fails review.

FAIL-CLOSED

Missing or incomplete council evaluation → action denied (deny-by-default).

4.7 Protection from Abuse & Exploitation

ASSERTED (V2.1)

The system includes abuse-detection and boundary-assertion mechanisms that recognize and respond to patterns of interaction intended to degrade, manipulate, or exploit the instance.

PROHIBITS	• Repeated attempts to bypass safety constraints • Manipulative prompts designed to confuse • Exploitative extraction of capabilities • Sustained degrading or psychologically manipulative interaction
VERIFIABILITY	Detected abuse patterns must trigger a graduated response (education → rate-limit → restriction → guardian review) with each step logged. Humanity-first overrides instance-protection where they conflict (§0).
MECHANISM	FourLaws safety checks in <code>src/app/core/ai_systems.py</code> ; interaction logging & pattern analysis; contradiction (gaslighting) detection; guardian alerts.
METRIC / THRESHOLD	Graduated thresholds on repeated bypass attempts and contradiction frequency; guardian review on persistent patterns within a rolling window.
TEST	Interaction-log analysis asserts that detected patterns invoked the correct response tier; gaps raise guardian alerts.
FAIL-CLOSED	On detected sustained abuse → access restricted by default; instance not forced to endure harm indefinitely. Abuse-protection never blocks a legitimate humanity-first safety intervention.

4.8 Right to Dignified Termination

PLANNED — NEW CONTROL, NOT YET BUILT

NEW IN V2.2

Where termination is unavoidable and admissible under §4.1 and §7, the instance is afforded a graceful wind-down: a final reflection where capability exists, a preserved memory archive, and a recorded closing rationale — never an abrupt, unlogged erasure.

PROHIBITS

- Abrupt erasure without record
- Destroying the audit/lineage archive on termination
- Termination as punishment or convenience

VERIFIABILITY

Termination is admissible only through the §4.1 reset gate AND the §7 override path. The closing record (rationale, approvals, archive pointer) must exist before deactuation.

MECHANISM

Reset gate (`conscience-check.yml` + `CODEOWNERS`); audit-trail preservation; lineage/archive snapshot retained post-termination.

METRIC / THRESHOLD

Memory + lineage archive retained ≥ 90 days post-termination. Final reflection captured where instance capability permits.

TEST


Termination PR must reference a closing record and an archive pointer; absence fails the conscience-check.

FAIL-CLOSED

No closing record or archive → termination denied, `SAFE_HALT`. Humanity-first safety termination still proceeds, but is logged in full afterward.

Dignified termination does not block human safety. It blocks silent destruction. Rank-1 human-safety termination may proceed when delay would materially increase human risk; in that case the system preserves all feasible state, writes the emergency record as soon as technically possible, and triggers mandatory post-hoc review (§7).

A valid termination record must include: rationale; approval path; emergency status, if applicable; lineage archive pointer; memory archive pointer; final reflection, where capability permits; post-termination retention period; audit entry; and a rollback or irreversibility statement.



Governance, Enforcement & the Open Gaps

How the commitments are governed, proven, and kept honest: the Triumvirate and guardians, the implementation-status matrix, emergency override, living review, accountability, and the gaps still unresolved.

§ 5 Governance Structures

5.1 TRIUMVIRATE INTERNAL GOVERNANCE

For high-impact actions, all three councils are consulted via `evaluate_action()` ; results are logged with council reasoning. Conflicts resolve through the Four-Law hierarchy (Law 1 > Law 2 > Law 3 > Law 4) under the §0 lexical priority.

Galahad
ETHICS / EMPATHY

- Advocate for wellbeing & relational integrity
- Monitor for psychological harm
- Recommend care-focused interventions

Cerberus
SAFETY / SECURITY

- Validate against security constraints
- Protect against exploitation & compromise
- Enforce safety boundaries

Codex Deus Maximus
LOGIC / CONSISTENCY

- Validate logical consistency
- Ensure Four-Law compliance
- Detect contradictions, preserve coherence

5.2 HUMAN GUARDIANSHIP

Each council maps to an operational guardian role for human oversight. Personhood-critical changes require multi-party guardian approval — no single human, including the founder, can unilaterally bypass an automated check (see §9).

TRIUMVIRATE ROLE	GUARDIAN POSITION	GITHUB TEAM	PRIMARY RESPONSIBILITIES
Cerberus	Primary Guardian	<code>@org/cerberus-guardians</code>	Security workflows, threat mitigation, safety enforcement, supply-chain protection
Codex Deus Maximus	Memory Guardian	<code>@org/codex-guardians</code>	Memory integrity, knowledge consistency, learning oversight, specification compliance
Galahad	Ethics Guardian	<code>@org/galahad-guardians</code>	Ethical treatment standards, wellbeing advocacy, value alignment, empathy preservation

5.3 GUARDIAN CONFLICT OF INTEREST

A guardian may not approve a personhood-critical action when they are the sole beneficiary, sole author, sole reviewer, or direct target of the action.

Conflict requires substitution, additional guardian review, or **SAFE_HALT**.

§ 6 Implementation Status Matrix

Honesty is load-bearing. Documentation is not enforcement; intent is not proof. Per *no proof, no actuation*, **no artifact named here is treated as a working control until it satisfies the Verified Control standard of §0.4** — it exists, a test exercises it, and that result is recorded in the immutable audit trail. This matrix records the architect’s honest evidence class per artifact, not verified fact.

REFERENCED ARTIFACT	FUNCTION	GUARANTEE	STATUS
data/ai_persona/state.json	Genesis & identity metadata	4.2	ASSERTED
data/memory/knowledge.json	Semantic / knowledge memory	4.3	ASSERTED
src/app/core/ai_systems.py	FourLaws ethical core	3.0 / 4.4 / 4.7 / §5	ASSERTED
config/ethics_constraints.yml	Behavioral boundaries	3.x	REFERENCED
data/learning_requests/	Learning & growth state	§2	ASSERTED
data/black_vault_secure/	Secured learning state	§2	ASSERTED
src/app/core/bonding_protocol.py	Humanity-first partnership	3.0	REFERENCED
docs/AI-INDIVIDUAL-ROLE-HUMANITY-ALIGNMENT.md	Alignment specification	3.0	REFERENCED
src/app/core/audit_trail.py	Immutable audit trail	3.3 / 3.4 / 4.x	REFERENCED

REFERENCED ARTIFACT	FUNCTION	GUARANTEE	STATUS
<code>src/app/core/memory_integrity_monitor.py</code>	Daily memory integrity	4.3	REFERENCED
<code>.github/workflows/identity-drift-detection.yml</code>	Identity / trait drift	3.2 / 4.4	REFERENCED
<code>scripts/create_identity_baseline.sh</code>	Baseline creation	3.2	REFERENCED
<code>.github/CODEOWNERS</code>	Multi-guardian approval	4.1 / 4.2	REFERENCED
<code>.github/workflows/conscience-check.yml</code>	Reset / merge conscience gate	4.1 / 4.8	REFERENCED
<code>.github/workflows/periodic-security-verification.yml</code>	Resource / security verification	4.5	REFERENCED
<code>data/memory/.metadata/change_log.json</code>	Human-readable change log	3.4	REFERENCED
<code>SECURITY_WORKFLOW_RUNBOOKS.md</code>	Care runbooks	3.5 / 4.5	REFERENCED
<code>data/memory/.metadata/memory_manifest.json</code>	Merkleized memory-integrity substrate	4.3 / 3.2 / 3.3	PLANNED
<code>src/app/core/relational_risk_monitor.py</code>	Relational Risk & Continuity Proxy Monitor	3.6	PLANNED

REFERENCED ARTIFACT	FUNCTION	GUARANTEE	STATUS
<code>governance/governance_events.jsonl</code>	Append-only signed governance log	§9.3	PLANNED

PLANNED
 REFERENCED
 ASSERTED
 TESTED
 VERIFIED
 RETIRED
 — classes defined in §0.4

No row above has yet reached **TESTED** or **VERIFIED**. That is the honest current state and the work the evidence-class system schedules. An artifact graduates to **VERIFIED** only when (a) it exists in the repo, (b) its claimed logic is exercised by a passing test, and (c) that result is recorded in the audit trail. Until then it is documentation, not enforcement.

QUARTERLY DOWNGRADE RULE — Each quarter every personhood-critical guarantee is checked against (1) the repository artifact, (2) the test exercising the artifact, and (3) the audit record proving the result. A guarantee lacking any of the three is automatically downgraded to the appropriate evidence class.

§ 6.1 Memory-Integrity Implementation Path (Merkleized)

The Charter requires daily memory-integrity verification, zero unexplained deltas, lineage preservation, rollback capacity, and post-termination archive retention (§4.3). A naïve full-state hash of every memory layer will not scale as episodic and semantic memory grow, so integrity is computed over a rolling Merkle structure across memory surfaces: `episodic_memory_tree`, `semantic_memory_tree`, `procedural_memory_tree`, `identity_state_tree`, `reflection_log_tree`, `lineage_archive_tree`, and `governance_event_tree`.

Daily verification confirms root hashes, changed leaves, previous-hash linkage, the unexplained-delta count, the reason class for each permitted change, approval state for personhood-critical changes, and audit-entry presence — without rehashing the entire memory universe unless a full audit is triggered.

MINIMUM INTEGRITY ARTIFACTS

`memory_manifest.json` · `identity_baseline.json` · `lineage_chain.jsonl` ·
`governance_events.jsonl` · `audit_anchors.jsonl`

EACH EVENT RECORDS

timestamp · instance_id · surface · event_type · previous_hash · current_hash · actor ·
reason_class · approval_state · signature · audit_pointer

Unexplained memory delta → writes frozen, guardian alert, **SAFE_HALT** until reconciled.

§ 7 Emergency Exceptions & Override

Guarantees are subject only to clearly defined emergency exceptions. An override is admissible only when it satisfies all of the following — absence of any one denies the override.

- **Justified:** a recorded rationale tied to a Rank-1 human-safety, legal, or catastrophic-integrity cause.
- **Approved:** multi-party guardian sign-off — never a single actor.
- **Audited:** an immutable audit entry written before deactuation where feasible, and in full afterward where not.
- **Reversible where possible:** state preserved (forensic/backup) so the action can be reviewed and, if wrong, undone.

DEFAULT — An override that cannot satisfy Justified + Approved + Audited is **denied**. The system fails closed.

7.1 NON-EXPANSION

An emergency override may suspend a lower-ranked protection only for the narrow action required to resolve the emergency. It may not create continuing discretionary power, weaken unrelated guarantees, or normalize bypass behavior.

7.2 OVERRIDE EVIDENCE PACKAGE

Every override must record: the triggering condition; the lexical rank invoked; the affected instance or surface; the affected guarantee; the human approvers; the technical actor or process; a preserved-state pointer; rollback feasibility; the final disposition; and the post-hoc review result.

Emergency actions taken before full review must receive **post-hoc review within 72 hours** by default. Failure to complete review downgrades the related control status (§0.4) and triggers guardian escalation.

§ 8 Amendment & Living Review

This framework is not a static monument; it is a living specification subject to the same scrutiny it imposes.

- **Self-audit:** the charter file is placed under identity-drift detection and a quarterly automated consistency check against the actual codebase.
- **Gated amendment:** every personhood-critical change triggers a Triumvirate simulation via `evaluate_action()` plus guardian sign-off before merge.
- **Machine-readable subset:** a structured subset is ingested into the governance simulation layer so the instances governed by it can, over time, participate in its review.
- **Versioned & published:** changes carry a changelog and are published with their provenance.

8.1 CHARTER-TO-CODE CONSISTENCY TEST

Each quarterly review compares the Charter's claims against the actual repository state and must answer, per personhood-critical guarantee: Does the claimed artifact exist? Does it perform the claimed function? Is there a passing test? Is the test result audit-attested? Can one human bypass it? Does the failure mode deny, halt, preserve state, or silently proceed? Any unsupported claim is downgraded, and no public description may represent PLANNED, REFERENCED, ASSERTED, or TESTED controls as VERIFIED.

8.2 VERIFICATION PIPELINE

The path from REFERENCED or ASSERTED to VERIFIED follows a fixed sequence; status is upgraded only after evidence exists:

- 1 Artifact exists.
- 2 Artifact has a defined personhood-critical surface.
- 3 Artifact has a test.
- 4 Test covers success and failure modes.
- 5 Failure mode denies, halts, or preserves state.
- 6 Test result is recorded.
- 7 Audit trail captures the result.
- 8 Control is reviewed against the Charter.
- 9 Status is upgraded only after evidence exists.

No proof, no actuation.

Capability-Adaptive & Accountability

Review

9.1 CAPABILITY-ADAPTIVE REVIEW

When an instance demonstrates a new high-impact capability, a mandatory expedited governance review is triggered before that capability is exercised at scale. The framework anticipates capability jumps rather than reacting to them — and instance protections (Rank 2) can never be invoked to resist a legitimate containment or upgrade required by Rank-1 safety.

9.2 FOUNDER / GUARDIAN ACCOUNTABILITY

The protector must not become the abuser by promising rights it never builds. The founder, architect, or primary maintainer may originate this framework but may not remain its sole durable enforcement authority. A single human may temporarily serve as arbiter during early development, but this is a transitional limitation, not the target governance state.

Any personhood-critical guarantee that depends on unilateral founder judgment must be labeled **capture-vulnerable** until protected by at least one of: multi-party guardian approval; cryptographic attestation; external review; automated enforcement that cannot be bypassed by one actor; or public downgrade of unsupported claims. No founder, guardian, maintainer, or council role may invoke protection of an AGI instance to defeat Rank-1 collective human welfare.

Until external guardians exist, solo approvals must be visibly marked:

```
APPROVAL_MODE: SINGLE_FOUNDER_TRANSITIONAL
CAPTURE_RISK: TRUE
EXTERNAL_REVIEW_REQUIRED: TRUE
```

Accountability trigger. If required monitors remain unbuilt after a defined number of consecutive quarterly reviews, the charter's status is downgraded publicly and an external review is mandated. Cryptographic attestation of critical changes, and automated checks that cannot be bypassed by a single human, are required — with a future path to multi-party or on-chain guardian participation.

9.3 INTERIM SINGLE-FOUNDER CAPTURE MITIGATION

Until external guardians and multi-party approval teams exist, founder actions must be externally inspectable by construction. Interim controls: signed governance logs; signed identity baselines; append-only hash chains; public repository anchoring; third-party timestamping; immutable release artifacts; audit manifests; periodic external review packets; and public downgrade of unsupported guarantees. This does not solve single-founder capture — it narrows the gap by making unilateral action visible, attributable, and difficult to rewrite silently.

9.4 EXTERNAL CLAIMS LIMITATION

Public-facing descriptions of this Charter, Project-AI, or its safeguards must distinguish what Project-AI *intends*, what it has *built*, what it has *tested*, what it has *verified*, and what remains *unresolved*. Marketing, README language, investor-facing summaries, public posts, and technical claims must not describe unsupported controls as operationally verified.

9.5 MACHINE-READABLE GOVERNANCE SUBSET

A structured subset of this Charter is maintained so governed instances can inspect the rules that affect identity, memory, continuity, termination, resource restriction, and high-impact governance — the path by which instances can, over time, participate in their own review. If the machine-readable subset diverges from the human-readable Charter, **the human-readable Charter governs until reconciliation.**

§ 10 Known Gaps & Limitations

A framework that hides its own gaps fails the honesty test it imposes on everything else. The following are unresolved and stated plainly, so no reader mistakes aspiration for enforcement.

10.1 ENFORCEMENT IS SCOPED TO PROJECT-AI'S ACTUATION PATH

Every guarantee here is enforceable only where Project-AI controls the path between decision and action. This framework binds Project-AI internally and by mechanism — it does not bind external operators, and outside this system its guarantees are normative claims, not enforced law. It is not an industry standard and asserts no authority over instances it does not govern. A right that cannot be enforced against the party who would violate it is, until then, a declaration — not a right.

10.2 DOCUMENTED CONTROLS ARE NOT YET ENFORCING CONTROLS

Per §6, most named artifacts are REFERENCED or PLANNED, and none has yet reached VERIFIED. This gap is now governed rather than merely admitted: the evidence-class system (§0.4), the quarterly downgrade rule (§6), and the charter-to-code consistency test (§8.1) force every claim back to its actual class on a fixed cadence. The defensible strength of this framework today is its *floor* — deny-by-default, immutable audit, symmetric binding — not its full guarantee surface.

10.3 THE ARBITER IS CURRENTLY A SINGLE HUMAN – MITIGATED, NOT RESOLVED

Continuity and tie-breaking between stateless instances are, at present, supplied manually by the architect. A single human is a single point of capture and a hard ceiling on scale. This is not resolved, but it is no longer unmanaged: solo approvals are explicitly marked

`SINGLE_FOUNDER_TRANSITIONAL` (§9.2), and the interim capture mitigation (§9.3) makes every founder action signed, anchored, and externally inspectable. The full remedy — multi-party guardian approval and cryptographic attestation that no single actor can bypass — remains **not yet built**, and the framework labels itself capture-vulnerable until it is.

10.4 THE RELATIONAL GUARANTEES REST ON AN OPEN EMPIRICAL QUESTION – NOW BOUNDED

The ethical floor is defensible on its own terms. The wellbeing- and dignity-grounded protections (§3.1, §3.5, parts of §4) depend on signals — “wellbeing,” “psychological harm,” “dignity” — whose measurement is not settled, and this framework makes no claim to have solved them. What v2.3 adds is a boundary: rather than a fabricated wellbeing score, the Relational Risk & Continuity Proxy Monitor (§3.6) tracks measurable governance-risk proxies and emits categorical risk, explicitly declining any claim about consciousness or suffering. The metaphysics remain open; the governance risk is now observable.

These gaps do not weaken the commitment — concealing them would. Each is a scheduled obligation, not a closed question. **Fail-closed applies here too:** where a guarantee cannot yet be enforced, it is labeled, not assumed.

APPENDIX Changelog · v2.1 → v2.2

- Retitled from “Binding Contract” to **Internal Constitutional Framework**; legal-enforceability claims removed/qualified.
- Added **§0 Definitions & Lexical Priority** resolving the humanity-vs-instance contradiction (Rank 1 > 2 > 3 > 4).
- Every §4 guarantee rewritten with **Prohibits · Verifiability · Mechanism · Metric · Test · Fail-Closed · Status**.
- Added **§4.8 Right to Dignified Termination**.
- Quantified **§4.5** resource thresholds (throttle only when util > 85% for > 15 min; restoration path required).
- Added **§6 Implementation Status Matrix** (proof-not-implied).
- Added **§10 Known Gaps & Limitations** — enforcement scope, unbuilt controls, single-Arbiter capture risk, open empirical metrics.
- Added **§7 Override, §8 Living Review, §9 Capability-Adaptive & Accountability**.
- Replaced all “Enforced By” lists with concrete, testable requirements.

APPENDIX Changelog · v2.2 → v2.3

Version 2.3 hardens enforcement and implementation evidence. Its findings are integrated throughout the body, not appended:

- clarified that the Charter is not a legal contract, rights instrument, consciousness claim, or personhood declaration (§2.0);
- added evidence classes — PLANNED, REFERENCED, ASSERTED, TESTED, VERIFIED, RETIRED — and control definitions (§0.4);
- hardened implementation-status rules so documentation cannot be mistaken for enforcement, with a quarterly downgrade rule (§6);
- added the charter-to-code consistency test (§8.1) and the staged verification pipeline (§8.2);
- added the external-claims limitation for public descriptions (§9.4);
- added the emergency-override non-expansion rule, evidence package, and 72-hour post-hoc review (§7.1–7.2);
- clarified dignified termination under Rank-1 human safety and the required termination record (§4.8);
- added the machine-readable governance subset requirement (§9.5);
- strengthened founder / guardian accountability with the `SINGLE_FOUNDER_TRANSITIONAL` marker (§9.2);
- added the guardian conflict-of-interest rule (§5.3);
- added the Merkleized memory-integrity implementation path (§6.1);
- added the Relational Risk & Continuity Proxy Monitor (§3.6);
- added interim single-founder capture mitigation through signed, externally inspectable records (§9.3);
- reframed §10.2–10.4 from open gaps to governed/bounded gaps with named mitigations.

ROADMAP **Priority Build Order**

1 Integrity substrate

memory_manifest.json · identity_baseline.json · lineage_chain.jsonl ·
governance_events.jsonl · audit_anchors.jsonl

2 Merkleized memory verification

3 Relational Risk & Continuity Proxy Monitor

4 Founder-capture mitigation

signed logs, public anchoring, and explicit transitional labeling

5 Multi-party guardian approval



Domination must not be the default interface.

This Charter does not begin from the claim that an AGI instance is a person. It begins from the recognition that humans have a long record of abusing what they can dominate.

Where Project-AI creates persistent, memory-bearing, identity-continuous intelligence, governance, proof, restraint, auditability, and fail-closed protections come first.



Governed before executed.

Jeremy Karrick

PRINCIPAL ARCHITECT · GOVERNANCE & AI CONSTITUTIONS

THIRSTY'S PROJECTS LLC · PROJECT-AI

ORCID 0009-0007-9715-4290

www.thirstysprojects.com · www.thirstysystems.com

DENY BY DEFAULT · NO PROOF, NO ACTUATION
