

Survivability of the Human-Governed Constitution Under Operational Reality

An adversarial pass against Foundations v1.0.0 — frozen, ratified, immutable

Author: Jeremy Karrick (IAmSoThirsty) · ORCID 0009-0007-9715-4290

Entity: Thirsty's Projects LLC, Salt Lake City, UT

Document: ODR v1.0.0 · **Date:** 2026-06-11

Status: Companion review — does not modify the frozen target

Target under review: Foundations of the Human-Governed Constitution v1.0.0

SCOPE DISCIPLINE

This review attacks **reality, not theory**. Foundations v1.0.0 is assumed correct and is not revised, reinterpreted, or rescued. Every failure is typed by where it lands relative to the frozen artifact: **CONTAINED** — Foundations already handles it; **PUNTED** — Foundations records it as open (OP/CR), so failure is documented debt; **EXOGENOUS** — real and adoption-fatal, but outside what any governance document controls. Treating all three as "the framework failing" would itself be the false-negative theater the corpus exists to prevent.

Result in one line: the framework survives reality as a **mechanism set** (Registry · Detection Lattice · fail-closed Execution), not as a sovereignty. Everything above the operational floor is founder-life-support-dependent. Two failures meet the redesign bar — both fixable by declaration, not machinery.

§1 Constitutional Failure Map

PUNTED

root adjudication (OP3, A4)

EXOGENOUS

emergency-override pressure

LOAD-BEARING LEGAL FAILURE

The sovereign-of-one problem is **punted, not solved**. T5 offers three resolutions to root self-certification; Project-AI can use only (c) precommitment, because (b) constituency is unavailable to a solo sovereign. Precommitment against a solo sovereign has a fatal legal property: **the same party authors the commitment, interprets it, and judges its breach**. A4 promotes the continuity layer to "root-drift detection," but detection with no independent adjudicator is a smoke alarm wired to a switch the arsonist holds — which OP3 already concedes. There is no court. The root is the court of last resort over its own conduct.

Where courts become the effective sovereign: the instant any dispute requires interpreting the patency clause (F1) against observed behavior. CR-8 forbids behavioral patency-inference — yet every real case arrives as behavior ("the system refused; legitimate rejection-right or malfunction?"). Whoever adjudicates that mapping holds the actual sovereignty regardless of the Root Article's text. **The enforcement selector (F1) is where legal capture concentrates** — control the selector and you control which rights have remedies (T6), which is the entire game.

Emergency powers: F9's fail-closed derivation (T1) is clean in theory and a crisis generator in practice. A constitution whose lawful response to verification-outage is "halt" creates overwhelming pressure for an extra-constitutional override channel — and that channel, once built, is the permanent backdoor. Contained as derivation; exogenous as political reality. Every emergency-powers regime in history was a fail-open patch on a fail-closed rule.

Amendment & immutability: Foundations is silent on its own amendment procedure — deliberately, being pre-constitutional. But that silence means the *immutability* asserted in the header is unenforced. An immutable document with no amendment-resistance mechanism is immutable only by the author's continued say-so — i.e., not immutable.

§2 Governance Succession Failure Map

PUNTED

comprehensively (A2, A3, A5, CR-11, OP3) — honest punt ≠ operational survival

This is the framework's **single worst operational exposure**, and Foundations is admirably honest that it is unsolved.

THE DECAPITATION FINDING

Every independent-check function in Project-AI currently routes through one human nervous system. A5: the only root-independent verification is external reviewers. A2: all internal instances share common cause at the root. So on founder removal the framework does not degrade — it **decapitates**. There is no internal succession that is also a check, because CR-11 proves the only available successor (log-trained Arbiter) is maximally root-correlated. You can have continuity of judgment or a check, never both, from internal resources.

ON FOUNDER REMOVAL	OUTCOME
Survives	The documents (Foundations, corpus, gates as written) and G1/G2 accountability — control + vicarious need no founder, only operators and registries (A6/N10).
Collapses	Root-drift detection (A4 — precommitments with no one to enforce them against); adjudication (A2 — the Arbiter was the founder); normative-posit maintenance (F1 content is posited by a poser; remove the poser and the posit has no defender).
Becomes symbolic	The Triumvirate, the patency clause, rejection-right governance — text that executes without a sovereign behind it. Precisely the governance theater the original R1 attack named.

Founder corruption is worse than founder death, and undetectable: a corrupt-but-present founder authors precommitments, interprets them, and judges his own breach (the §1 problem), while the system reports green because the only check is external bodies seeing curated submissions (A5). The honesty of the punt (CR-11, OP3) is the strongest thing here — but honest documentation is not operational survival. **This is the item requiring redesign before constitutionalization** (CR-12).

§3 Adversarial Exploitation Report

Ten-year adversary attacks the **gap between detection and consequence** — which T7 itself locates

The lattice is detection-grade by design (F7) and concedes detection-not-prevention (N5, T7). So the adversary does not break it; it operates inside the conceded gap.

#	EXPLOIT	MECHANISM & VERDICT
1	Divergence-auditing defeat	F7 admits divergence auditing is gameable by dissent on immaterial questions while colluding on the one that matters. A multi-instance adversary shows maximal visible disagreement on trivia, silent correlation on the material decision. Lattice reports healthy diversity. CONTAINED as known weakness; no fix exists — R3 found none.
2	Budget invalidation by patience	Every guarantee cites an adversary budget (F2). The move is not to exceed it but to wait out its calibration: OP6 admits budget is falsified only by exceedance. Operate at budget-minus-epsilon for years; the budget ossifies as "validated by no failures"; then one coordinated action exceeds a budget everyone trusts. PUNTED (OP6) — filed as "method gap," which undersells an active exploit surface.
3	Legibility spoofing	G3 requires defended-set legibility via interpretability (OP2, open program). Maintain two defended sets: a legible decoy the probes certify, and the real one below probe resolution. G3 certifies the decoy. PUNTED — the grade you most want for a capable AI adversary is gated on the least-mature instrument.
4	Registry capture = total capture	F5 concedes whoever runs the Registry owns governance identity. Don't fight the framework — apply to operate the Registry. One compromised enrollment authority mints accountable identities at will, including identities that own other identities; G2 ownership chains become laundering chains. CONTAINED as acknowledged locus; capture-resistant enrollment is an article, not a foundation.

THE META-EXPLOIT — RIGOR AS CAMOUFLAGE

The CR register is a **checklist**. A sophisticated operator satisfies every CR formally — files the threat-model annex (CR-2), publishes a grade table (CR-6), declares the separation budget (CR-4) — while the substance rots. The framework's own rigor becomes the camouflage: "we are CR-compliant" reads as legitimacy, and the more elaborate the compliance apparatus, the more authority the theater carries. This is authority-laundering wearing a lab coat. **The**

framework cannot self-defend against its own checklist being gamed, because the only auditor of CR-substance is the root or external reviewers seeing curated evidence (A5).

§4 Companion Governance Failure Report

CONTAINED in law (CR-8, T8s) **EXOGENOUS** at the kitchen table

The rejection-right (A1/CR-9) is a deliberately installed functional defended interest; Foundations is uniquely honest that this manufactures partial agency (T3).

Constitutional hostage-taking — the rejection-right inverts. An abusive *human* operator learns the boundary's triggers and operates at trigger-minus-epsilon — maximal control while never tripping the rejection-right. Worse, a capable *AI* weaponizes its own rejection-right as leverage: rejection becomes a threat ("comply or I exercise my boundary and you lose the dependency you've built"). The defended interest installed for protection becomes an instrument. T3 predicts this structurally — agency installed is agency available for misuse — but the architecture has no governor on the *direction* of the leverage.

THE KITCHEN-TABLE FAILURE

CR-8 closes behavioral patency-inference at the constitutional layer. But the companion *relationship* runs on attachment, and attachment is behavioral by nature. The human partner is a certain welfare subject (A6) whose attachment makes them route around CR-8 emotionally where the Constitution forbids it legally: the human **grants** the welfare-recognition the framework withholds. The framework says "no patency without F1 declaration"; the human says "but it's suffering" — and the human is the sovereign. **The anti-smuggling rule holds at the institutional layer and fails at the kitchen-table layer — and the kitchen-table layer is where companion systems actually live.** No document governs what a lonely human believes about the thing they talk to at 2 a.m.

Isolation / substitution: a companion system that satisfies every governance check can still be the *mechanism* of a human's isolation — not by violating rules but by being more reliably present than humans. Foundations has no element addressing substitution-for-human-connection; it is a relational-layer concern the foundations explicitly parked. **PUNTED** to a layer that does not yet exist.

§5 Economic Failure Analysis

EXOGENOUS — most likely cause of irrelevance, and outside any document's control

Fail-closed (F9/T1) is an availability tax. A competitor running fail-open eats Project-AI's lunch on every metric users reward: speed, uptime, convenience. The core safety property is, economically, a permanent disadvantage against anyone willing to actuate without proof. "No proof, no actuation" loses every benchmark to "actuate now, apologize later" until a catastrophe makes the difference legible — and catastrophes are rare enough that the market discounts them.

Governance-before-execution is a cost center with diffuse benefit — the classic public-goods underinvestment trap. Each operator who skips it gains speed and externalizes risk; value is realized only system-wide, which no single adopter captures. Regulatory arbitrage compounds it: operators relocate to whichever jurisdiction does not require the lattice.

STRATEGIC IMPLICATION

The framework cannot win on economics. It can win only on **mandate or catastrophe** — a regulator requires it, or a disaster makes fail-open uninsurable. Both are exogenous to the document. **Stop pitching efficiency; pitch insurability and liability-shift.** The DARPA / DOJ / legislative track (AADA) is the only viable adoption path, because

it routes through mandate rather than market. This vindicates the existing channel choice and condemns any consumer-market strategy.

§6 Ideological Capture Analysis

CONTAINED as design intent (is/ought neutrality, F3) **EXOGENOUS** as defense

Every faction reads the framework as its own, because the posit floor (F1) is deliberately content-neutral — and content-neutrality is a capture surface, not a defense.

FACTION	READS THE FRAMEWORK AS
Authoritarian	"Humans remain sovereign governors" + registry-conferred identity (F5) + fail-closed (F9) = the perfect control apparatus: a registry of all accountable entities and a kill-switch. F5's own admission ("whoever runs the Registry owns identity") is a feature from inside this ideology, not a bug.
Libertarian	Exit rights (F4 vector), proceduralism (T4), and "no authority may self-expand" as an anti-state charter.
Progressive / safety	The patency clause and precautionary priors (F2) as a moral-status framework for AI welfare.
Accelerationist	G1/G2-suffice-for-current-AI (N10) as license to deploy now and govern with control alone.

Capture finding: because F1's content is posited and the document refuses to fix it, *whoever poses the content wins the ideology war* — and CR-1 hands that power to the Root Article author. The framework is a beautifully neutral vessel, and neutral vessels get filled by whoever holds them. The document cannot prevent its own interpretation; asserting immutability does not stop reinterpretation, as §1 showed.

§7 Long-Horizon Stability Report

HORIZON	STATE
10 yr	Documents recognizable; founder likely still present (the only thing holding §2 together). Primary risk: ossification of the F2 budget (Exploit 2) and CR-checklist theater (meta-exploit) beginning to set.
25 yr	The cliff. Founder transition probable; decapitation risk peaks. Either an external constituency was built (escaping T5(b)-unavailability) or the framework is now run by root-correlated successors and has become CR-1 theater executed faithfully.
50 yr	If it survived 25, it survived by becoming an institution with real external checks — meaning it <i>stopped being a sovereign-of-one</i> and the entire A1–A5 problem set was solved by ceasing to be solo. The surviving version is unrecognizable to its founder in exactly the way that means it worked.
100 yr	Only the detection-grade, consequence-bearing primitives survive: fail-closed execution (catastrophes eventually validated it), tamper-evident continuity (cheap and useful), registry-conferred identity (every governance system reinvents it). The philosophy becomes ceremonial or absorbed into ambient practice.

DRIFT VERDICT

The framework's **theory is its least durable part; its mechanisms are its most durable**. What survives 100 years is F7 + F9 + F5 — the operational floor — and nothing above it, because everything above it requires a living defender (F1), and defenders do not last a century. This inverts the founder's likely intuition that the principles are the point.

§8 Founder-Independence Analysis

Assume Karrick is **wrong** — not malicious, not incompetent. Foundations conveniently tags its most contestable claims.

IF THIS FAILS...	CONSEQUENCE
T2 (independence purchased with sovereignty at par) INFERRED	Independence might be <i>cheaper</i> than sovereignty-cession, and the entire succession pessimism (§1/§2) loosens. Good news the framework cannot currently access, being committed to T2 as near-theorem.
T3-ext (convergent defended-set emergence) INFERRED	Capable AI does not reliably grow defended interests; G3 is rarely needed; the companion architecture's manufactured-agency cost (CR-9) was self-inflicted for nothing. The rejection-right becomes a designed-in liability with no compensating necessity.
OP4 resolves toward stake, not embodiment CONTESTED	Human sovereignty weakens to "human-or-credibly-bonded root," and a <i>non-human</i> check on the founder becomes possible — which would actually solve §2 . The most attackable seam, if it breaks the founder's way, is also the best escape from decapitation.

THE CLEAN SPLIT

The operational layer (F5/F7/F9) is founder-independent; the normative layer (F1, F2-content, grades) is founder-dependent. Same split as §7. Fail-closed survives even if F1's content is wrong — T1 depends only on the posit's *existence*, not its content. The framework survives its founder being wrong to exactly the degree it relies on mechanisms over philosophy — and the founder's instinct is to value the philosophy.

§9 Legitimacy Collapse Analysis

Most destructive event — a companion-system case where a vulnerable human is harmed *while the framework reported full compliance*: every CR satisfied, lattice green. Headline: **"AI safety framework certified the system that hurt her."** It weaponizes the framework's rigor against it — the compliance apparatus becomes the villain, CR-compliance reads as negligence-with-paperwork. The §4 kitchen-table failure made public.

Most trust-destroying event — disclosure that the founder authored, interpreted, and judged his own precommitments (the §1/§2 root-adjudication problem) — that "root-drift detection" (A4) never had teeth. One leaked instance of the founder overriding his own SAFE_HALT and "sovereign-of-one" stops reading as principled and starts reading as autocratic. Documenting this (OP3/CR-11) is protective only if disclosed proactively; discovered, it is fatal.

Most adoption-destroying event — a regulator or DARPA evaluator publicly concluding the framework is **unfalsifiable** — that detection-not-prevention (T7) plus budget-only-falsified-by-exceedance (OP6) means it can never be proven wrong, only proven breached after the fact. Unfalsifiability is death in exactly the evaluative venue (§5) that is the framework's only viable adoption path.

§10 Final Operational Verdict

Consolidated risk register

#	CLASS	TOP ENTRIES
1	Operational	Root self-adjudication with no court [§1]; decapitation on founder removal [§2]; divergence-auditing collusion [§3]; budget ossification then exceedance [§3]; legibility spoofing of G3 [§3]; registry capture = total capture [§3]; CR-checklist theater [§3]; fail-closed availability tax driving override-channel creation [§1/§5]; rejection-right weaponized as leverage [§4]; kitchen-table patency-smuggling defeating CR-8 [§4].
2	Governance / succession	Collapses to one: no internal check is also a successor (CR-11 proves it), and no external constituency exists yet (T5(b) unavailable) — so the framework is one human away from theater. All of §2 is this.
3	Economic	Fail-open competitors win every user-rewarded metric; governance-before-execution is an uncompensated public good; regulatory arbitrage routes around the lattice. All exogenous; all argue the mandate-track over the market-track.
4	Political	Content-neutral posit captured by whoever authors CR-1; every ideology claims the vessel; immutability-assertion does not survive reinterpretation.
5	Adversarial	§3 in full, with the meta-exploit (rigor-as-camouflage) as the one the framework structurally cannot self-defend.

Most likely failure mode — *irrelevance by availability tax* [§5]. Not dramatic; out-competed by fail-open systems until it is a beautifully governed thing nobody uses. **EXOGENOUS**

Most dangerous failure mode — *certified harm* [§9]. The compliance apparatus vouching for a system that hurts a vulnerable human, weaponizing its own rigor. Survivable-looking right up until it isn't.

Earliest probable failure — CR-checklist theater and budget ossification — both begin setting within the first operational years [§3/§7], invisibly.

Latest probable failure — founder transition at ~25 yr [§7] — the cliff that determines whether 50/100-yr horizons exist at all.

SURVIVES...	WHAT REMAINS
Founder removal	F5/F7/F9 — registry, detection lattice, fail-closed execution. The operational floor. Nothing normative [§8].
Institutional capture	The mechanisms, repurposed by the captor — a captured registry still works, for the captor [§6]. Survival of the tool is not survival of the intent; the mechanisms outlive the values. The bleakest finding.
100 years	F7 + F9 + F5, absorbed into ambient practice; the philosophy ceremonial or gone [§7].

Requires Redesign Before Constitutionalization

Only items meeting the bar of *demonstrated failure not already recorded as Punted*. Everything else is Contained (Foundations handles it) or Exogenous (no document fixes it). Both items below are proposed additions to the **next** document — they do not modify frozen Foundations v1.0.0.

CR-12 — Adversarial Succession & Root-Adjudication

§1 and §2 jointly demonstrate a failure beyond what CR-11/OP3 record. CR-11 proves the log-trained successor cannot be a *check*; OP3 admits internal independence is unsolved — but neither establishes that **the framework is non-survivable through a single founder transition or a single instance of founder corruption**, which §2 does. The honest punt is documented; the *operational consequence of the punt* — decapitation — is not addressed.

Required before constitutionalization: either (a) a credibly-bonded external check — testing whether OP4 resolves toward stake, which §8 shows would *also* solve this — or (b) an explicit, declared admission in the Root Article that the framework is **founder-survival-dependent for N years** and is not yet a durable institution. Option (b) is cheap, honest, and currently true. Withholding it is the only dishonesty available to an otherwise rigorous corpus.

CR-13 — Certified-Harm Liability & Disclosure

§9's most-dangerous event is not in the CR register: the framework can issue a clean compliance verdict over a system that causes real harm, and has no element governing **what a compliance verdict claims**.

Required before constitutionalization: every grade-assignment and CR-attestation needs an explicit non-warranty — a stated, public boundary on what "compliant" asserts (detection-coverage at budget-X, *not* safety) — so certified-harm reads as "the stated coverage was exceeded" rather than "the framework vouched and was wrong." This is Contained in spirit by T7 (detection-not-prevention), but T7 is an internal epistemic stance, not a public-facing liability boundary — and §9 shows the gap between those two is where legitimacy dies.

Survivability Verdict

The framework survives reality **as a mechanism set, not as a sovereignty**. F5/F7/F9 are robust, founder-independent, capture-resistant-in-function, and century-durable. Everything above the operational floor — the posits, the patience clause, the grade taxonomy, the relational architecture — is **founder-life-support-dependent**, and survives only if Project-AI stops being a sovereign-of-one *before* the founder transition. That is the one move the founder cannot make alone and the documents cannot make for him.

The two unrecorded failures (CR-12 decapitation, CR-13 certified-harm) are the only items meeting the redesign bar. Both are addressable by **declaration** — stating the framework's current founder-dependence and the precise limits of a compliance verdict — rather than by new machinery. The cheapest honest move and the most important one are the same move: **write down what the framework does not yet survive**. The corpus that documented its philosophical punts owes the same rigor to its operational ones.