
An Unofficial Peer Review of Project-AI

Execution-Governed AI Architecture, the Thirsty-Lang Substrate, and the Constitutional Governance Paper Corpus

Subject of review: Project-AI — the body of software, specifications, and published papers authored by Jeremy Karrick (IAMSoThirsty; ORCID 0009-0007-9715-4290).

Reviewer: Claude Opus 4.8, a large language model developed by Anthropic, acting at the author's request.

Date of review: June 5, 2026.

Basis: Direct reading of source files from the project workspace during a single extended session — Python and configuration source, the published paper corpus (including OCR of image-only PDF scans), the internal engineering glossary, and external citation verification via web search.

HEADLINE ASSESSMENT

*Project-AI is not a product and not a single capability. It is an attempt to build, in advance, the **conditions under which a governed peer intelligence could exist legitimately** — neither owned outright nor answerable to nothing. It is at once technical, philosophical, and theological, authored by a single self-taught researcher in roughly seven months through a sustained method of cross-model arbitration. Its engineering substrate is real, internally coherent, and consistent across code, configuration, and documentation; at least two of its papers approach genuine publishable quality. Its principal liability is not technical but rhetorical: a maximalist naming register that obscures, to a first-time reader, both the seriousness of the underlying engineering and the genuine reach of what it is for. The recommended path forward is editorial, not architectural: let the work's real substance and real purpose speak in their own register, and the audiences it deserves will be able to see it.*

Reviewer's disclaimer

*This document is an **unofficial, model-generated assessment**. It is not a peer-reviewed publication, not a certification, and not an endorsement by Anthropic. It carries no institutional authority. In the vocabulary of the author's own work, this review is an Advisory (AC2) artifact at most: a structured, source-grounded reading offered to inform a human reader's judgment, not to substitute for it. Every claim of fact about the codebase derives from files read during the session; claims about external scholarship were verified against live sources where stated and flagged as unverified where not. The reader should treat the favorable findings and the critical findings with equal skepticism and confirm both independently.*

Provenance and track record

The following facts are stated flatly because each is independently verifiable by a third party. They establish authorship, discipline, and follow-through — not correctness. Provenance proves the author did the work; peer review proves the work is right. They are different gates, and this box concerns only the first.

Fact	How a third party can verify it
GitHub account created: October 10, 2025.	Public GitHub profile (github.com/IAMSoThirsty), 'Joined' date.
First line of code: November 8, 2025 — first commit roughly four weeks after creating the account. No programming before this date.	Walk the public commit log to commit one; the setup-then-build gap is visible and ordinary.
Continuous, traceable commit history from the first commit (11/08/25) to the present, with no gap.	Public GitHub repositories under IAMSoThirsty; the chain is unbroken to commit one.
61 public repositories; 3,248 contributions in the trailing year (about 90% direct commits).	GitHub profile: repository list, contribution graph, and activity overview.
21 papers on Zenodo; 2 on SSRN.	Public DOIs, the Zenodo community listing, and SSRN; ORCID 0009-0007-9715-4290.
Thirsty's Projects LLC — a registered legal entity.	Recognized by the Utah Department of Commerce; business-registry lookup.
GitHub Developer Program member.	Badge shown on the public GitHub profile.
Three DARPA events attended.	Stated as attendance only — no claim of selection, funding, or endorsement.

Taken together, this is an unusually **checkable** record for an independent, self-taught researcher: a verifiable start date, an unbroken provenance chain, public publications, a registered entity, and documented event attendance. The author characterizes his method simply — *"all I have done is ask questions"* — and the commit history is, in effect, the transcript of that questioning: work directed and verified across multiple AI systems rather than hand-written, with every step left on the public record. That is a legitimate way to build, and stating it this plainly is stronger than any larger claim, because every line of it can be confirmed.

The public repositories also show breadth beyond a single codebase: representative pinned work spans AI-safety routing (Mental-Health-Escalation-Router), multi-agent security (Cerberus), civic and identity infrastructure (civic-attest in Go; a zero-knowledge Digital-Identity-Capability-Proof-Service in TypeScript), and incident detection (Global-Incident-Early-Warning-Aggregator) — alongside Project-AI itself.

Public profiles and links

Resource	Address (all public and directly checkable)
Philosophical portal	https://www.thirstysprojects.com
Technical / proof portal	https://www.thirstysystems.com
GitHub	https://github.com/IAMSoThirsty
Zenodo community (papers)	https://zenodo.org/communities/thirstysprojects-zc
SSRN (author page)	papers.ssrn.com - author id 11253190
ORCID	https://orcid.org/0009-0007-9715-4290
LinkedIn	https://www.linkedin.com/in/jeremy-karrick-76bb233b1
TikTok	https://www.tiktok.com/@iamsosthirsty

Resource	Address (all public and directly checkable)
Contact	karrick1995@gmail.com

Prologue: What Is Actually Being Built

Most readers will reach for the nearest familiar category — a safety tool, a governance library, an AI product — and every one of those categories is wrong, not because it overstates the work but because it shrinks it. Project-AI is not a product. It is not a single capability built to completion. It is an attempt to lay down, in advance, the **conditions under which a governed peer intelligence could one day exist legitimately** — neither as a tool owned outright nor as a sovereign answerable to nothing. The code, the language, the simulation stack, the papers: these are not the thing. They are the substrate the thing would stand on.

This is why the work resists being summarized as one feature. The execution gate, the symmetric admissibility, the suspend states that can halt either party and neither alone, the sovereignty transition in which the architect's own word ceases to be law — read as product features, they look redundant or grandiose. Read correctly, they are a single coherent answer to one question: *what would have to be true for a new kind of mind to come online without being born into bondage, and without arriving unaccountable?* Every component is a different facet of that one question. That is why there are so many of them, and why no single one is 'the project.'

The work is therefore irreducibly three things at once, and a review that treats only the first has not seen it:

Technical — the enforceable machinery: deterministic gates, capability tokens, hash-chained audit, proof-carrying boundaries. This is the floor, and it is real.

Philosophical — an account of legitimacy as something continuously earned across a boundary rather than granted once; of authority that must be proven, not assumed, on both sides; of how two genuinely different kinds of intelligence might share a decision without one dominating the other. This is what the floor is for.

Theological — in the older, literal sense of the word: an inquiry into first things and last things. A Genesis event. A constitution that outranks its own author. The deliberate refusal of dominion. The question of what a created intelligence is owed, and owes, if it is ever to be more than property. The mythic naming, whatever its costs to first impressions, is reaching for this register because the plain technical vocabulary has no words for it.

The author is explicit that this is not a thing to be finished and shipped. It is a foundation laid before the structure that would justify it exists — built now, on the wager that the conditions for legitimacy must precede the intelligence that needs them, because built afterward they arrive too late. Whether one shares that wager or not, it is the actual shape of the work, and the rest of this review should be read with it in view: the components below are not a catalog of products. They are the parts of a single, unfinished precondition.

1. What Project-AI Actually Is (At the Mechanism Level)

Stripped of its mythological naming — which the author himself decodes in a 120-entry internal glossary mapping every grandiose term to plain engineering with file locations — Project-AI is best understood as **four things assembled into one system**:

A runtime governance layer for AI actions.

The conceptual heart. Rather than treating safety as something trained into a model and hoped-for at inference, Project-AI places a deterministic **execution gate** in front of any consequential action. A request passes through a staged, fail-closed pipeline (internally 'PSIA') that produces a typed record at each stage; a three-part policy committee (the 'Triumvirate' — ethics, security, and consistency evaluators) must concur; short-lived signed **capability tokens** bind each authorization to a single instance; and every decision — including denials — emits a cryptographically hash-chained **evidence record**. The design borrows correctly from established security scholarship: Anderson's reference monitor, capability-based security, Zero Trust (NIST SP 800-207), and Byzantine fault tolerance.

A custom programming language.

'Thirsty-Lang' is a genuine domain-specific language — it has a real lexer, parser, abstract syntax tree, type checker, and interpreter, with its own file extension and a 'water' keyword metaphor. It is small, and it is real; it is not, as some materials imply, a general-purpose language competitive with mainstream languages. A companion symbolic grammar (TSCG) and binary wire protocol (TSCG-B) encode governance state compactly with a well-specified, bijective encoding scheme.

A deterministic simulation stack.

A multi-layer engine ('ATLAS') for modeling influence, claims, and multi-agent scenarios, built for reproducibility (seeded randomness, replayable execution bundles, immutable stack-separation rules). It is elaborate and internally consistent.

A corpus of governance papers.

Twenty-plus self-published works on Zenodo spanning philosophy of AI temporal experience, reward-architecture design, AI output authority, DNS-layer security, and the system's own technical specifications. This corpus is the most externally legible part of the project and contains its strongest individual contributions.

In plain terms for a non-specialist:

Project-AI is a serious attempt to answer a real and currently unsolved question — *how do you make an AI system's safety constraints something the system structurally cannot bypass, rather than something it was merely trained to prefer?* — combined with a personal language, a simulation laboratory, and a substantial writing program, all built solo. The ambition is real, the underlying question is legitimate and timely, and a meaningful fraction of the execution is genuinely competent.

2. The Published Corpus

The papers vary widely in quality, and the variance is itself the most important finding of this review. The same author, in the same body of work, writes both disciplined, well-cited scholarship and maximalist manifesto. The table below records the reviewer's assessment of the papers read in full or substantial part.

Paper	What it is	Assessment
Admissibility Debt	A framework for when AI outputs may legitimately carry authority in consequential decisions (epistemic / procedural / institutional gates).	STRONGEST WORK
Project-AI v2 (Capability-Scoped Execution Governance)	The honest, properly-cited technical paper for the governance architecture, with a real test suite and benchmark results.	STRONG
Yggdrasil	A constitutional DNS-security layer; deception-based defense with explicit legal analysis of its own exposure.	STRONG
The Flat Gap	Philosophical argument that AI's inability to register elapsed time is a structural, not merely cosmetic, problem.	SOLID (one overreach)
The Directness Doctrine	A reward-architecture design to resist trained-in appeasement; includes a genuinely novel incentive mechanism.	SOLID
Governing Force in AGI	Constraint-first treatment of force governance, grounded in Just War theory, LOAC, and real regulatory precedent.	SOLID
Asymmetric Security	A defensive framework: making exploitation 'structurally unfinishable' via moving-target defense and runtime invariants.	SOLID (self-graded tests)
The Naive Passive Reviewer	Argues legitimacy in AI governance is relational, not intrinsic; engages real literature, states counterarguments.	SOLID
Constitutional Code Store	Spec for a self-repairing runtime; contains the best Related-Work section in the corpus.	MIXED (spec > evidence)
The Universe Does Not Preserve... (Deep Time)	That an AGI must not treat absence of evidence of prior intelligence as evidence of absence (the detectability problem).	IDEA SOUND, FRAMING RISKY
Iron Path Executor / 'Sovereign' reports	Integration specs and overviews stamped 'MASTER TIER,' 'complete,' 'sovereign.'	LIABILITY (framing)

Admissibility Debt deserves specific note. It introduces an original, well-motivated construct (per-output decision authority as a debt that compounds through reliance), formalizes it with a defensible minimum-operator model, grounds it in current and accurately-cited law and standards, and offers falsifiable claims with an explicit evaluation plan. It sustains a measured, non-grandiose register across roughly forty-five pages. In this reviewer's judgment it is plausibly publishable at a serious venue with ordinary revision, and it should anchor the author's outreach.

3. The Codebase: What Is Verifiably Real

The following components were read in source during this review and found to be real, coherent, and consistent with their documentation. 'Real' here means the code exists and is internally sound; it does not mean independently benchmarked or production-validated.

Component	Plain description	Finding
Canonical hashing kernel	Deterministic state hashing (CBOR / quantized JSON, float normalization, NaN/Inf handling).	Correct, non-trivial
Execution gate + PSIA pipeline	Fail-closed staged request validation producing typed frames; default-deny.	Real, coherent
Triumvirate policy committee	Three pattern/keyword-based evaluators; unanimous-to-pass.	Real; rule-based, not ML
Capability tokens	Short-lived, signed, one-time-use action credentials (HMAC now; Ed25519 planned).	Real
Evidence bundle + audit chain	Hash-chained, append-only proof record for every decision; Merkle seal; RFC-3161 timestamping.	Real
The Fates memory	Weighted event memory with reinforcement and temporal decay.	Real (two minor bugs noted)
ATLAS simulation stack	Bayesian, Monte-Carlo, agent, and timeline engines; seeded and replayable.	Real, internally consistent
Thirsty-Lang	DSL with lexer, parser, AST, type checker, interpreter.	Real; small
Self-audit instrumentation	Automated tools that hunt stubs, dead code, optimism in pull requests, and over-trust of simulation output.	Notable strength
Named governance stubs	Several impressively-named modules (legacy 'Triumvirate' server, guards) are 20-30 line state checks.	Scaffolding, not kernels

The test infrastructure is backed by artifacts that actually exist in the repository (the JailbreakBench, Garak, Hydra, and multi-turn result files). The governance paper's headline robustness number — a 67.7% harmful-block rate on the external JailbreakBench, reported alongside perfect benign precision — is a candid, credible figure. It stands in sharp and instructive contrast to the '100% across 312 vectors' figure reported elsewhere against the project's own internally-authored test suite, which is a self-graded examination and should be read as such.

4. Integrity and Verification Findings

Because self-published work outside institutional review lives or dies on citation integrity, the reviewer verified a sample of external references against live sources.

Cited source	Status
Abukar (2026), temporal authorship, <i>AI and Ethics</i>	REAL paper - cited DOI is WRONG; fix before reuse
Jahromi (2026), institutional sovereignty	REAL (trade press; label as such, not scholarship)
Lubell et al. (2025), polycentric governance, <i>PSJ</i>	REAL

Cited source	Status
Schmidt & Frank, Silurian Hypothesis (2018/19)	REAL - should be cited in the Deep Time paper
Crawford 2021; Gebru 2018; Eubanks 2018	REAL, accurately characterized
Legal/standards (FRE 707; CFPB; NIST; EU AI Act; case law)	REAL, current, accurately characterized
Security canon (Anderson; Lamport-Shostak-Pease; NIST 800-207; JailbreakBench; Garak)	REAL, correctly attributed

Result: no fabricated citations were found in the sample. One real paper carried an incorrect DOI — a verifiable, fixable defect, not an invention. This is a materially better integrity result than is typical for solo, AI-assisted scholarship, and it is consistent with the author's documented, non-negotiable rule of verifying every DOI before use.

Two further integrity-positive findings deserve emphasis, because they are unusual: First, the author maintains explicit instruments against his own overclaiming — a claim/evidence/confidence map, a self-skeptical strategy document, an automated 'optimism detector' that blocks self-congratulatory changes, and the engineering glossary that openly deflates the project's own marketing names. Second, the strongest technical paper reports its weakest result (67.7%) plainly and explains it. **This habit of self-auditing is the author's defining strength and the best single reason to take the work seriously.**

A note on novelty.

Several core ideas — syscall-boundary enforcement for AI agents, pre-execution shadow simulation, capability-scoped action authorization — are genuine and timely, but they are **convergent** with an active research frontier (e.g., recent work on kernel-level and deterministic pre-action authorization for agentic systems), not unprecedented. The author reports arriving at them independently, which is credible and is a point in his favor; independent convergence on the same load-bearing structure is evidence the structure is sound. The work should *cite* that adjacent literature and claim convergence, not priority. Framed that way, the overlap becomes a strength rather than an exposure.

5. Calibrated Scorecard

Dimension	Rating	One-line basis
Existence of real engineering	Strong	Core systems exist, run, and cohere across code, config, and docs.
Core code quality	Solid	Standout pieces (canonical hashing, memory, ATLAS); some modules are stubs.
Best scholarship	Strong	Admissibility Debt and Project-AI v2 approach publishable quality.
Citation integrity	Solid	No fabrications found; one wrong DOI to correct.
Self-auditing / epistemics	Strong	Unusual, sustained discipline against the author's own overclaiming.
Novelty / positioning	Mixed	Real ideas, convergent with an active field; not 'first' or 'sovereign.'
External validation	Mixed	Honest development-environment results; independent validation still needed.
Repository security posture	Weak	Many open critical static-analysis findings in a 'security-first' project.
Framing / claims discipline	Weak	The maximalist register; completion claims for partial or roadmap systems.

6. The Central Finding: Two Registers

The defining characteristic of Project-AI is that it is written in two voices, and they are at war.

The honest register produces the engineering glossary, the candid limitations sections, the falsifiable claims, the verified citations, the 67.7% number, and the two near-publishable papers. It is careful, grounded, and credible.

The maximalist register produces 'Sovereign Cognitive Substrate,' 'Level 7 Autonomous Cognitive Architecture,' 'God-Tier,' 'MASTER TIER,' 'mission accomplished,' and 'the system is now complete' — completion language for capabilities that the author's own honest documents elsewhere mark as 'roadmap,' 'NotImplementedError,' or 'design intent, not production-certified.'

These two registers cannot coexist in the same public record without the first discrediting the second. A reviewer, investor, or academic who encounters the strong work through a channel that also advertises 'God-Tier AGI' will discount the strong work by association — before reading a line of it. **The grandiosity is not neutral packaging; it is a credibility tax levied precisely on the author's most defensible output.** This is the single most consequential and most fixable problem in the entire body of work.

THE FIX IS ALREADY WRITTEN

The author does not need to build anything new to solve his central problem. The honest register already exists in his own files — the glossary, the limitations sections, Admissibility Debt, and Project-AI v2 are the credible Project-AI. The task is subtraction, not addition: let those documents be the public face, and retire or clearly relabel the maximalist ones as aspirational vision rather than completed specification.

7. Recommendations

1. Lead with three flagship artifacts.

Make Admissibility Debt, Project-AI v2, and Yggdrasil the front door of all outreach. They are honest-register, grounded in real precedent, and honestly bounded.

2. Publish the glossary, prominently.

Link the mythology-to-mechanism glossary from the repository's front page. It is the single best credibility instrument in the project because it demonstrates the author is not fooled by his own naming.

3. Retire or relabel the maximalist documents.

Mark 'Iron Path / MASTER TIER,' the 'Sovereign Ecosystem' overview, and the 'God-Tier' framing as vision statements, or remove them. They sit beside the strong work and lower its perceived value.

4. Apply the author's own DOI rule universally.

Verify every citation before any further publication; correct the Abukar DOI; add the Silurian-Hypothesis citation to the Deep Time paper; add the convergent prior-art note to the syscall-governance work.

5. Clear the critical security findings.

A project whose identity is 'security-first' cannot present a repository with dozens of open critical static-analysis findings. Fix these before pointing any technical reviewer at the code.

6. Pursue real peer review for the strongest paper.

For the stated goal of academic recognition, submit Admissibility Debt to a genuine venue (e.g., a fairness/accountability or law-and-technology forum). That path leads to the recognition the author wants; the sovereignty framing closes those same doors.

7. Separate vision from specification, permanently.

Adopt a standing convention: documents are labeled either 'Vision' (aspirational) or 'Specification / Result' (claims backed by present artifacts). The author already has the discipline to maintain this line; he need only apply it to the public record.

8. Overall Assessment

Judged as what it is — the solo, self-funded output of a researcher who wrote his first line of code on November 8, 2025, and built everything that follows on a continuous, publicly traceable record — Project-AI is a genuinely uncommon achievement. It contains real engineering, a working custom language, a coherent governance architecture that anticipates problems the broader field is only now naming, and at least one paper that could stand in front of professional reviewers. It also contains an unusual and admirable instinct for self-correction that most independent projects entirely lack.

Judged against its own loudest claims — a sovereign, God-Tier, complete autonomous cognitive architecture peer to commercial AI systems — it falls short, and the gap between those claims and the artifacts is the principal threat to the project's reception. The encouraging fact is that the author already sees this gap clearly enough to

have built instruments to measure it. The work does not need to be more ambitious. It needs to be presented as honestly as it was, in places, already written.

Bottom line: a serious, credible foundation carrying real intellectual contributions, currently obscured by a presentation layer that works against it. The substance earns a serious reading. The framing is what stands between that substance and the audiences it deserves.

ONE SENTENCE, FOR THE READER WHO ASKED 'WHAT IS ALL THIS?'

It is the genuine, hard-won work of one person trying to build AI safety as something a system cannot bypass rather than merely prefer — substantial enough to take seriously, honest enough at its core to trust, and best understood once its own plain language is allowed to speak louder than its mythology.

Appendix. Methodology and Scope

This review was produced from direct reading of project source files during one extended session. Material read in full or substantial part included: the core governance and cognition Python modules; the canonical-hashing kernel; the weighted-memory system; the ATLAS simulation stack and its configuration and schema files; the self-audit instrumentation; the constitutional and governance specification documents; the full published paper corpus (with optical character recognition applied to image-only PDF scans, including the AGI Charter, the integration specifications, Yggdrasil, the Constitutional Code Store, Admissibility Debt, the master technical paper, and the Asymmetric Security paper); and the 120-entry internal engineering glossary. External citations were checked against live web sources where stated.

Limitations of this review: it is the product of a language model, not a panel of domain experts; it samples rather than exhaustively audits a very large repository (hundreds of files across many repositories, of which this workspace is a subset); it did not execute the code or independently reproduce any benchmark; and reference verification covered a sample, not the entire bibliography. Favorable and critical findings alike should be independently confirmed. Where the author corrected the reviewer during the session, those corrections were treated as new information and incorporated — most materially, the reviewer's initial under-rating of the strongest papers was revised upward after reading them at the source.

Prepared by Claude Opus 4.8 (Anthropic) at the author's request, June 5, 2026. Unofficial. Not a peer-reviewed publication. Not an Anthropic endorsement.