

Reverse BAA / Program Challenge Statement

Governed Execution Substrates for High-Consequence Autonomy

High-consequence autonomous systems should not be able to produce actuator-level effects unless authority, provenance, restraint, and continuity of command are machine-checkable at the moment of execution.

1. Problem Statement

Current autonomy assurance efforts focus heavily on design-time testing, ethical evaluation, simulation, readiness scoring, and doctrinal frameworks. While these are necessary and valuable, they leave a critical gap at the runtime execution layer. When autonomous systems must act at machine timescales, after-the-fact audit and human review cannot preserve meaningful human command over the specific actions taken.

2. The Gap

When systems act faster than human judgment can meaningfully intervene, the absence of a runtime enforcement mechanism means that command authority can erode or become unverifiable without detection or prevention. After-the-fact logging and doctrinal statements are insufficient. The missing architectural layer is a governed execution substrate capable of vetoing or blocking actuator-level actions in real time when valid human command cannot be cryptographically or formally proven at the exact moment of intended execution.

3. Desired Technical Capability

A governed execution substrate that enforces the following four predicates as machine-checkable preconditions before any high-consequence actuator-level effect is permitted:

- **Authority** — A binding, accountable link between the contemplated action (or the policy generating it) and a specific, identifiable human commander or authorizing chain.
- **Provenance** — An auditable, tamper-evident derivation of the decision from declared policies, rules of engagement, mission objectives, training artifacts, and configuration state.
- **Restraint** — Enforced compliance with codified legal, ethical, and safety constraints (including IHL/LOAC and ROE) within clearly quantified risk bounds, with the ability to block or modify actions that would exceed those bounds.
- **Continuity of command** — Persistence and revocability of human authorization across model updates, changing conditions, degraded communications, and extended autonomous operation, such that actions remain attributable to a current, valid human command.

4. Research Questions

1. Can command authority be represented as a machine-checkable runtime object that can be verified under distribution, partial failure, and adversarial conditions?
2. Can autonomous actions carry cryptographic or formal proof of valid authority that must be successfully verified before any actuator-level effect is permitted?
3. Can restraint conditions (legal, ethical, safety) be enforced as hard, blocking predicates at runtime rather than evaluated only after the fact in simulation or test?
4. Can continuity of command be maintained and enforced under model updates, mission drift, degraded or denied communications, and deliberate adversarial manipulation?

5. Metrics

A successful governed execution substrate would be evaluated on its ability to:

- Block actions when valid authority, provenance, or restraint cannot be verified at runtime.
- Permit actions only when all four predicates are satisfied and current.
- Enforce command revocation in real time with measurable latency bounds.
- Maintain a complete, inspectable provenance chain for every actuator-level action.
- Operate with bounded runtime overhead acceptable for the target platform class.
- Default to a safe halted or degraded mode on verification failure or ambiguity.

6. Demonstration

A simulated or hardware-in-the-loop autonomous platform (e.g., UAV, UGV, maritime system, or decision-support node) in which actuator-level actions are structurally impossible unless the runtime governed execution substrate successfully verifies all four predicates at the precise moment the action is contemplated. The demonstration should include cases of valid command, revoked command, degraded provenance, and attempted override under stress or adversarial conditions.

7. Ask to DARPA

Does DARPA view this as a program-worthy gap in the current autonomy portfolio — distinct from existing evaluation, benchmarking, and assurance efforts? If so, which office, program manager, ERIS topic area, or BAA pathway would be the appropriate recipient for a short concept paper further developing this program challenge?

This document is issued as a reverse BAA — a program challenge statement identifying a missing execution-governance problem that the author believes DARPA is structurally positioned to formalize. It is not a formal proposal, white paper, or response to an existing solicitation.

Jeremy Karrick — Principal Architect, Thirsty's Projects
Karrick1995@gmail.com | 385-522-9743 | www.thirstysprojects.com