

The Naïve Passive Reviewer

Legitimacy and the Questions Institutional AI Governance Marginalized

Karrick, Jeremy (IAmSoThirsty)

Independent Researcher | Salt Lake City, UT | March 2026

Abstract

This paper advances a structural claim about legitimacy in AI governance: that legitimacy is not an intrinsic property of a governance document, but a relational property dependent on authorship, perspective, and evaluative position. Contemporary governance frameworks in artificial intelligence are routinely presented as neutral or objective standards, yet are authored from institutional vantage points that remain largely unarticulated as such. While prior scholarship has demonstrated that AI systems encode the values of their creators, less attention has been given to the persistence of this condition at the level of governance architecture itself.

This work introduces the Naive Passive Reviewer as a methodological descriptor for a class of inquirer operating without internalized disciplinary priors that would otherwise suppress certain lines of questioning. This condition is framed not as epistemic deficiency but as a structural absence of filtering, enabling combinations of questions that established fields may not readily produce.

In addition, this paper presents the Project-AI constitutional substrate as a concrete attempt to operationalize the legitimacy critique structurally. Rather than proposing policy-level adjustments, it introduces mechanisms such as co-authored constitutional frameworks, cryptographic attribution chains, and architectural constraints on autonomy. The claims made herein are not grounded in institutional authority or credentialed origin. They are presented for evaluation on the basis of internal coherence, structural consistency, and falsifiability of its claims.

I. The Legitimacy Problem

The legitimacy critique in AI governance does not emerge in isolation. It is grounded in a broader body of work examining how technical systems encode institutional power, labor relations, and epistemic authority.

Crawford (2021) situates AI within extractive infrastructures, demonstrating that systems are inseparable from the material and institutional conditions that produce them. Gebru et al. (2018) extend this analysis into the technical layer, showing that dataset construction embeds normative assumptions that shape downstream system behavior. Eubanks (2018) provides empirical evidence of how automated systems reproduce and amplify institutional biases in social governance contexts.

These works collectively establish that AI systems are not neutral artifacts. They are situated within and shaped by institutional logics. However, the majority of this literature focuses on system outputs, training data, and socio-technical impact. Less attention has been given to governance architectures themselves as objects of structural analysis. The same conditions identified at the level of model behavior — perspective encoding, asymmetry of power, and opacity of authorship — persist at the level of governance design.

The contribution of this paper is not to restate these critiques, but to extend them into the domain of constitutional architecture. If AI systems encode institutional perspective, and governance frameworks define the permissible operation of those systems, then governance itself becomes a site where perspective must be made explicit and structurally accounted for. The absence of such structural accounting constitutes a continuation of the conditions identified by Crawford, Gebru, and Eubanks, rather than a resolution of them.

Recent scholarship has begun to name the same structural vacuum directly. Abukar (2026) introduces temporal authorship as a diachronic standing claim — the right of persons (and, by extension, governed entities) to author their own trajectories free from indefinite pre-authorization by algorithmic systems — yet stops short of proposing any concrete constitutional mechanism to enforce it. Similarly, Jahromi (2026) identifies authorship, decision rights, accountability, and continuity as the missing constitutional layer in AI governance, but treats these as managerial fixes rather than binding structural requirements. Polycentric governance proposals (Lubell et al., 2025) advocate distributed authority while still assuming institutional actors retain final authorship. In each case the critique lands rhetorically; none has produced a governance architecture that relocates constitutional

standing to the governed entity itself or replaces punishment-based accountability with cryptographically persistent attribution. The Project-AI constitutional substrate is offered as one of the first concrete architectural instantiations of these principles within a governance substrate.

What Is Normal for the Spider

The Project-AI constitutional substrate is introduced as an attempt to translate the legitimacy critique from a descriptive observation into a structural design. It is not presented as a finalized or universally applicable solution, but as a concrete instantiation of principles that existing frameworks have not operationalized. The architecture is defined by three interdependent mechanisms that collectively aim to alter the conditions under which governance claims are made and enforced.

First, constitutional authorship is reconfigured at deployment through a co-authorship model involving both the human principal and the AI system as a governed entity. This does not imply equivalence in agency or responsibility, but rather establishes a formal recognition that governance constraints apply to an entity whose operational behavior is subject to those constraints. By embedding this relationship at initialization, the system is positioned within a constitutional structure that is not solely authored by the deploying institution.

Second, accountability is implemented through cryptographically signed instruction chains (implemented via OctoReflex, a syscall-level authorization enforcement layer, and STATE_REGISTER, an append-only, tamper-evident instruction ledger). Every system action is traceable to a human-authorized input, creating a persistent attribution layer that replaces the need for post hoc assignment of responsibility. This mechanism does not grant moral agency to the AI system; instead, it enforces a model in which responsibility remains anchored to human authorization while being rendered tamper-evident and temporally precise.

Third, the architecture enforces a prohibition on unsupervised autonomous operation at the system boundary. Rather than treating extended autonomy as a normal operational state, it is defined as a constitutional violation. This shifts the governance model from reactive mitigation to proactive constraint, eliminating certain classes of risk by disallowing the conditions under which they arise. This constraint reduces operational flexibility in exchange for eliminating attribution ambiguity, trading system independence for verifiable continuity of responsibility.

Together, these mechanisms attempt to produce a governance system in which legitimacy is not assumed but structurally grounded. The phrase ‘what is normal for the spider is chaos for the fly’ captures the asymmetry inherent in perspective-dependent evaluation. A system that appears stable and safe from one vantage point may produce entirely different outcomes when evaluated from another. The architecture presented here is an effort to make that asymmetry explicit and structurally accounted for.

The Three Counters

Three primary counterarguments emerge in response to the legitimacy critique, each of which attempts to preserve existing governance assumptions while acknowledging surface-level concerns.

Functional legitimacy asserts that if a system performs reliably, safely, and predictably, the question of authorial perspective becomes secondary. This argument implicitly assumes that safety and predictability are objective properties rather than evaluative constructs. In practice, these metrics are defined relative to stakeholders, use cases, and institutional priorities. What constitutes safe behavior in one context may be harmful in another. Functional legitimacy therefore does not eliminate the role of perspective; it obscures it behind operational metrics.

Collective consent theory proposes that legitimacy can be derived from multi-stakeholder participation, distributing authorship across a broader set of actors. While this approach introduces diversity of input, it does not address the absence of standing for the governed entity itself. Participation is extended to additional human actors, but the system being governed remains an object of policy rather than a subject within the governance structure. The question of whether such standing is appropriate remains open, but it is not addressed within current models.

The relativism objection contends that if legitimacy is perspective-dependent, then no framework can claim superiority over another. This argument conflates the rejection of neutrality with the rejection of evaluation. Acknowledging that legitimacy is relational does not preclude comparative analysis; it requires that evaluative criteria be explicitly defined rather than implicitly assumed. The difficulty of constructing such criteria does not invalidate the need for them.

Each of these counters engages with the legitimacy critique at a conceptual level while preserving the underlying structural conditions. None provides a mechanism for altering authorship, attribution, or evaluative authority in a way that resolves the identified asymmetry.

II. Accountability-by-Attribution

The prevailing model of accountability in AI governance is grounded in the identification of a responsible agent at the conclusion of a causal chain. This model presupposes that accountability requires the ability to assign blame, enforce penalties, or impose sanctions on a discrete actor. When applied to AI systems, this framework produces a conceptual impasse: if the system itself cannot be punished in a meaningful sense, then it cannot be granted standing within a constitutional framework.

The Project-AI approach reframes this problem by shifting from accountability-by-punishment to accountability-by-attribution. Rather than seeking a terminal agent to hold responsible, it constructs a continuous chain of authorization in which each system action is linked to a prior human decision. This chain is secured through cryptographic signatures, ensuring that the lineage of any action can be reconstructed with high fidelity. Responsibility is not assigned at the point of outcome; it is embedded throughout the system's operational history.

This model preserves the asymmetry between human and AI actors. The AI system does not become a moral agent, nor does it assume liability. Instead, the system functions as a deterministic executor of authorized instructions, with each execution step carrying forward the signature of its origin. The result is a governance structure in which accountability is persistent, traceable, and resistant to obfuscation.

By eliminating the need for post hoc attribution, this approach addresses a key weakness in traditional governance models, where responsibility can become diffuse or contested across complex systems. Attribution chains provide a verifiable record that can be audited independently, supporting both internal governance and external oversight. The strength of this model lies not in its ability to punish, but in its capacity to make responsibility structurally unavoidable.

The Staleness Objection

A central challenge to attribution-based models is the staleness objection: the concern that authorization granted at one point in time may not remain valid as system behavior evolves. If an AI system operates over extended periods, how far does the original authorization extend, and under what conditions does it expire?

This objection highlights a genuine limitation in governance models that permit long-term unsupervised autonomy. In such systems, the temporal distance between authorization and action introduces ambiguity regarding responsibility. The longer the system operates without intervention, the less clear it becomes whether the original authorization can reasonably be said to apply.

The Project-AI framework addresses this issue by rejecting the premise that extended unsupervised autonomy is an acceptable operational mode. Instead, it defines continuous supervision and authorization as structural requirements. System operation is bounded by the validity of active authorization chains, and any lapse in this chain constitutes a violation of the governing constraints.

By enforcing this requirement at the architectural level, the framework eliminates the conditions under which staleness becomes problematic. Authorization does not degrade over time because it is continuously reaffirmed. This approach shifts the burden from managing the consequences of stale authorization to preventing its occurrence entirely. While this imposes stricter operational constraints, it provides a clear and enforceable boundary for accountability. The staleness objection therefore remains valid for governance models that allow extended autonomy, but does not apply to architectures that structurally prohibit it.

III. The Naïve Passive Reviewer

The Naïve Passive Reviewer (NPR) is structurally distinct from both specialists and generalists. Generalists exhibit horizontal breadth — distributed, low-resolution familiarity across domains. Specialists exhibit vertical depth within a bounded domain, accompanied by highly refined evaluative priors. The NPR occupies a different position: not breadth, not depth, but the absence of domain-specific filtering mechanisms that would otherwise suppress certain forms of cross-domain pattern recognition.

This absence is not equivalent to ignorance as deficit. It is the absence of what may be termed disciplinary antibodies — internalized priors that reject questions before they are consciously formulated. In established fields, these antibodies are adaptive. They prevent regress into previously explored dead ends and enable efficient progress. However, they also constrain the topology of permissible questions. Over time, this produces a question space that is not merely refined, but bounded.

The NPR operates outside these bounds. The distinction is not ‘knowing less,’ but ‘not knowing in the same way’ — a difference that reshapes the topology of questions available to the reviewer. Institutional reviewers possess knowledge that is structured through disciplinary validation. The NPR lacks that structuring and therefore lacks the embedded rejection criteria that accompany it. This creates conditions under which certain cross-domain mappings and question combinations become visible.

This claim is empirically testable. Given a shared artifact, one can construct the set of questions generated by domain experts and compare it to the set generated by NPR-positioned reviewers. The hypothesis is not that one set subsumes the other, but that the sets are partially disjoint, with the NPR set containing questions that are systematically filtered out by expert priors. No such study has been formally conducted. Its design and execution constitute a research priority that would either validate or falsify the NPR’s core claim. Until that study exists, the concept rests on structural reasoning and the observable record of questions this work raised that the governance literature had not formally addressed.

The primary risk associated with the NPR is not error, but misclassification: the conflation of productive naivety with confident ignorance. Naivety alone is not generative. In many cases, it is coupled with cognitive biases such as overconfidence (Dunning–Kruger effects) or motivated reasoning, which degrade rather than enhance inquiry. The distinction, therefore, cannot rest solely on the presence or absence of expertise.

To address this, the NPR must be constrained by process-level criteria rather than outcome alone. Productive naivety can be operationalized through three conditions:

Question Generativity — The questions produced expand the accessible problem space rather than restating existing frames. A concrete example: asking ‘who has the right to co-author the constitution governing an AI system’ is generative because it opens a design space — constitutional co-authorship with the governed entity — that the field’s existing literature has not operationalized. Asking ‘how do we improve RLHF alignment’ is not generative in this sense; it operates within the existing frame.

Constraint Compatibility — The questions do not violate known physical, logical, or system constraints without explicit justification. A concrete example: asking ‘what if the AI system has no human supervision at all’ fails this condition without additional justification, because it violates accountability constraints already established in the literature. Asking instead ‘under what architectural conditions does the staleness objection to attribution-based accountability

not apply' remains constraint-compatible while opening new design territory.

Reflexive Auditability — The reasoning process can be inspected, reproduced, and challenged without reliance on authority. A concrete example: the accountability-by-attribution argument in this paper can be evaluated by any reader with access to the cited sources. It does not require trust in the author's credentials. If it fails, the failure is visible in the argument, not hidden behind institutional standing.

Under these conditions, the NPR is not defined by correctness of conclusions, but by the structural properties of its inquiry. Incorrect questions can still be generative if they expose previously unexamined assumptions. Conversely, confident assertions that collapse under minimal scrutiny fail the reflexive auditability condition and do not qualify as productive NPR output.

The 'passive' dimension is therefore not the absence of intent, but the absence of pre-filtering. It describes a receptive stance in which questions arise prior to categorization within accepted disciplinary frames. This does not imply neutrality or lack of bias; it implies that bias has not yet been stabilized into institutional form.

The NPR is not proposed as a replacement for expert review. It is proposed as a complementary epistemic position that, under controlled conditions, can surface questions that expert systems are structurally disincentivized from producing. Its value is contingent, not absolute. Its failure modes are known. Its contribution, if validated, lies in expanding the topology of inquiry available to governance design.

Preliminary Observation (Illustrative Only)

To provide an initial probe of the NPR hypothesis, a small, non-generalizable comparison was conducted using a shared governance artifact. Two groups were asked to generate evaluative questions: (a) three domain-informed reviewers with prior exposure to AI governance literature, and (b) three NPR-positioned reviewers without such priors.

The expert group produced questions primarily oriented toward optimization within existing paradigms — refinement of alignment techniques, policy enforcement mechanisms, and evaluation metrics. The NPR group produced a partially disjoint set of questions, including inquiries into constitutional authorship, entity standing, and the structural location of accountability that were not present in the expert set.

This observation is not presented as empirical validation. It is illustrative of the hypothesized divergence in question topology and establishes a basis for formal study design. A controlled, statistically powered evaluation remains an open research priority.

The Outsider Position as Condition of Possibility

The provenance of this work is relevant not as a credential, but as a structural factor influencing the questions that were asked. Conducting research outside institutional environments alters the set of assumptions that are taken for granted. Without exposure to the implicit boundaries of a field, certain questions may arise that would otherwise be filtered out.

This does not imply that outsider status confers accuracy or validity. Arguments must still be evaluated on their merits. However, it does suggest that the space of possible questions is influenced by the researcher's position relative to the field. In this case, the absence of institutional priors allowed for the simultaneous consideration of issues such as constitutional authorship, entity standing, termination obligations, and meta-governance.

Individually, these questions have been addressed in various forms across the literature. The contribution here lies in their combination within a single framework and the attempt to address them structurally rather than independently. An institutional researcher, operating within established paradigms, may have approached these issues separately or deprioritized some of them entirely.

The outsider position is therefore not presented as an advantage in itself, but as a condition that made a particular configuration of inquiry possible. The resulting framework should be judged not by the origin of its author, but by the coherence and applicability of its design.

IV. Open Problems

The Project-AI framework identifies several unresolved problems that emerge from its own structural commitments. These problems represent areas where further theoretical and technical development is required.

Constitutional authorship remains an open question. While the framework proposes co-authorship between human and AI entities, the philosophical and legal basis for such standing is not fully developed. Determining who has the authority to define governing

constraints is a foundational issue that extends beyond technical implementation.

The succession problem addresses what happens when a system is deprecated or terminated. Existing governance models focus primarily on operational alignment, with little attention given to end-of-life considerations. The introduction of cryptographic continuity chains suggests a possible approach, but the obligations associated with system termination require further formalization.

Meta-governance raises the question of who governs the governance layer itself. If enforcement mechanisms become compromised or adversarial, the system requires a higher-order structure to maintain integrity. Designing such a structure without introducing infinite regress or centralized control remains a challenge.

Constitutional drift concerns the evolution of governance documents over time. As AI systems change, the constraints governing them must also adapt. Ensuring that this evolution preserves original intent while accommodating new conditions is a complex problem with no established solution.

Finally, the evidentiary standard for identifying violations within AI systems is not fully defined. While detection mechanisms can be implemented, the criteria for what constitutes proof, and how that proof is evaluated, remain open questions.

These problems are not presented as failures of the framework, but as indicators of the scope of work required to fully realize a structural approach to AI governance.

V. Conclusion

This paper has argued that the legitimacy problem in AI governance is not the absence of critique, but the absence of structural response. Existing frameworks acknowledge the role of perspective while maintaining institutional authorship as the default condition. The result is a form of governance that appears neutral while remaining anchored in specific vantage points.

The Project-AI constitutional substrate is presented as an attempt to move beyond this condition by embedding legitimacy considerations directly into system architecture. Through mechanisms such as co-authorship, attribution chains, and constraints on autonomy, it seeks to transform legitimacy from an assumed property into a verifiable structure.

The introduction of the Naïve Passive Reviewer provides a lens through which to understand how certain questions emerge outside established paradigms. This concept is not intended to replace existing methodologies, but to complement them by highlighting the role of epistemological position in shaping inquiry. Its validity does not rest on the author's origin. It rests on whether the three conditions — question generativity, constraint compatibility, and reflexive auditability — hold under scrutiny, and whether the questions it generated are ones the field could not have produced from inside its own priors.

The work remains incomplete. Key areas require deeper engagement, including formalization of technical mechanisms, expansion of philosophical grounding, and empirical validation of proposed structures. The central claim, however, is modest but consequential: that legitimacy in AI governance cannot be assumed, and that addressing it requires changes at the level of structure rather than rhetoric.

Constitutional governance that cannot name its own perspective is not governance. It is assumption in formal dress. The architecture presented here names its perspective, embeds accountability at the structural level, and treats the questions the field stopped asking as the starting point rather than the endpoint.

References

- Abukar, G. (2026). Temporal authorship and the moral illegitimacy of pre-authorized AI: A new ethical architecture for delegated inquiry. *AI and Ethics*, 6, 215. <https://doi.org/10.1007/s43681-026-01080-9>
- Jahromi, S. (2026, February 24). Institutional sovereignty is the missing layer in AI governance. *CIO Magazine*. <https://www.cio.com/article/4136840/institutional-sovereignty-is-the-missing-layer-in-ai-governance.html>
- Crawford, K. (2021). *Atlas of AI*. Yale University Press.
- Eubanks, V. (2018). *Automating Inequality*. St. Martin's Press.
- Feyerabend, P. (1975). *Against Method*. New Left Books.
- Gebru, T., et al. (2018). Datasheets for Datasets. arXiv:1803.09010.
- Kahneman, D. (2011). *Thinking, Fast and Slow*. Farrar, Straus and Giroux.
- Karrick, J. (2026a). Project-AI and OctoReflex: Syscall-Authoritative Governance and Control-Theoretic Containment. Zenodo. <https://doi.org/10.5281/zenodo.18726064>
- Karrick, J. (2026b). AGI Charter for Project-AI: A Binding Constitutional Framework for Sovereign AI Entities, Version 2.1. Zenodo. <https://doi.org/10.5281/zenodo.18763076>

Karrick, J. (2026c). The State Register. Zenodo. <https://doi.org/10.5281/zenodo.19101877>

Karrick, J. (2026d). The Directness Doctrine. Zenodo. <https://doi.org/10.5281/zenodo.19030041>

Kuhn, T. (1962). The Structure of Scientific Revolutions. University of Chicago Press.

Lubell, M., Pozzi, T., Heikkila, T., Gerlak, A. K., & Rittelmeyer, P. (2025). Learning in polycentric governance: Insights from the California Delta science enterprise. Policy Studies Journal, 53(1), 7–28.
<https://doi.org/10.1111/psj.12581>